

Final Exam for CS 165 (Fall 2021)

Dec 7, 2021

Instructions (write down your name): Steven Truong

* Be brief in your answers. You will be graded for correctness, not on the length of your answers.

I. (1 point *10) Answer the following multiple choice questions (one or more correct answers).

1. What can be said about control flow hijacking vulnerabilities? **__a, c__**

a. They are inherently difficult to eradicate as long as we can continue to use unsafe programming languages such as C.

b. C++ will not allow control flow hijacking because it is object-oriented.

c. Format string vulnerability is an example vulnerability type that can lead to control flow hijacking.

2. Which of the following are true about attack surface? **__a,b,c__**

b. To a remote attacker, the attack surface a system should include all TCP and UDP ports listening on the system.

a. Entry points of a program can be defined as any external inputs consumed through syscalls, e.g., files.

c. Defining attack surface requires reasoning about an adversary's capabilities (threat model).

3. What role does the access control play in improving security? **__b, c__**

a. It eliminates buffer overflow vulnerabilities.

b. It reduces attack surface.

c. It guarantees that an attacker is unable to access certain resources.

4. Why is security hard to achieve in general? **a, b, c__**

a. Because defenses can be bypassed.

b. Because software is complex and hard to avoid loopholes.

c. Because computer systems and networks were designed initially without much consideration of security and it is now difficult to retrofit security (due to backward compatibility).

5. Which of the following are good applications/analogies of security policies? **__a,c__**

a. "A king does not eat anything prepared by untrusted people" is an application of the Biba policy (the integrity policy).

b. "A web server should not process trusted input from arbitrary users" is an application of the BLP policy (the secrecy policy).

c. “A mail server should not have read from a file that can be modified by a regular user” is an application of the Biba policy (the integrity policy).

6. Which of the following are true about information flow analysis? **__a, b__**

- a. It is possible to use information flow to describe security policies.
- b. Information flow can be viewed as a building block for the BLP policy (secrecy policy).
- c. Information flow can be viewed as a building block for the Biba policy (integrity policy).

7. What are the pros and cons between static and dynamic analysis for vulnerability discovery?

__a, b__

- a. Static analysis has the advantage of “covering” all program paths that are not necessarily exercised dynamically.
- b. Static analysis is more flexible because it can start the analysis abstractly from any function (e.g, not necessarily the main function).
- c. Dynamic analysis is effective only when the program is smaller.

8. Network firewalls and intrusion detection systems can be deployed on the host as well. What are the reasons that the network version be more desirable? **__b, c__**

- a. Host-based IDS has significant maintenance and runtime overhead.
- b. Network-based IDS has the advantage of covering more hosts.
- c. Network-based IDS has a global view of the network and therefore in principle more knowledge on whether certain attacks (e.g., DDoS) are happening.

9. Which of the following are true regarding vulnerability research? **__a, b__**

- a. Fuzzing is highly effective in finding bugs but those bugs are not necessarily exploitable vulnerabilities.
- b. One can apply either static analysis or dynamic analysis to find bugs.
- c. Exploiting a vulnerability is always easier than finding the vulnerability in the first place.

10. Patching looks like fairly straightforward process to fix bugs and vulnerabilities. Why would it be considered a problem? **__b, c__**

- a. IoT devices may not receive patches because devices may outlive the lifespan of manufacturers.
- b. The constant stream of new bugs being discovered can be overwhelming for developers to fix.
- c. In a complex ecosystem, if a bug is found in the upstream, there can be multiple layers/hops of patch propagation, e.g., from Linux mainline all the way down to Android devices.

II (1.5 points) An important security principle is called defense-in-depth. It advocates multiple layers of defenses as it will still provide some protection even if one defense layer is bypassed. Can you explain how Android applies this principle to defend against malware? Please give at least three major defenses employed by Google.

- 1) **Apps have to be audited before being placed in the google play store to make sure it's safe**
- 2) **Android users don't have root access by default**
- 3) **Google play protect which acts like an antivirus**

III. (2 points) One principle of secure programming is called “fail-safe”, where the system or program should maintain security even when it encounters failure or errors. Consider the following general code for allowing access to a resource:

```
int ret = IsAccessAllowed(...);  
  
if (ret == ERROR_ACCESS_DENIED) {  
    // Security check failed.  
    // Inform user that access is denied.  
  
} else {  
    // Security check OK.  
  
}
```

a. Explain the security flaw in this program from the fail-safe perspective. Hint: no race condition is involved. (1 point)

The program only checks if the ret is ERROR_ACCESS_DENIED. It doesn't check for other types of errors. It also doesn't use an exact condition to check if ret is ACCESS_ALLOWED. As a result, if ret is an error besides ERROR_ACCESS_DENIED, then security check would not fail because execution goes to the else block.

b. Rewrite the code to avoid the flaw. (1 point)

```
int ret = IsAccessAllowed(...);

if (ret == ERROR_ACCESS_DENIED) {

    // Security check failed.

    // Inform user that access is denied.

} else if (ret == ACCESS_ALLOWED) {

    // Security check OK.

}

else {

    // Security check failed.

    // Inform user that access is denied.

}
```

IV. Below is an interactive program where the main function runs in a while loop and processes commands through command line one after another, i.e., doCmd() is called multiple times with a *cmd* value and *arg* string decoded from the input. Assuming the decoding is done correctly and safely, can you spot any bugs with the program and address them? (4 points)

```
int *data_ptr; // global variable
int size; // global variable
char doCmd (int cmd, char* arg) {
    if (cmd == CREATE) {
        size = atoi(arg);
        data_ptr = malloc(size);
    }
    if (cmd == TEST) {
        for (unsigned int i = 0; i < size; ++i) {
            printf("%d\n", data_ptr[i]);
        }
    }
}
```

```

        if(cmd == UPDATE_SIZE) {
            size = atoi(arg);
        }
    }
}

```

1. Describe clearly four bad things that could happen.

(2 points)

- 1) If cmd is CREATE and arg is a number bigger than the maximum integer size, then size will become a large negative number because it overflowed.
- 2) data_ptr is uninitialized, so if the very first command that the program runs is TEST, then there may be a segmentation fault because it accesses an uninitialized data_ptr.
- 3) The if statement for when cmd is UPDATE_SIZE does not update the size of data_ptr.
- 4) One can use the cmd UPDATE_SIZE to make the size bigger, then use the cmd TEST to print contents of data_ptr that are out of bounds.

2. Describe how you might use static analysis in terms of control flow and data flow to discover them.

(2 points)

I can use control flow analysis to discover uninitialized pointers that are dereferenced and look for accesses to data_ptr that are out of bounds. Similar to what I did in project 4 part 2, I would use two std::vector in parallel. One vector can store the identifier for the pointer and the other vector stores its state: whether it is initialized or not. By storing this information, I can determine whether an uninitialized pointer is dereferenced.

For data flow analysis, I can compare the buffer size of data_ptr and the stored integer in size to make sure they match. This is similar to project 4 part 1 where I checked for memory overflow by comparing buffer size.

V. Thinking like an attacker (1 point * 6)

1. In project 2, we were able to crack programs easily by modifying its intended behaviors at the binary level. What is the fundamental reason that we can crack a program? Alternatively, would you be able to design a program that prevents a hacker from cracking it (or at least make it substantially more difficult)? Answering either of the question properly would get you the full point.

Programs can be cracked because there is no general defense mechanism that protects an executable binary from being modified.

2. In Intrusion Detection Systems (IDS), one main challenge is to keep up with the new malware samples that appear every day. One specific difficulty is to collect such malware samples. To solve this problem, one solution is to deploy “honeypots” where virtual machines are intentionally configured to be vulnerable and publicized to attract attackers. Once the machine is compromised, all new executables downloaded onto the machine will be considered malware samples. This whole process can be fully automated for malware sample collection. What is the potential risk of this solution?

Since the virtual machine is configured to be vulnerable, the can attacker can exploit a vulnerability to find a way to access the host's network. Attackers may also exploit a vulnerability within the VM software to gain access to the host machine.

3. Modern operating systems have already provided the basic isolation mechanism --- each process has its own independent virtual address space. By design, one process is not allowed to read/write the memory of another process. However, there are cases where a single process can have multiple threads which share the same address space. In such cases, what techniques do you think are suitable to solve this problem and why? (please use the knowledge given in the lectures only).

One way to solve this problem is to use an isolation mechanism. This can be done by allowing each thread a certain address range for which they can

read/write within the shared address space. A reference monitor can enforce this policy by checking if the reads/writes are within a thread's allowed range.

4. For project 3, explain how you might use the vulnerability of the binary to cause trouble of other teams without being easily noticed.

I can spawn a shell with admin privileges with ROP, giving me write access to the admin directory. Then I can create a lot of directories with long names. This will cause trouble for other teams because it will be hard for them to look for the directory they created when they use ls. I can also delete directories created by other teams, which would cause trouble.

5. Google operates a fuzzer that continuously fuzzes Linux kernels and on average finds 1000+ bugs in a single year. They even publish all the bug reports in real time publicly. Why do you think they do this? Wouldn't it allow attackers to take an early look at these bugs and try to exploit them? Please construct an argument that this is in some way helpful for security.

By publishing the bugs, it allows anybody to look at the bugs and attempt to fix them. This is better for security because it can help bugs get fixed faster. Moreover, most bugs are benign and don't really pose any security risks.

6. Linux is known for intentionally obfuscating the nature of a kernel change, i.e., git commit. This means that even if a commit is to fix a serious security vulnerability, the commit message would not indicate as such. This concealing of information is a direct contradiction to what Google practices making everything (e.g., bug reports) public. The intention is to avoid helping attackers identify which patches fix important security vulnerabilities and develop exploits for unpatched systems. Can you come up with an argument this pro-security practice can actually end up hurting security?

If there is no commit message then one can't identify the version that might have made the kernel less secure by introducing a vulnerability. This hurts security because updating to the latest kernel version can make a system more vulnerable in this case.

VI (1.5 points). In the 2016's off-path TCP attack, the presence of a connection between two arbitrary hosts can be inferred through side channels by a completely off-path attack. Please answer the following questions.

1. We know side channels always require some shared resources between an attacker and victim. What is the shared resource in this case? (0.5 point)

The shared resource is the global rate limit on a server, which limits the number of challenge ACKs to 100 per second.

2. Compared to the off-path attack in 2012 (where malware on the client is used to assist the attack), it generally takes a lot longer to complete the attack proposed in 2016. What makes it more difficult? (1 point)

It's more difficult because the port number, sequence number, and ack number has to be guessed. But for the 2012 off-path attack, the malware can reveal this information so there is no guessing needed.

VII (1.5 points). We introduced a side-channel-enabled DNS cache poisoning which is a serious attack that can cause incorrect mappings between domain names and IP addresses cached in a DNS resolver. The fundamental flaw is that the source port used in a DNS request can be learned by an off-path attacker through a side channel. Since we learned that any side channel would involve some form of shared resource, can we simply eliminate the shared resources such that the side channel can no longer work? Would there be any potential side effects? If so, what would be an alternative to eliminating the sharing?

If the shared resource was eliminated, then there would be no global rate limit. The side effect of this would be denial of service attacks that send a lot of packets to a server to overwhelm the cpu. The alternative would be to keep the shared resource, but randomize it so that an attacker can't infer anything.

VIII. (6 points) Design question: Consider the following alternative design to defend against buffer overflows and control flow hijack in general:

a. Buffer overflow seems to be possible because of the fact that the buffer and the stack grow in opposite directions (stack growing from top to bottom while buffer grow from bottom to top). Imagine that if the CPU and OS are redesigned so that stack now grows from bottom to top (same as buffer). Would this prevent any control flow hijacking attacks completely? If your answer is yes, state your reasoning. If not, give one counterexample. (2 points)

Yes, because a buffer overflow can no longer overwrite the return address since it grows the same direction the stack grows.

b. To prevent return address from being overwritten, we design a specialized hardware register %SRA (Secure Return Address) that always remembers the latest saved return address on the stack. Whenever the call instruction is executed, besides pushing the return address on the stack, the CPU now automatically saves the return address in %SRA as well. When a function returns, the CPU will automatically compare the return address saved on the stack and the value in %SRA. If they differ, then a buffer overflow attack is detected. If they are the same, then the return address will be moved to %EIP. At the same time, the saved return address for the previous stack frame will be located and saved to %SRA. Is this design secure or not? If not, please elaborate your reasoning. (2 points)

No, because %SRA has to be saved on the stack later on when a function calls another function. The %SRA saved on the stack can be overwritten with a buffer overflow.

c. Consider another design: for every call instruction, the program pushes an encrypted return address (along with a signature) onto the stack using a secret key stored in a specialized register (which is initialized to a random value every time the program runs). When the return instruction is about to be executed, the program decrypts the return address on the stack using the same secret key stored in the register before actually returning. If the attacker tampers with the return address, it will not match the

signature. In other words, the program will catch the tampering of saved return address on the stack. Can this really defend against tampering of saved return address? If so, can you explain how this can be applied to protecting the function pointers? If not, please elaborate your reasoning. (2 points)

This design can defend against tampering of a saved return address. To apply this to function pointers, two encrypted addresses, representing a range that is legal for the function pointers, can be pushed onto the stack. The location of the function pointer can be compared with the saved address range on the stack.