

# CS165 – Computer Security

## Control Flow Defenses

Oct 14, 2021

# Agenda

Control Flow Hijacks



Common Hijacking Methods

- Buffer Overflows
- Exploits (shell code) Construction
- Integer Overflows
- Heap Overflows
- Format String Vulnerability



What's new since 2000



# What's new since 2000?

## **Assigned Reading:**

*Smashing the stack in 2011*

by Paul Makowski

<http://paulmakowski.wordpress.com/2011/01/25/smashing-the-stack-in-2011/>

# A lot has happened...

- Heap-based buffer overflows also common
- [not mentioned] fortified source by static analysis (e.g., gcc can sometimes replace strcpy by strcpy\_chk)

Additional materials:

- Canary (e.g. ProPolice in gcc)
- Data Execution Protection/No eXecute
- Address Space Layout Randomization

# A lot has happened...

- Heap-based buffer overflows also common
- [not mentioned] fortified source by static analysis (e.g., gcc can sometimes replace strcpy by strcpy\_chk)

Additional materials:

- Canary (e.g. ProPolice in gcc)
- Data Execution Protection/No eXecute
- Address Space Layout Randomization

```
alias gcc732='gcc -m32 -g3 -O1 -fverbose-asm -fno-omit-frame-pointer  
-mpreferred-stack-boundary=2 -fno-stack-protector -fno-pie -fno-PIC  
-D_FORTIFY_SOURCE=0'
```

# But little has changed...

Method to gain entry remains the same

- buffer overflows
- format strings

What's different is shellcode:



return-Oriented  
Programming

# Agenda

Control Flow Hijacks



Common Hijacking Methods

- Buffer Overflows
- Exploits (shell code) Construction
- Integer Overflows
- Heap Overflows
- Format String Vulnerability



What's new since 2000



# Reading list

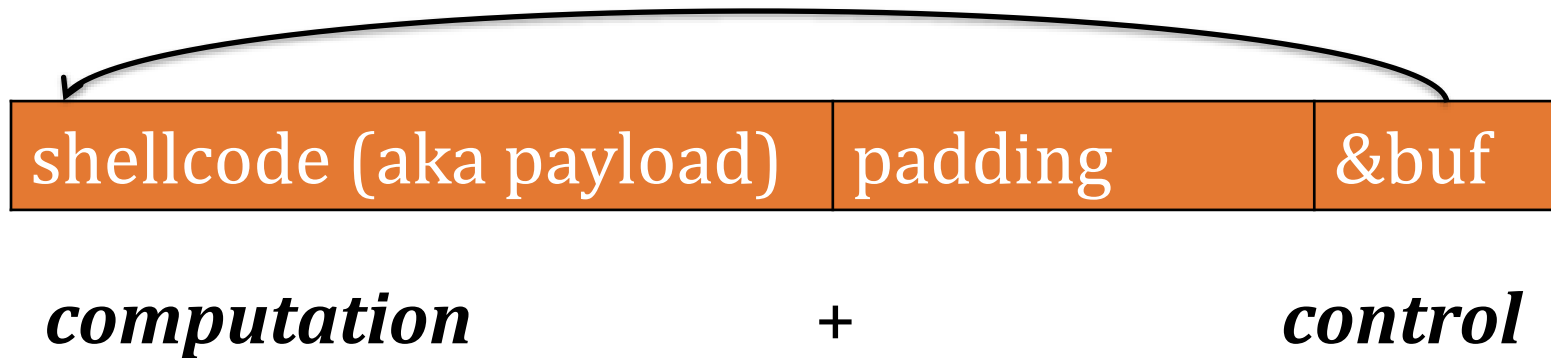
- “Smashing The Stack For Fun And Profit”
  - <http://www.phrack.org/issues.html?issue=49&id=14#article>
- “Exploit the format string vulnerabilities”
  - <http://www.utdallas.edu/~zhiqiang.lin/file/formatstring.pdf>
- Smashing the Stack in 2011
  - <http://paulmakowski.wordpress.com/2011/01/25/smashing-the-stack-in-2011/>



# Control flow hijack defenses

# Control Flow Hijack:

## Always control + computation



- code injection
- return-to-libc
- heap metadata overwrite
- return-oriented programming
- ...

Same principle,  
different  
mechanism

# Control Flow Hijacks

*... happen when an attacker gains control of*

# Control Flow Hijacks

*... happen when an attacker gains control of  
the instruction pointer.*

Two common hijack methods:

- buffer overflows
- format string attacks

# Control Flow Hijack Defenses

# Control Flow Hijack Defenses

**Bugs are the root cause of hijacks  
(hard/costly)!**

- Find bugs with analysis tools
- Prove program correctness

# Control Flow Hijack Defenses

**Bugs are the root cause of hijacks  
(hard/costly)!**

- Find bugs with analysis tools
- Prove program correctness

**Mitigation Techniques (simple/cheap):**

- Canaries
- Data Execution Prevention/No eXecute
- Address Space Layout Randomization

# Proposed Defense Scorecard

Aspect	Defense
Performance	• Smaller impact is better
Deployment	• Can everyone easily use it?
Compatibility	• Doesn't break libraries
Safety Guarantee	• Completely secure to easy-to-bypass




# Agenda

Canary / Stack Cookies

Data Execution Prevention (DEP)  
/No eXecute (NX)

Address Space Layout  
Randomization (ASLR)

# Agenda



Canary / Stack Cookies

Data Execution Prevention (DEP)  
/No eXecute (NX)

Address Space Layout  
Randomization (ASLR)

*Wikipedia*: “the historic practice of using canaries in coal mines, since they would be affected by toxic gases earlier than the miners, thus providing a biological warning system.”

## Canary / Stack Cookies

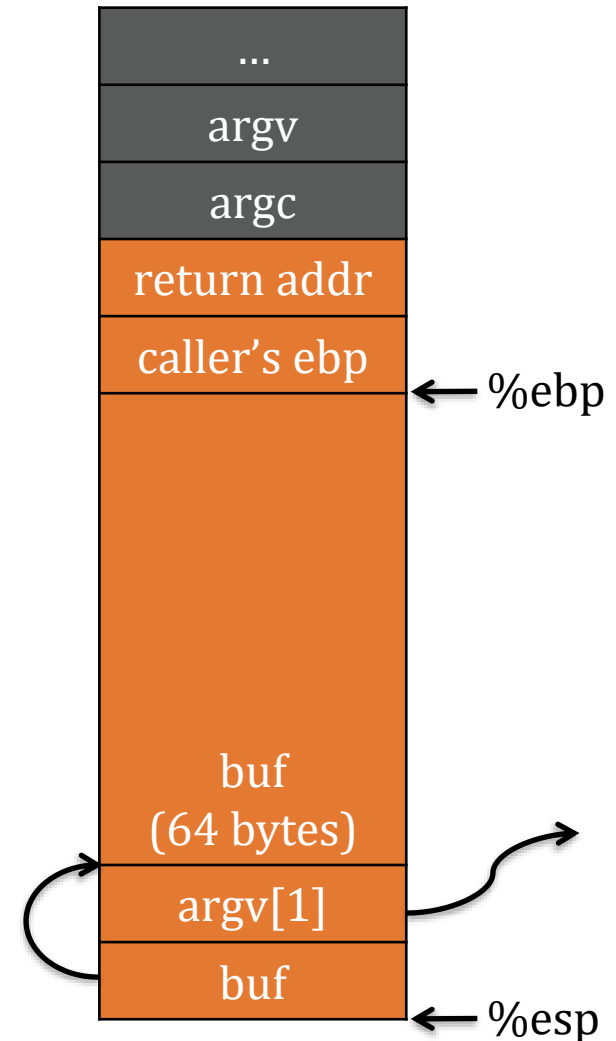


# “A”x68 . “\xEF\xBE\xAD\xDE”

```
#include<string.h>
int main(int argc, char **argv) {
    char buf[64];
    strcpy(buf, argv[1]);
}
```

Dump of assembler code for function main:

```
0x080483e4 <+0>: push    %ebp
0x080483e5 <+1>: mov     %esp,%ebp
0x080483e7 <+3>: sub     $72,%esp
0x080483ea <+6>: mov     12(%ebp),%eax
0x080483ed <+9>: mov     4(%eax),%eax
0x080483f0 <+12>: mov     %eax,4(%esp)
0x080483f4 <+16>: lea     -64(%ebp),%eax
0x080483f7 <+19>: mov     %eax,(%esp)
0x080483fa <+22>: call    0x8048300 <strcpy@plt>
0x080483ff <+27>: leave
0x08048400 <+28>: ret
```

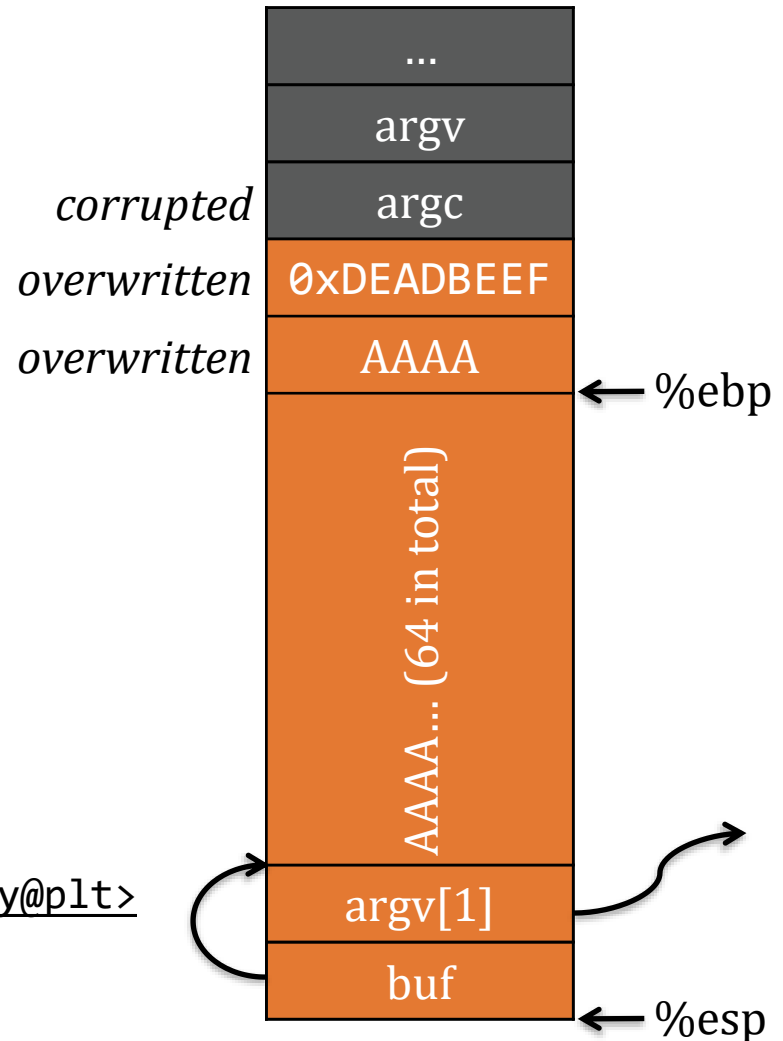


# “A”x68 . “\xEF\xBE\xAD\xDE”

```
#include<string.h>
int main(int argc, char **argv) {
    char buf[64];
    strcpy(buf, argv[1]);
}
```

Dump of assembler code for function main:

```
0x080483e4 <+0>: push    %ebp
0x080483e5 <+1>: mov     %esp,%ebp
0x080483e7 <+3>: sub     $72,%esp
0x080483ea <+6>: mov     12(%ebp),%eax
0x080483ed <+9>: mov     4(%eax),%eax
0x080483f0 <+12>: mov     %eax,4(%esp)
0x080483f4 <+16>: lea     -64(%ebp),%eax
0x080483f7 <+19>: mov     %eax,(%esp)
0x080483fa <+22>: call    0x8048300 <strcpy@plt>
0x080483ff <+27>: leave
0x08048400 <+28>: ret
```

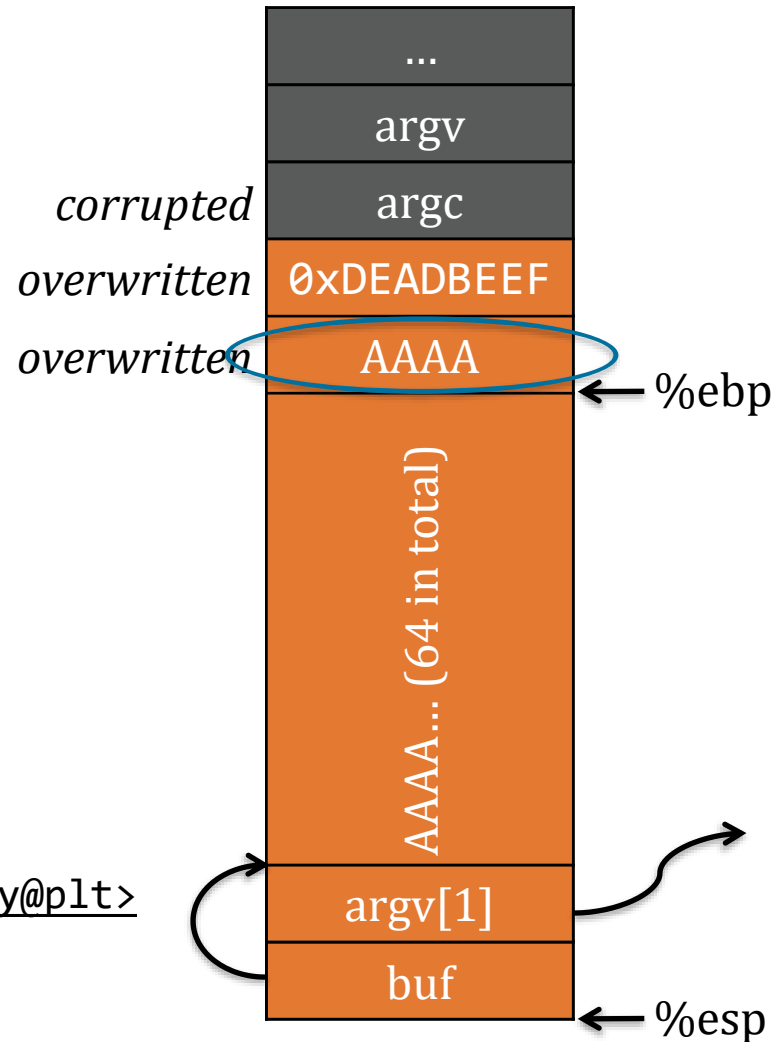


# “A”x68 . “\xEF\xBE\xAD\xDE”

```
#include<string.h>
int main(int argc, char **argv) {
    char buf[64];
    strcpy(buf, argv[1]);
}
```

Dump of assembler code for function main:


```
0x080483e4 <+0>: push    %ebp
0x080483e5 <+1>: mov     %esp,%ebp
0x080483e7 <+3>: sub     $72,%esp
0x080483ea <+6>: mov     12(%ebp),%eax
0x080483ed <+9>: mov     4(%eax),%eax
0x080483f0 <+12>: mov     %eax,4(%esp)
0x080483f4 <+16>: lea     -64(%ebp),%eax
0x080483f7 <+19>: mov     %eax,(%esp)
0x080483fa <+22>: call    0x8048300 <strcpy@plt>
0x080483ff <+27>: leave
0x08048400 <+28>: ret
```



# StackGuard [Cowen et al. 1998]

## Idea:

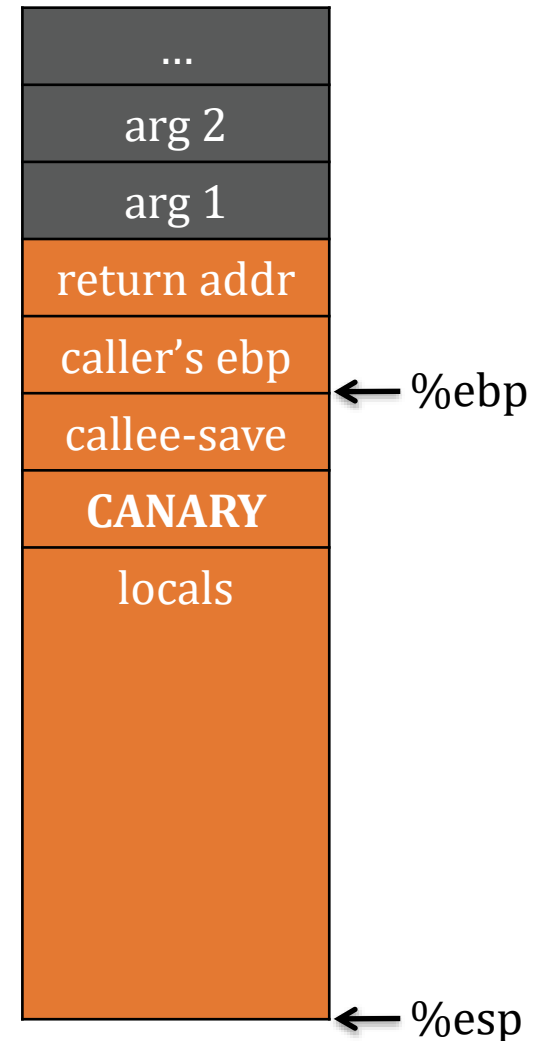
- prologue introduces a *canary word* between return addr and locals



# StackGuard [Cowen et al. 1998]

## Idea:

- prologue introduces a *canary word* between return addr and locals
- epilogue checks canary before function returns



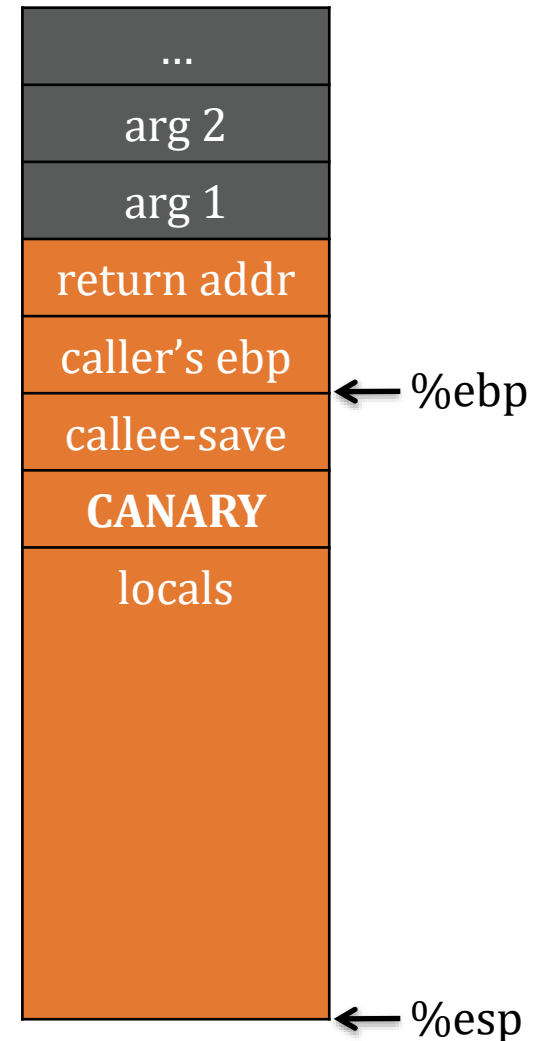


# StackGuard [Cowen et al. 1998]

## Idea:

- prologue introduces a ***canary word*** between return addr and locals
- epilogue checks canary before function returns

Wrong Canary => Overflow




# gcc Stack-Smashing Protector (ProPolice)

Dump of assembler code for function main:

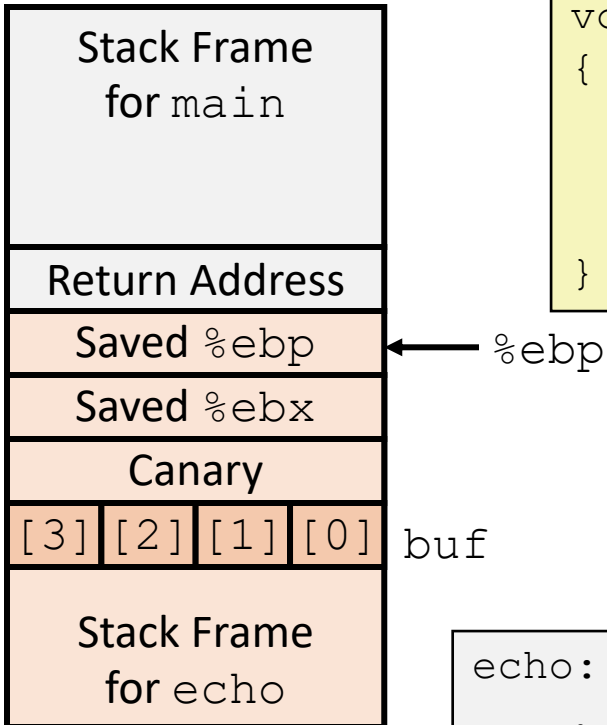
```
0x08048440 <+0>: push    %ebp
0x08048441 <+1>: mov     %esp,%ebp
0x08048443 <+3>: sub     $76,%esp
0x08048446 <+6>: mov     %gs:20,%eax
0x0804844c <+12>: mov     %eax,-4(%ebp)
0x0804844f <+15>: xor     %eax,%eax
0x08048451 <+17>: mov     12(%ebp),%eax
0x08048454 <+20>: mov     4(%eax),%eax
0x08048457 <+23>: mov     %eax,4(%esp)
0x0804845b <+27>: lea     -68(%ebp),%eax
0x0804845e <+30>: mov     %eax,(%esp)
0x08048461 <+33>: call    0x8048350 <strcpy@plt>
0x08048466 <+38>: mov     -4(%ebp),%edx
0x08048469 <+41>: xor     %gs:20,%edx
0x08048470 <+48>: je      0x8048477 <main+55>
0x08048472 <+50>: call    0x8048340 <__stack_chk_fail@plt>
0x08048477 <+55>: leave
0x08048478 <+56>: ret
```

**Compiled with v4.6.1:**  
`gcc -fstack-protector -O1 ...`



# Setting Up Canary

*Before call to gets*

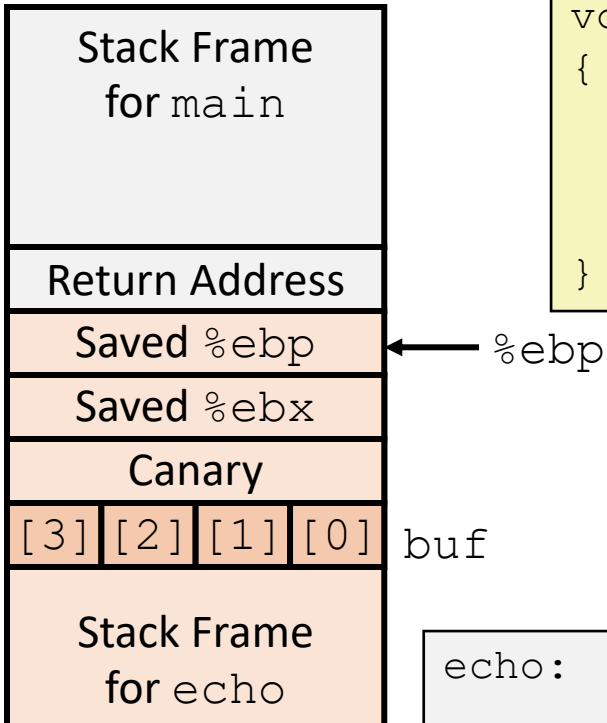


```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
echo:  
    . . .  
    movl    %gs:20, %eax    # Get canary  
    movl    %eax, -8(%ebp)  # Put on stack  
    xorl    %eax, %eax     # Erase canary  
    . . .
```

# Checking Canary

*Before call to gets*

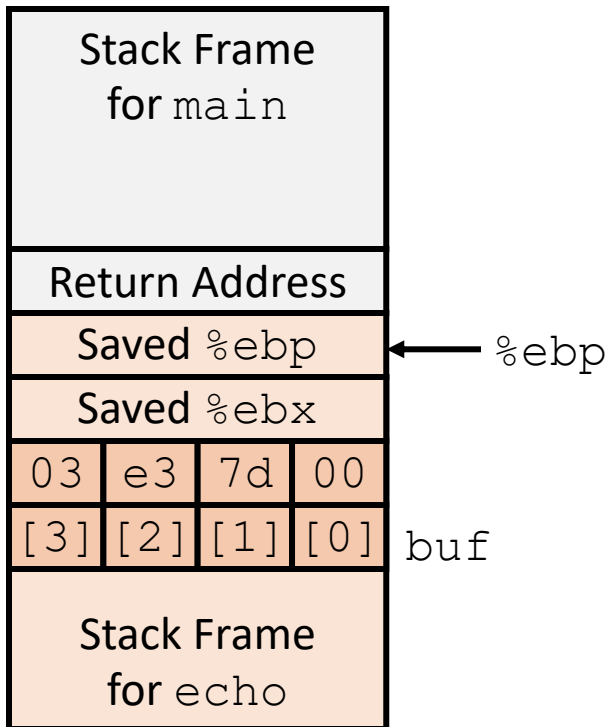


```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

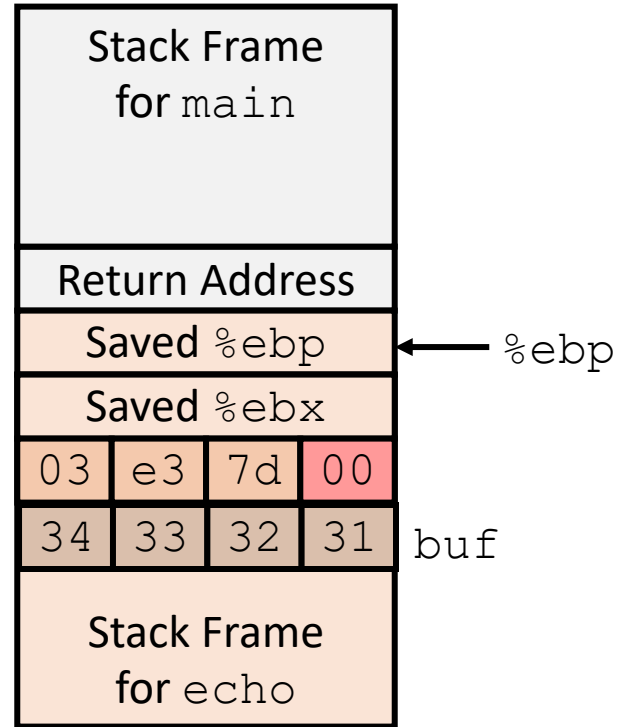
```
echo:
    . . .
    movl    -8(%ebp), %eax    # Retrieve from stack
    xorl    %gs:20, %eax     # Compare with Canary
    je      .L24              # Same: skip ahead
    call    __stack_chk_fail # ERROR
.L24:
    . . .
```

# Canary Example

*Before call to gets*



*Input 1234*



```
(gdb) break echo
(gdb) run
(gdb) stepi 3
(gdb) print /x *((unsigned *) $ebp - 2)
$1 = 0x3e37d00
```

Benign corruption!  
(allows programmers to make  
silent off-by-one errors)

# Canary should be **HARD** to Forge

- Random Canary
  - 4 random bytes chosen at load time
  - stored in a guarded page
  - need good randomness

# Canary Scorecard


Aspect	Canary
Performance	<ul style="list-style-type: none"><li>• several instructions per function</li><li>• time: a few percent on average</li><li>• size: can optimize away in safe functions (but see MS08-067 *)</li></ul>
Deployment	<ul style="list-style-type: none"><li>• recompile suffices; no code change</li></ul>
Compatibility	<ul style="list-style-type: none"><li>• perfect—invisible to outside</li></ul>
Safety Guarantee	<ul style="list-style-type: none"><li>• <i>not really...</i></li></ul>

\* <http://blogs.technet.com/b/srd/archive/2009/03/16/gs-cookie-protection-effectiveness-and-limitations.aspx>

# Bypass: Data Pointer Subterfuge

Overwrite a data pointer *first*...

```
int *ptr;  
char buf[64];  
memcpy(buf, user1);  
*ptr = user2;
```





# Canary Weakness

Check does ***not*** happen until epilogue...

- func ptr subterfuge
- C++ vtable hijack
- exception handler hijack
- ...

Code Examples:

[http://msdn.microsoft.com/en-us/library/aa290051\(v=vs.71\).aspx](http://msdn.microsoft.com/en-us/library/aa290051(v=vs.71).aspx)

VS 2003: /GS

# Function Pointer Subterfuge

Overwrite a function pointer to point to:

- program function (similar to ret2text)
- Other non-randomized functions

```
/*please call me!*/
int secret(char *input) { ... }


int chk_pwd(char *input) { ... }

int main(int argc, char *argv[]) {
    int (*ptr)(char *input);
    char buf[8];

    ptr = &chk_pwd;
    strncpy(buf, argv[1], 12);
    printf("[ ] Hello %s!\n", buf);

    (*ptr)(argv[2]);
}
```

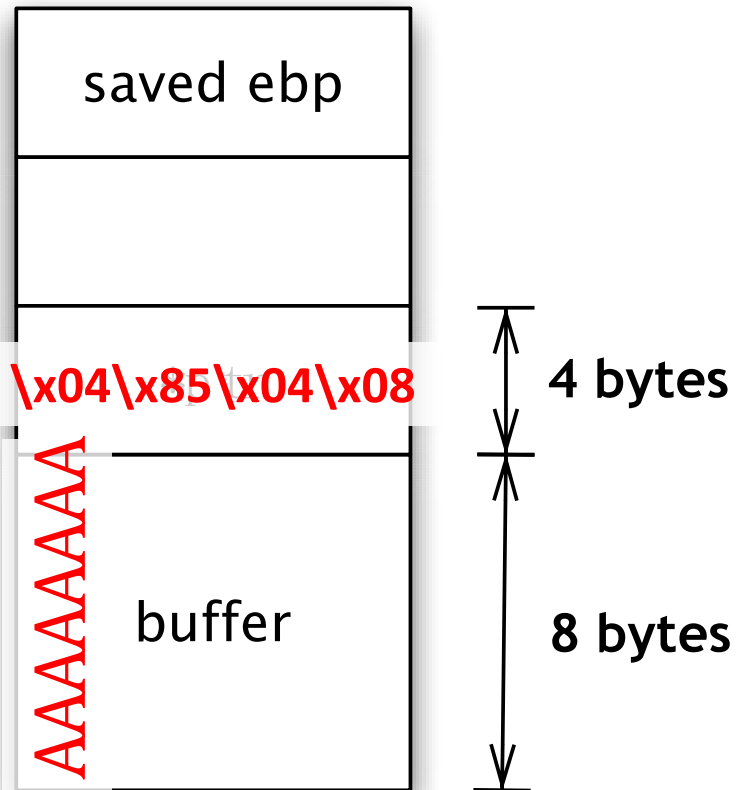
# Function Pointers



```
08048504 <secret>:  
8048504: 55  
8048505: 89 e5  
8048507: 83 ec 18  
804850a: 8b 45 08  
804850d: 89 44 24 04  
8048511: c7 04 24 30 87 04 08  
8048518: e8 df fe ff ff  
804851d: c7 44 24 0c 00 00 00
```

```
ptr = &chk_pwd;  
strncpy(buf, argv[1], 12);  
printf("[ Hello %s!\n", buf);  
  
(*ptr)(argv[2]);
```

# Function Pointers



```
08048504 <secret>
8048504: 55
8048505: 89 e5
8048507: 83 ec 18
804850a: 8b 45 08
804850d: 89 44 24 04
8048511: c7 04 24 30 87 04 08
8048518: e8 df fe ff ff
804851d: c7 44 24 0c 00 00 00
```

```
ptr = &chk_pwd;
strncpy(buf, argv[1], 12);
printf("[ Hello %s!\n", buf);

(*ptr)(argv[2]);
```

# Canary Weakness

Check does ***not*** happen until epilogue...

- func ptr subterfuge } PointGuard
  - C++ vtable hijack
  - exception handler hijack } SafeSEH  
SEHOP
  - ...
- ProPolice  
puts arrays  
above others  
*when possible*

Code Examples:

[http://msdn.microsoft.com/en-us/library/aa290051\(v=vs.71\).aspx](http://msdn.microsoft.com/en-us/library/aa290051(v=vs.71).aspx)

VS 2003: /GS

# Canary Weakness

Check does ***not*** happen until epilogue...

- func ptr subterfuge } PointGuard
  - C++ vtable hijack
  - exception handler hijack } SafeSEH  
SEHOP
  - ...
- ProPolice  
puts arrays  
above others  
*when possible*
- struct is fixed;  
& what about heap?

Code Examples:

[http://msdn.microsoft.com/en-us/library/aa290051\(v=vs.71\).aspx](http://msdn.microsoft.com/en-us/library/aa290051(v=vs.71).aspx)

VS 2003: /GS

# Agenda

Canary / Stack Cookies



Data Execution Prevention (DEP)  
/No eXecute (NX)


Address Space Layout  
Randomization (ASLR)

# Agenda

Canary / Stack Cookies




Data Execution Prevention (DEP)  
/No eXecute (NX)




Address Space Layout  
Randomization (ASLR)




# How to defeat exploits?




# How to defeat exploits?



# How to defeat exploits?



# Data Execution Prevention



Mark stack as  
non-executable  
using NX bit

# Data Execution Prevention




# Data Execution Prevention



(still a Denial-of-Service attack!)

W ^ X



Each memory page is  
*exclusively* either  
writable *or* executable.

(still a Denial-of-Service attack!)

# DEP Scorecard

Aspect	Data Execution Prevention
Performance	<ul style="list-style-type: none"><li>• with hardware support: no impact</li><li>• otherwise: reported to be &lt;1% in PaX</li></ul>
Deployment	<ul style="list-style-type: none"><li>• kernel support (common on all platforms)</li><li>• modules opt-in (now enabled by default)</li></ul>
Compatibility	<ul style="list-style-type: none"><li>• can break legitimate programs<ul style="list-style-type: none"><li>- Just-In-Time compilers</li><li>- unpackers</li></ul></li></ul>
Safety Guarantee	<ul style="list-style-type: none"><li>• code injected to NX pages never execute</li><li>• <i>but code injection may not be necessary...</i></li></ul>

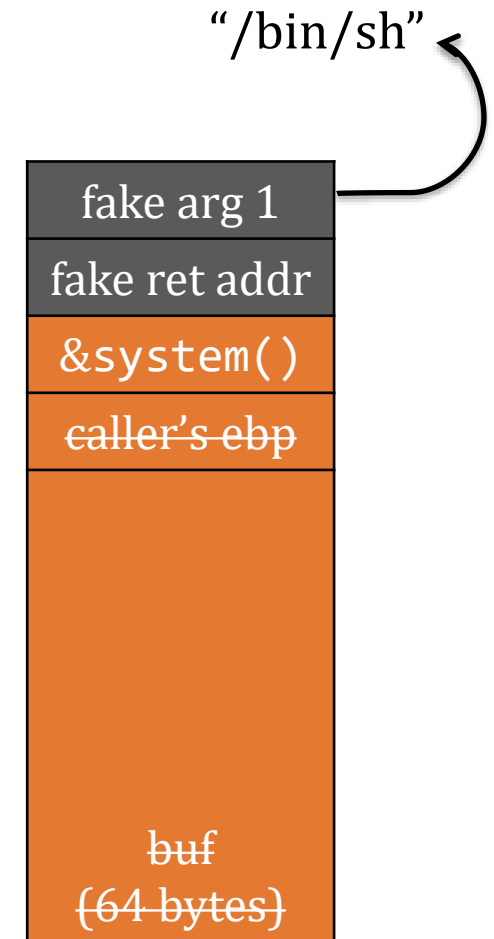


# Return-to-libc Attack

Overwrite return address by address of a libc function

- setup fake return address and argument(s)
- ret will “call” libc function

**No injected code!**



Reading:

The Geometry of Innocent Flesh on the Bone: Return-into-libc without Function Calls (on the x86), CCS 2007

More to come later

return-Oriented  
PROgramming

# Agenda

Canary / Stack Cookies



Data Execution Prevention (DEP)  
/No eXecute (NX)



Address Space Layout  
Randomization (ASLR)

# Agenda


Canary / Stack Cookies



Data Execution Prevention (DEP)  
/No eXecute (NX)



Address Space Layout  
Randomization (ASLR)




# Address Space Layout Randomization (ASLR)

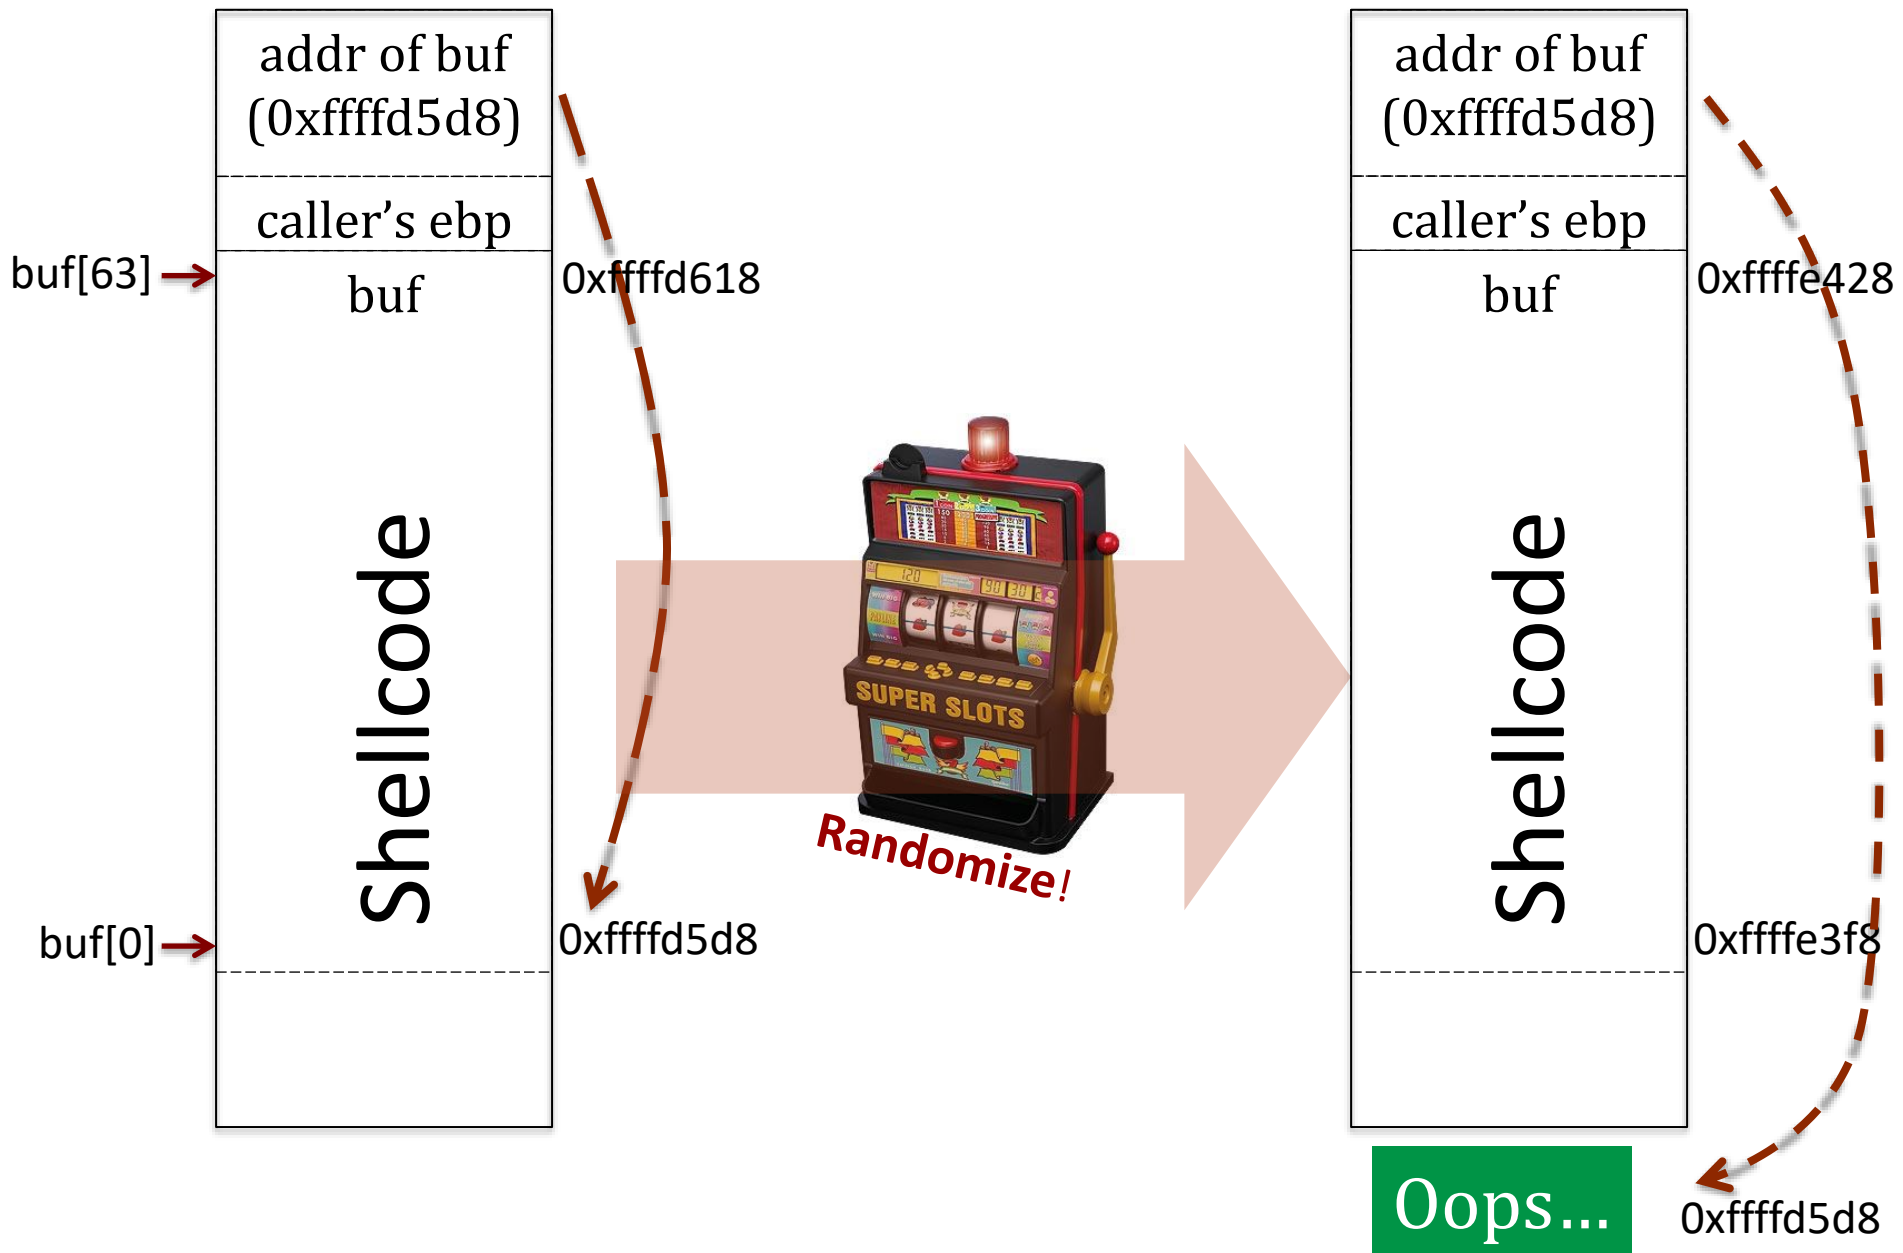
## **Assigned Reading:**

*ASLR Smack and Laugh Reference*

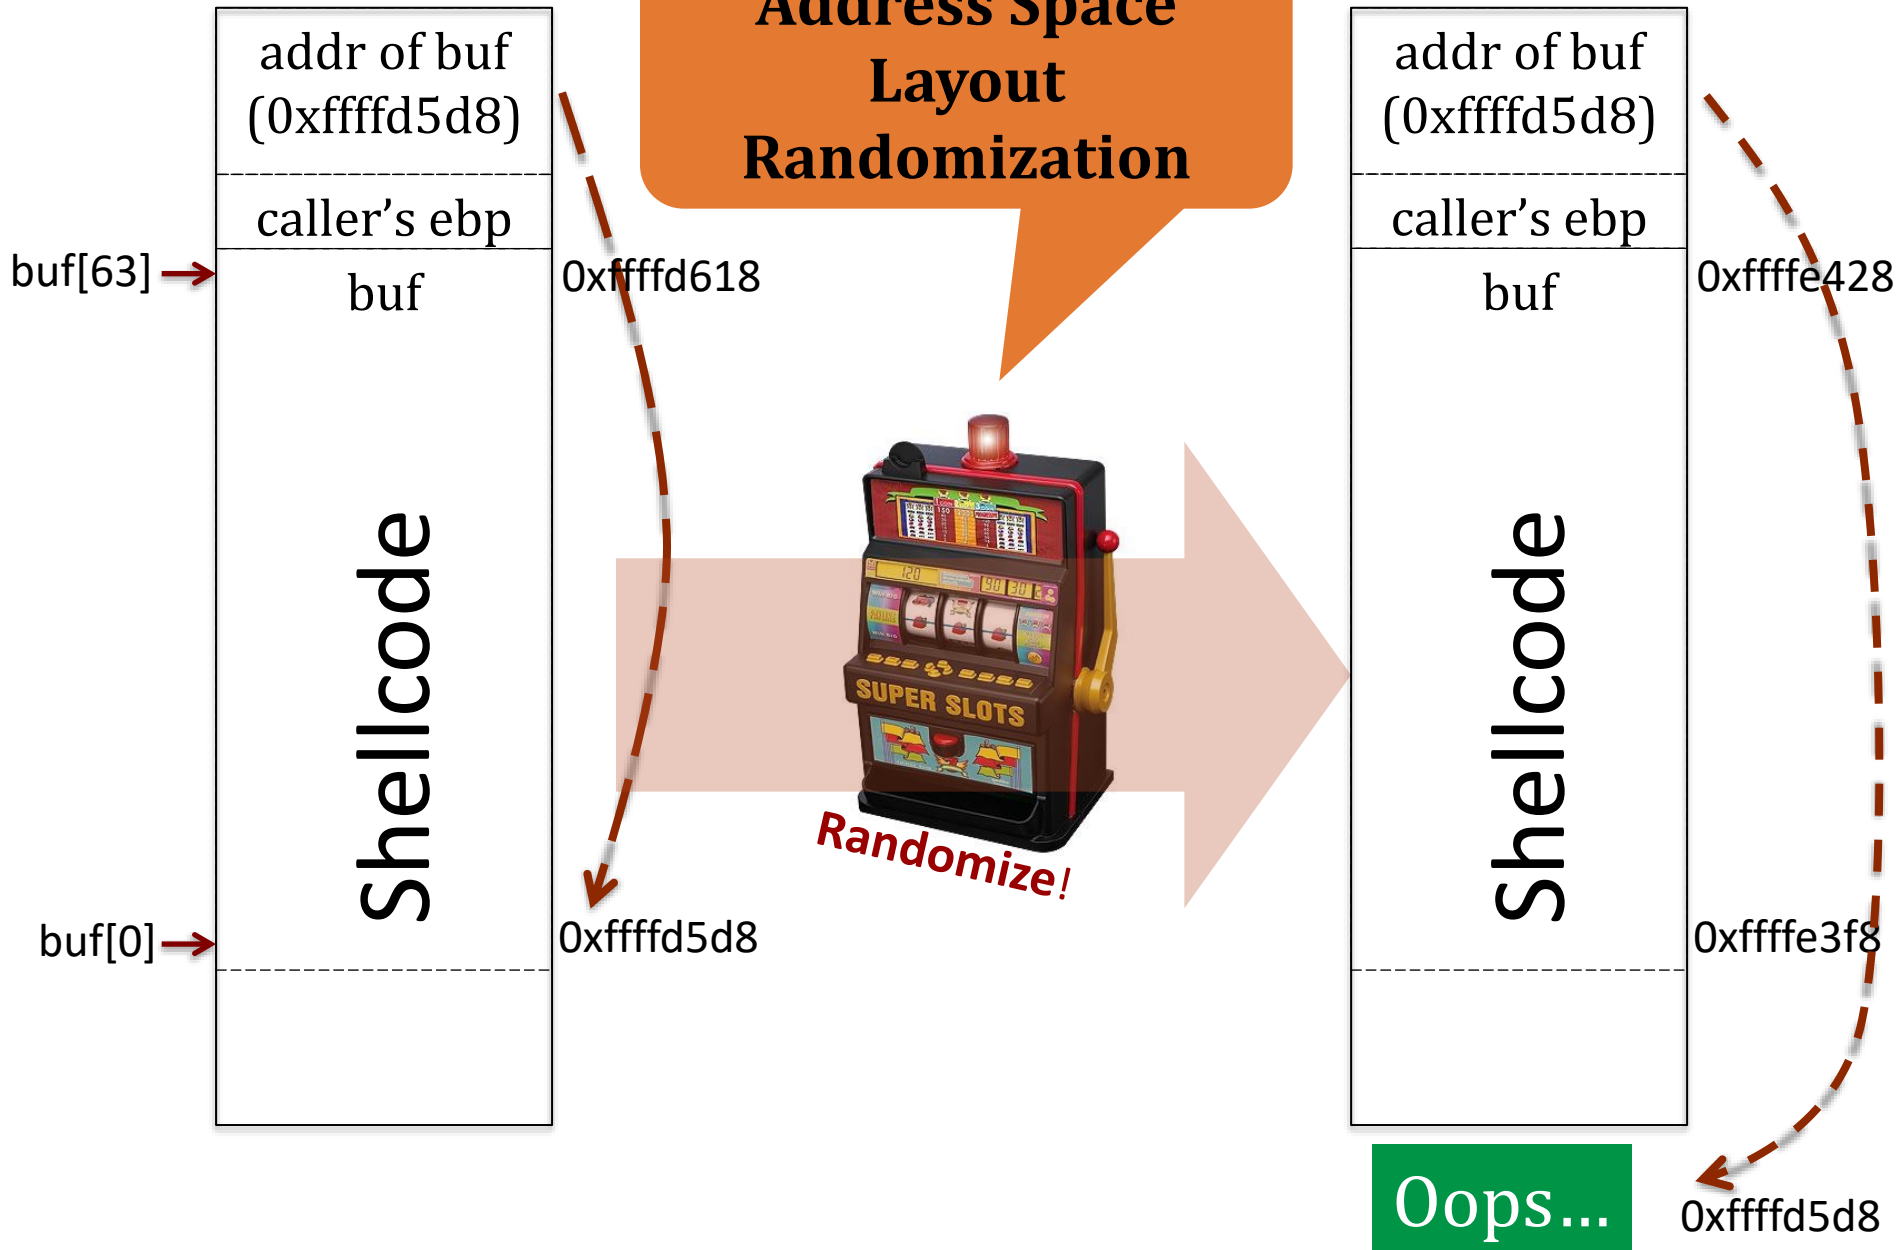
by Tilo Muller

[http://www.cs.ucr.edu/~zhiyunq/teaching/  
cs165/resources/paper/aslr\\_smack.pdf](http://www.cs.ucr.edu/~zhiyunq/teaching/cs165/resources/paper/aslr_smack.pdf)





## Address Space Layout Randomization





# ASLR

Traditional exploits need precise addresses

- *stack-based overflows*: location of shell code
- *return-to-libc*: library addresses

- **Problem:** program's memory layout is fixed
  - stack, heap, libraries etc.
- **Solution:** randomize addresses of each region!

# Running cat Twice

- Run 1

```
exploit:~# cat /proc/self/maps | egrep '(libc|heap|stack)'
```

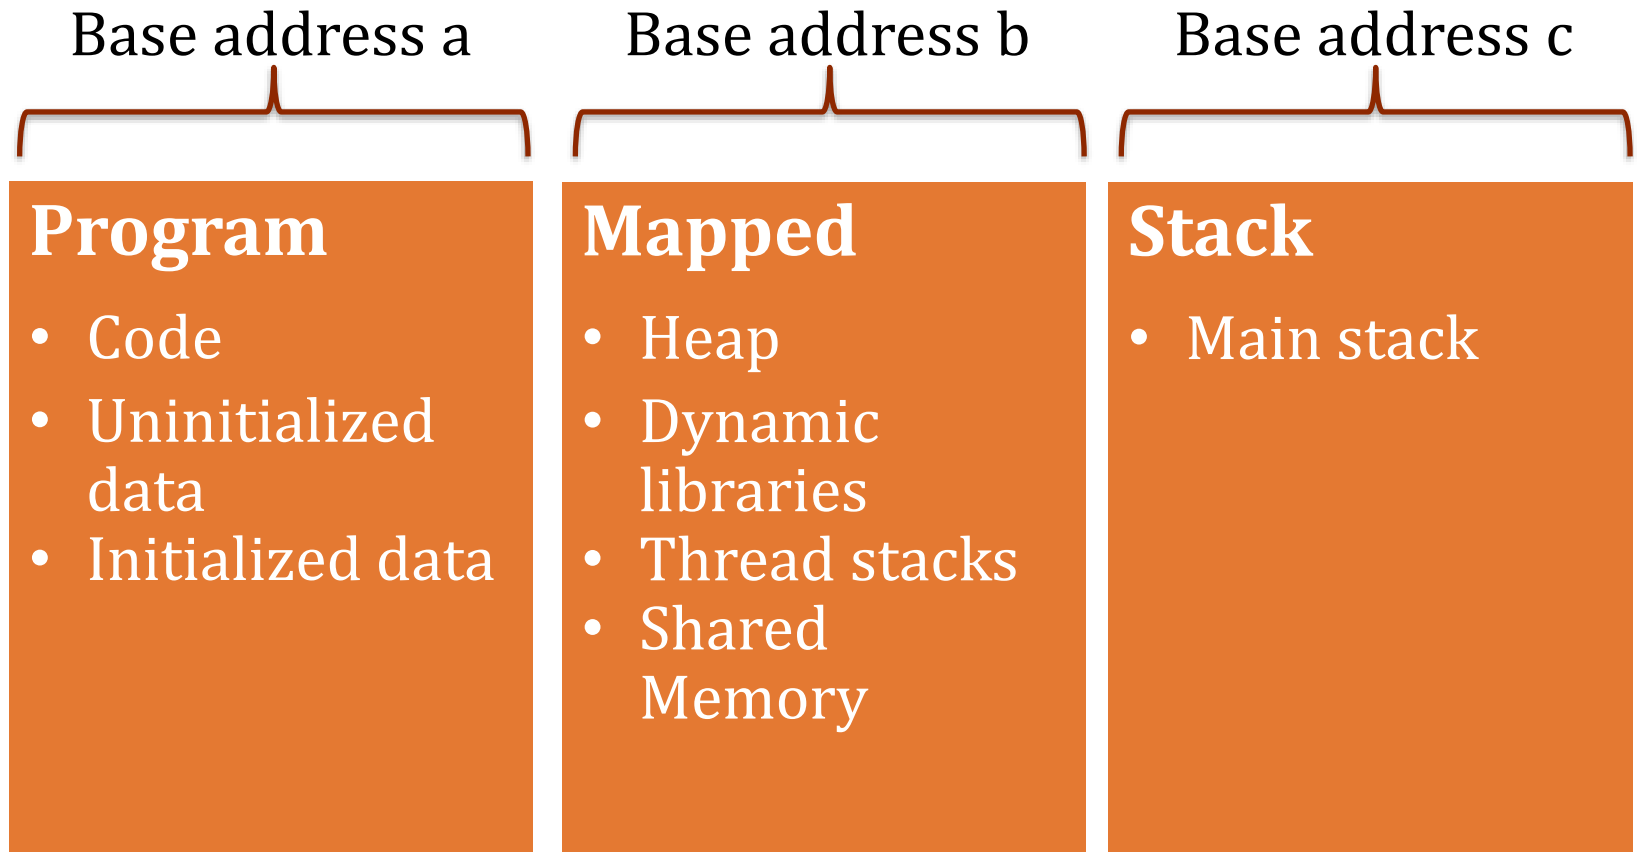
082ac000-082cd000	rw-p	082ac000	00:00	0	[heap]
b7dfe000-b7f53000	r-xp	00000000	08:01	1750463	/lib/i686/cmov/libc-2.7.so
b7f53000-b7f54000	r--p	00155000	08:01	1750463	/lib/i686/cmov/libc-2.7.so
b7f54000-b7f56000	rw-p	00156000	08:01	1750463	/lib/i686/cmov/libc-2.7.so
bf966000-bf97b000	rw-p	bf966000	00:00	0	[stack]

- Run 2

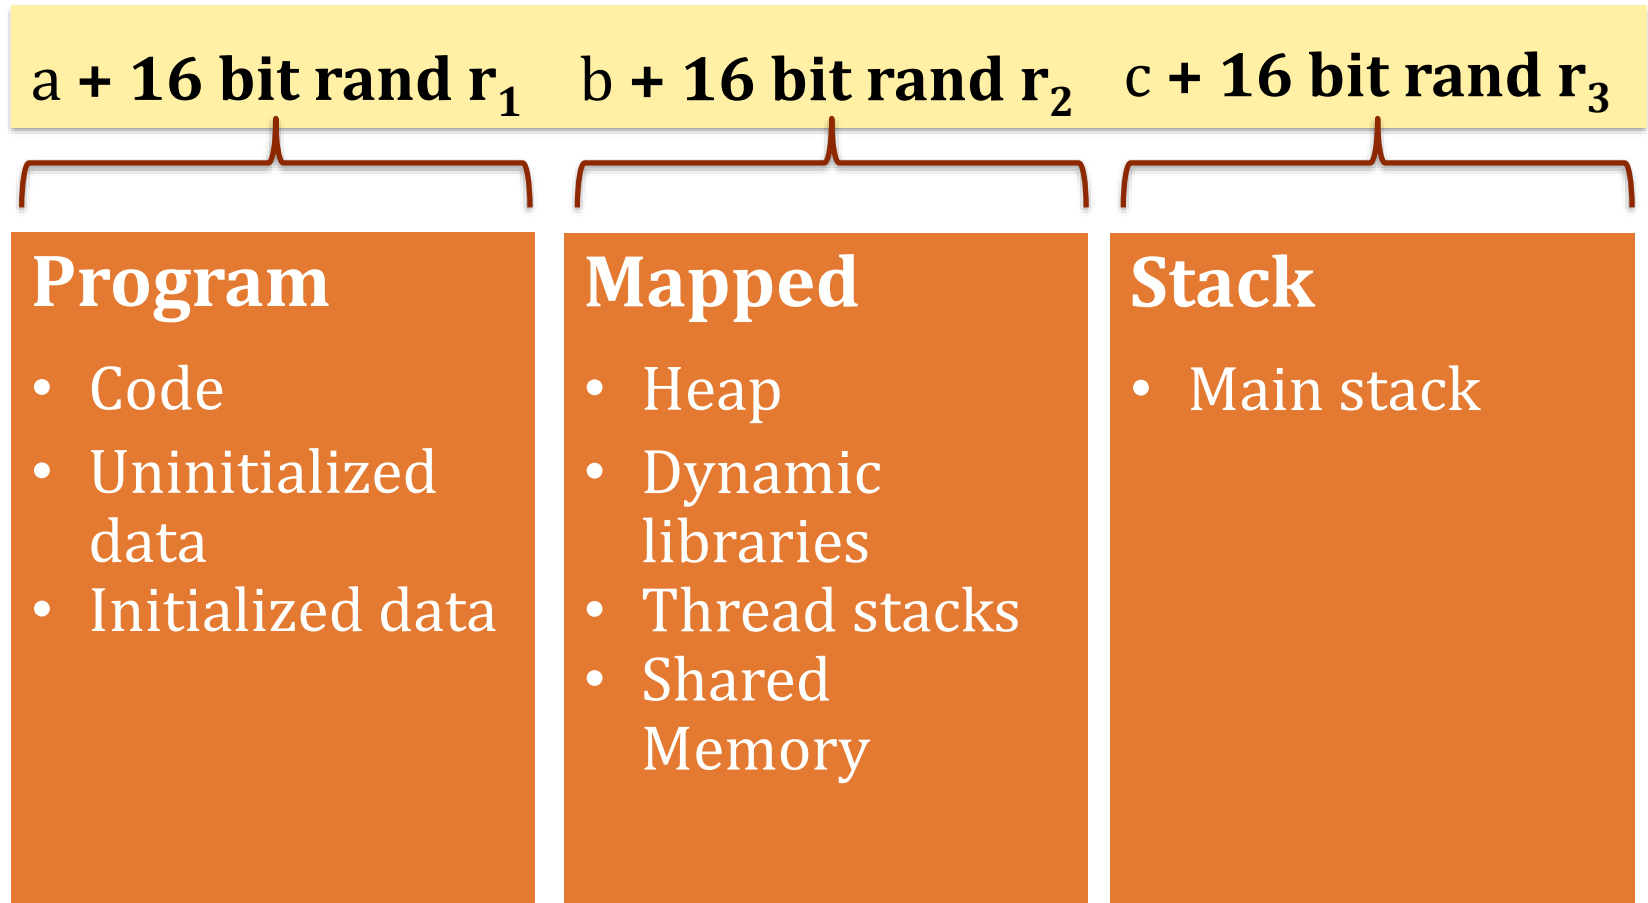
```
exploit:~# cat /proc/self/maps | egrep '(libc|heap|stack)'
```

086e8000-08709000	rw-p	086e8000	00:00	0	[heap]
b7d9a000-b7eef000	r-xp	00000000	08:01	1750463	/lib/i686/cmov/libc-2.7.so
b7eef000-b7ef0000	r--p	00155000	08:01	1750463	/lib/i686/cmov/libc-2.7.so
b7ef0000-b7ef2000	rw-p	00156000	08:01	1750463	/lib/i686/cmov/libc-2.7.so
bf902000-bf917000	rw-p	bf902000	00:00	0	[stack]

# Memory



# ASLR Randomization



\*  $\approx$  16 bit random number of 32-bit system. More (up to 32) on 64-bit systems.

# ASLR Scorecard

Aspect	Address Space Layout Randomization
Performance	<ul style="list-style-type: none"><li>• excellent—randomize once at load time</li></ul>
Deployment	<ul style="list-style-type: none"><li>• turn on kernel support (Windows: opt-in per module, but system override exists)</li><li>• no recompilation necessary</li></ul>
Compatibility	<ul style="list-style-type: none"><li>• transparent to safe apps (position independent)</li></ul>
Safety Guarantee	<ul style="list-style-type: none"><li>• not good on x32, much better on x64</li><li>• <i>possible to leak?</i></li></ul>

# Ubuntu - ASLR

- ASLR is **ON** by default [Ubuntu-Security]
  - cat /proc/sys/kernel/randomize\_va\_space
    - Prior to Ubuntu 8.10: **1** (*stack/mmap* ASLR)
    - In later releases: **2** (*stack/mmap/brk* ASLR)
  - stack/mmap ASLR: since kernel 2.6.15 (Ubuntu 6.06)
  - brk ASLR: since kernel 2.6.26 (Ubuntu 8.10)
  - exec ASLR: since kernel 2.6.25
    - Position Independent Executable (PIE) with “-fPIE -pie”

# How to attack with ASLR?

## Attack

Brute  
Force

Non-  
randomized  
memory

Stack  
Juggling

GOT  
Hijacking

ret2text

Func ptr

ret2eax


ret2got

# Brute Force


Shell code 




# Brute Force



# Brute Force



# Brute Force



# How to attack with ASLR?

## Attack

**Brute  
Force**

**Non-  
randomized  
memory**

**Stack  
Juggling**

**GOT  
Hijacking**

**ret2text**

**ret2eax**

**ret2got**

**Func ptr**


All assume some memory non-randomized

# ret2text

- `text` section has executable program code
  - but not typically randomized by ASLR except PIE
- can hijack control flow to unintended (but existing) program function

# ret2text

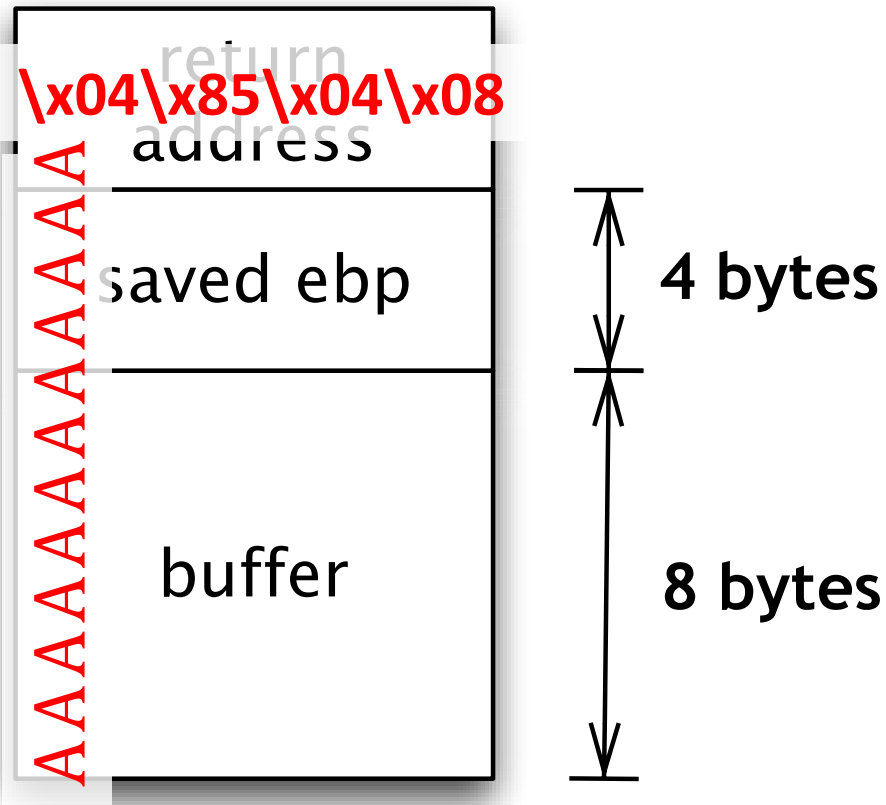
.text not  
randomized



```
08048504 <secret>:
8048504: 55
8048505: 89 e5
8048507: 83 ec 18
804850a: 8b 45 08
804850d: 89 44 24 04
8048511: c7 04 24 f0 86 04 08
8048518: e8 df fe ff ff
804851d: c7 44 24 0c 00 00 00
8048524: 00
8048525: c7 44 24 08 22 87 04
804852c: 08
804852d: c7 44 24 04 28 87 04
8048534: 08
8048535: c7 04 24 2c 87 04 08
804853c: e8 9b fe ff ff
8048541: b8 01 00 00 00
8048546: c9
8048547: c3
```

# ret2text

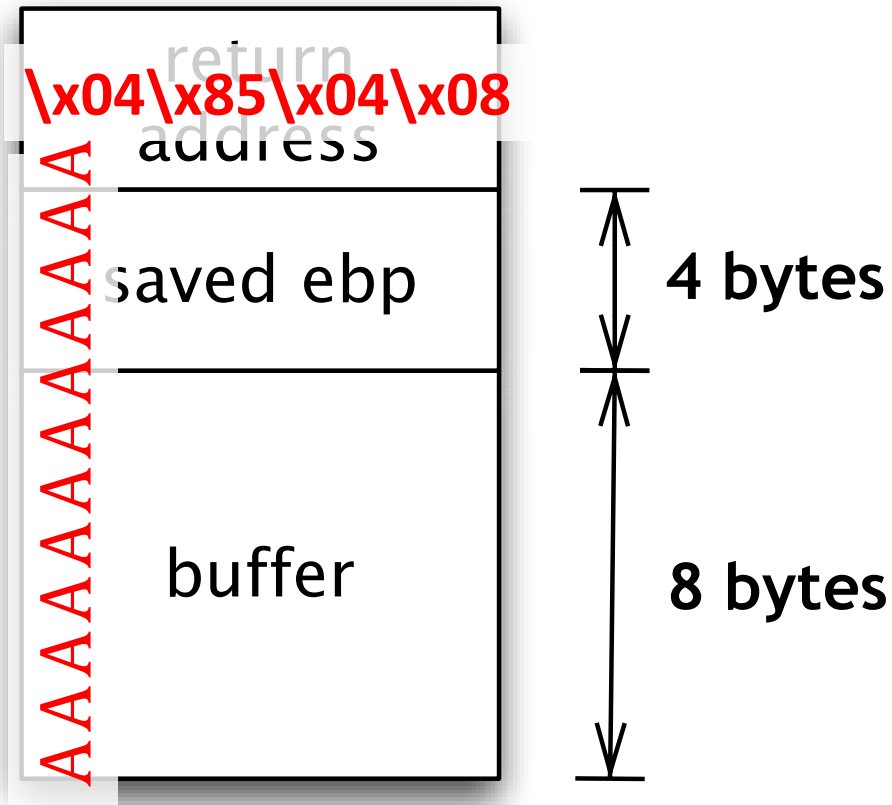
.text not  
randomized



08048504	<secret>
08048504:	55
8048505:	89 e5
8048507:	83 ec 18
804850a:	8b 45 08
804850d:	89 44 24 04
8048511:	c7 04 24 f0 86 04 08
8048518:	e8 df fe ff ff
804851d:	c7 44 24 0c 00 00 00
8048524:	00
8048525:	c7 44 24 08 22 87 04
804852c:	08
804852d:	c7 44 24 04 28 87 04
8048534:	08
8048535:	c7 04 24 2c 87 04 08
804853c:	e8 9b fe ff ff
8048541:	b8 01 00 00 00
8048546:	c9
8048547:	c3

# ret2text

.text not  
randomized



08048504	<secret>
08048504:	55
8048505:	89 e5
8048507:	83 ec 18
804850a:	8b 45 08
804850d:	89 44 24 04
8048511:	c7 04 24 f0 86 04 08
8048518:	e8 df fe ff ff
804851d:	c7 44 24 0c 00 00 00
8048524:	00
8048525:	c7 44 24 08 22 87 04
804852c:	08
804852d:	c7 44 24 04 28 87 04
8048534:	08
8048535:	c7 04 24 2c 87 04 08
804853c:	e8 9b fe ff ff
8048541:	b8 01 00 00 00
8048546:	c9
8048547:	c3

Same as running a “secret”  
function in project 3



# How to attack with ASLR?

## Attack

**Brute  
Force**

**Non-  
randomized  
memory**

**Stack  
Juggling**

**GOT  
Hijacking**

**ret2text**

**ret2eax**

**ret2got**

**Func ptr**

All assume some memory non-randomized

# ret2eax

```
void msglog(char *input) {  
    char buf[64];  
    strcpy(buf, input);  
}
```

returns pointer to  
buf in eax  
eax = buf

```
int main(int argc, char *argv[]) {  
    if(argc != 2) {  
        printf("exploitme <msg>\n");  
        return -1;  
    }
```

```
    msglog(argv[1]);
```

```
    return 0;  
}
```

# ret2eax

```
void msglog(char *input) {  
    char buf[64];  
    strcpy(buf, input);  
}
```


returns pointer to  
buf in eax  
eax = buf

```
int main(int argc, char *argv[]) {  
    if(argc != 2) {  
        printf("exploitme <msg>\n");  
        return -1;  
    }
```

```
    msglog(argv[1]);  
  
    return 0;  
}
```

A subsequent  
call \*eax  
would redirect  
control to buf

# ret2eax

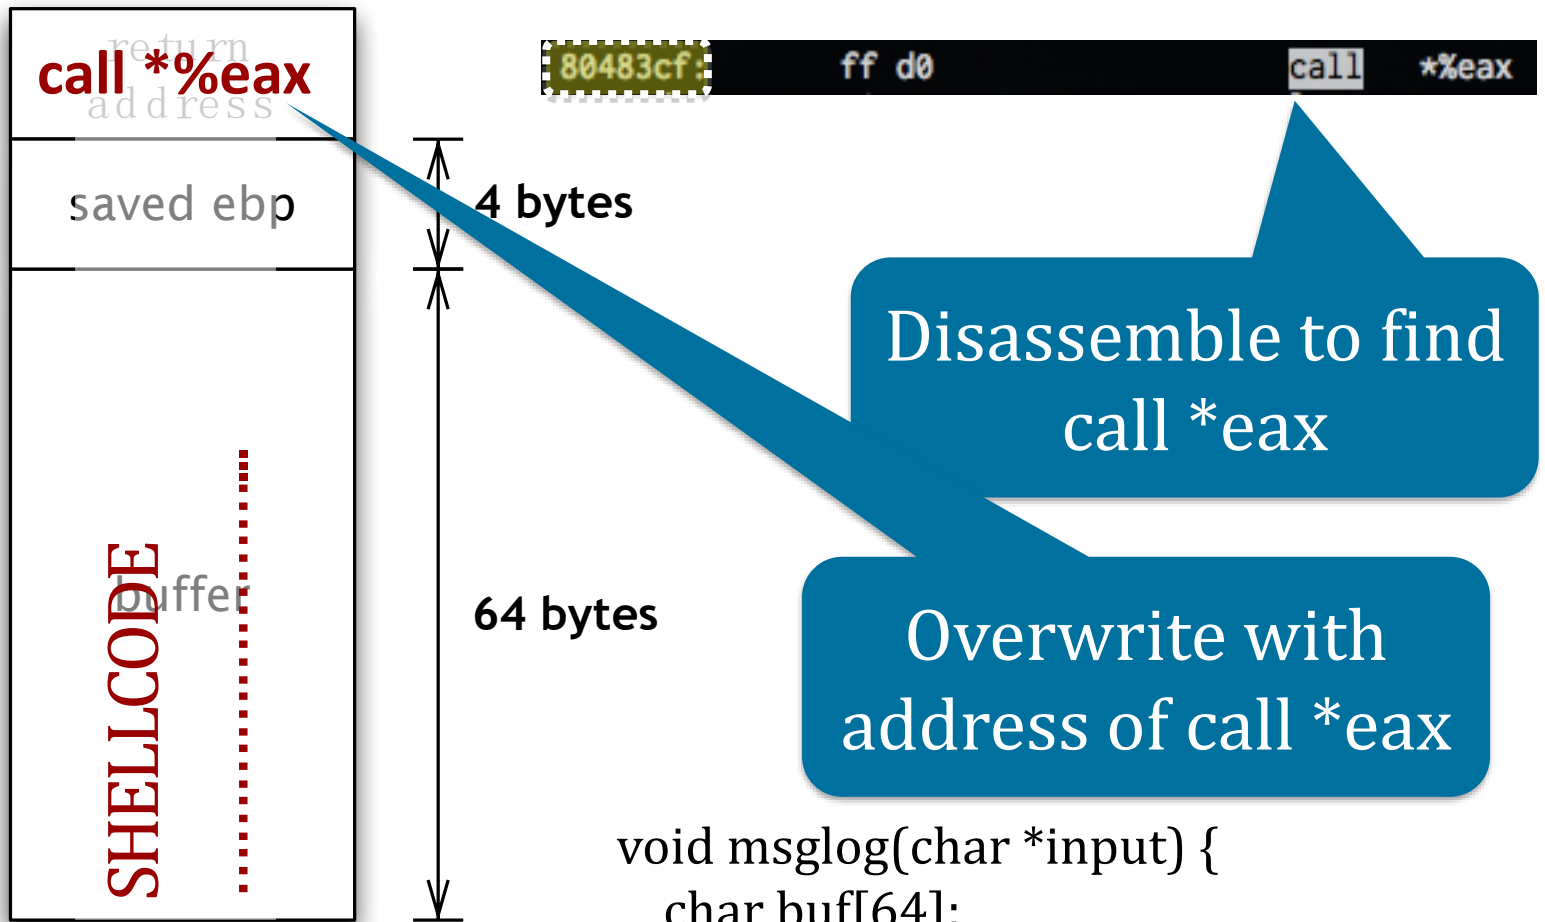


80483cf: ff d0 call \*%eax

Disassemble to find  
call \*eax

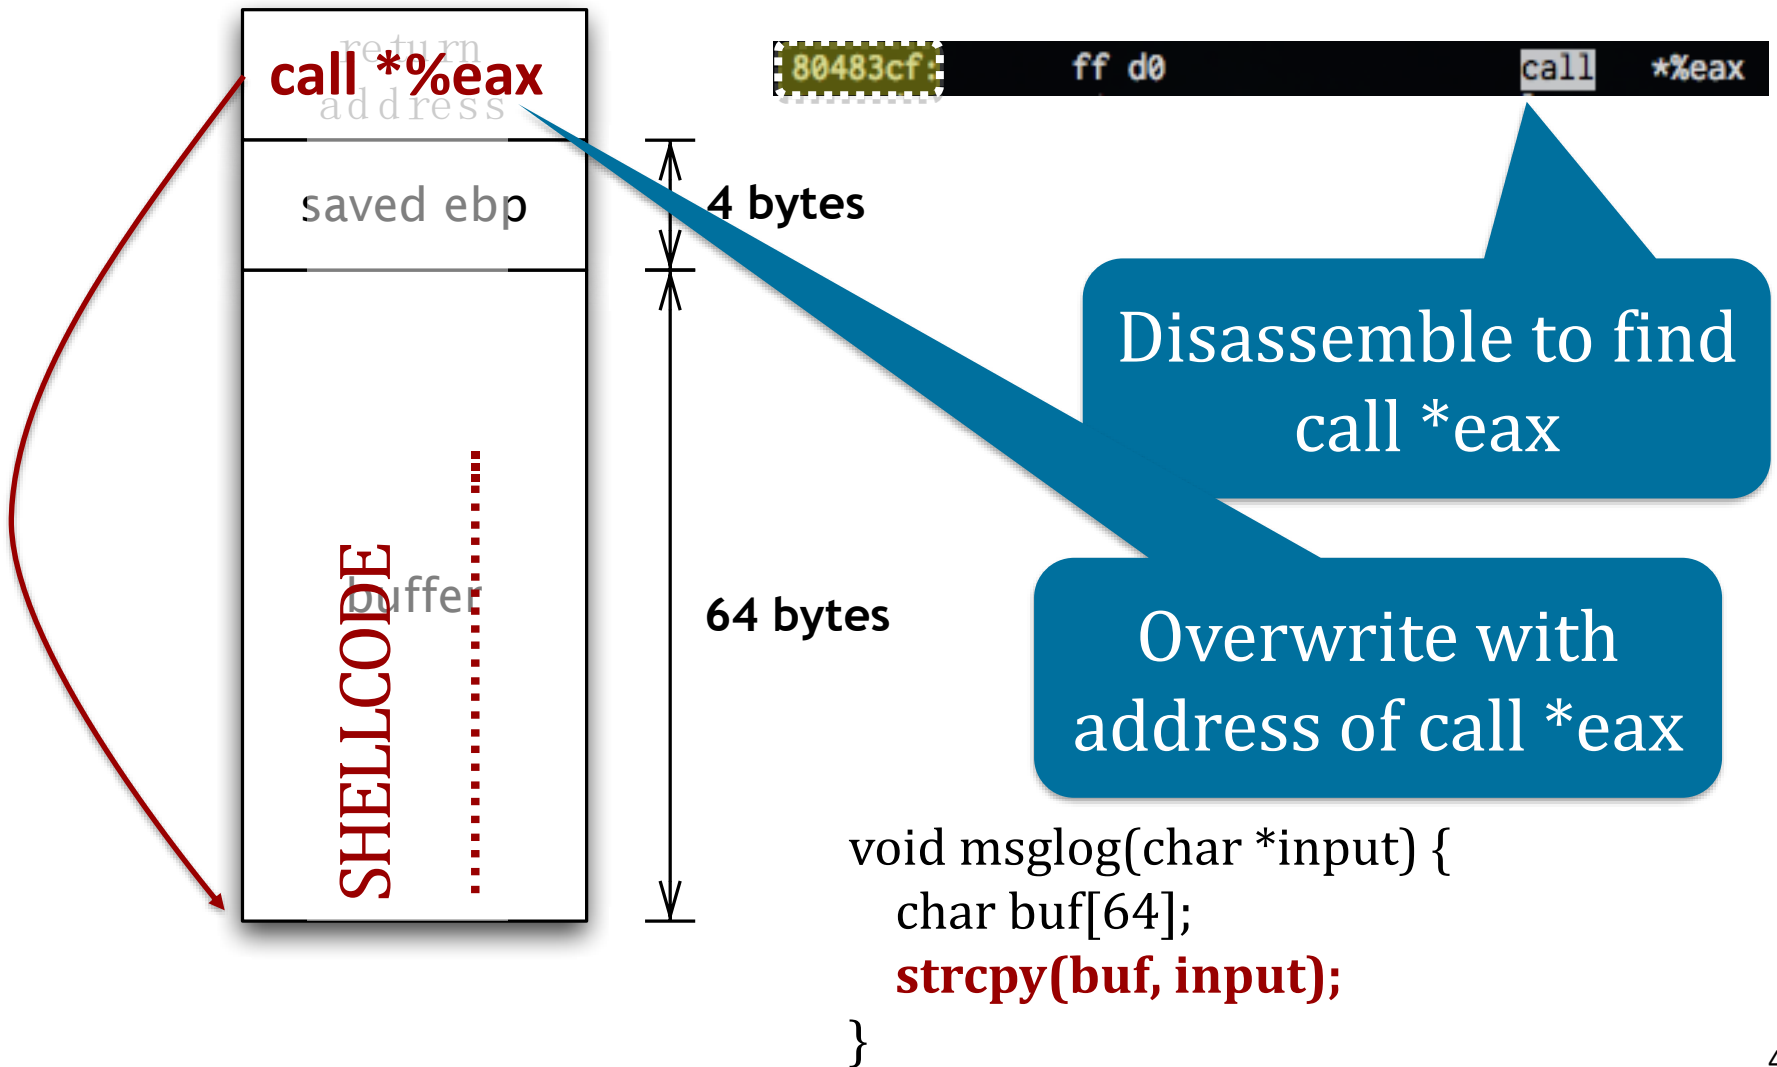
```
void msglog(char *input) {  
    char buf[64];  
    strcpy(buf, input);  
}
```

# ret2eax



```
void msglog(char *input) {  
    char buf[64];  
    strcpy(buf, input);  
}
```

# ret2eax



# How to attack with ASLR?

## Attack

**Brute  
Force**

**Non-  
randomized  
memory**

**Stack  
Juggling**

**GOT  
Hijacking**

**ret2text**

**ret2eax**

**ret2got**

**Func ptr**

All assume some memory non-randomized

# Questions

