

CS165 – Computer Security

Attack surface and access control
Oct 26, 2021

Security Problems

- ... could be anywhere in a program
 - Given the types of **vulnerability** we have seen, does that give us any insight into where we should look for vulnerabilities?

Security Problems

- ... could be anywhere in a program
 - Given the types of **vulnerability** we have seen, does that give us any insight into where we should look for vulnerabilities?
 - Software flaw
 - Accessible to an adversary
 - Who can exploit the vulnerability

Security Problems

- ... could be anywhere in a program
 - Given the types of **vulnerability** we have seen, does that give us any insight into where we should look for vulnerabilities?
 - Software flaw
 - Accessible to an adversary
 - Who can exploit the vulnerability
- Typically, we look for **software flaws** (e.g., control flow hijacking), but let us first consider “**adversary accessibility**”

Attack Surface


- After Microsoft faced several large-scale vulnerability exploits in the early 2000s
- They began to consider how to prevent such vulnerabilities

Attack Surface

- After Microsoft faced several large-scale vulnerability exploits in the early 2000s
- They began to consider how to prevent such vulnerabilities
- Michael Howard of Microsoft defined the term “attack surface”
 - A program’s **attack surface** consists of the **entry points** that are **accessible to an adversary**
- Entry point: where **untrusted inputs** come in
- Example attack surfaces?

Example attack surface?


- Relative to a threat model



Goal: execute code remotely



Goal: change the behavior of HTTP server



Goal: change the behavior of HTTP server

Anything a target depends on (directly or indirectly) should be included

Attack surface of autonomous vehicle



Relative Attack Surface Metric (system)

- Howard proposed the notion of a relative attack surface quotient (RASQ) metric
 - The idea is that we can use the metric to compare **systems** -- which has larger relative attack surface
- The metric lists a set of entry points that you should be concerned about minimizing as a system distributor

Relative Attack Surface Metric (system)

- Open (TCP/UDP) sockets - descriptors
- Open RPC endpoints - descriptors
- Open named pipes - descriptors
- Services - daemons
- Services running by default - daemons
- Services running as SYSTEM (or root) - daemons
- Active Web handlers – web server components
- Active ISAPI filters – web server add-ins
- Dynamic web pages – files
- Executable vdirs – directories for scripts

Relative Attack Surface Metric (system)

- Enabled accounts – accounts
- Enabled accounts in admin group – accounts
- Null sessions to pipes and shares – anonymous connections allowed
- Guest account enabled – accounts (special)
- Weak ACLs in FS – files allowing “full control” to everybody
 - “Full control” is the moral equivalent of UNIX rwxrwxrwx permissions
- Weak ACLs in Registry – registry keys that allow “full control” to everybody
- Weak ACLS on shares – Directories that can be shared by remote users that allow “full control” to everybody
- VBScript, JavaScript, Active X enabled – applications enabled to execute Visual Basic Script, JavaScript or Active X controls

Relative Attack Surface Metric (system)

- Essentially, you would count the number of unsafe instances
 - Also combined with weights per item, but numeric weights that are meaningful are often hard to predict effectively
- Windows systems saw a gradual reduction in attack surface metric values in the 2000s
 - But, attacks kept coming, exploiting new vulnerabilities
- Can we say something about programs individually with respect to their attack surfaces?

Program Attack Surface

- Can we say something about programs **individually** with respect to their attack surfaces?
- What do we need to identify to determine the adversary-accessible entry points of a program?

Program Attack Surface

- Can we say something about programs **individually** with respect to their attack surfaces?
- What do we need to identify to determine the adversary-accessible entry points of a program?
 - Identify the relevant subset of **system resources** that can be used in an attack (are or could be controlled by an adversary)
 - Identify when such **resources may be used by the program (program entry points)**
- Is it possible to compute such information?

Program Entry Points

- What's a program entry point?
- Programs obtain information from external sources (e.g., files and network sockets), and the **program statements that access such external sources** are entry points
 - What's an example of an entry point?

Program Entry Points

- What's a program entry point?
- Programs obtain information from external sources (e.g., files and network sockets), and the program statements that access such external sources are entry points
 - What's an example of an entry point?
 - System calls provide the sources for gaining most external information
 - But, for attack surfaces, we focus on the statements that a program makes to the individual library/system calls

System Calls as Attack Surface

- Why should we use system calls for the attack surface?

System Calls as Attack Surface

- Why should we use system calls for the attack surface?
 - A Program has to use library calls to access external resources
 - Wrappers in libraries, e.g., `fopen()` vs. `open()`
 - Program statements that invoke each call
 - Only a subset of these may be adversary accessible
- E.g., consider the “open” system call
 - May be invoked via “open” or “fopen” library call
 - `fopen(input_pathname, ...)` vs. `fopen(“/bin/sh”)`

System Calls as Attack Surface

- Why should we use system calls for the attack surface?
 - A Program has to use library calls to access external resources
 - Wrappers in libraries, e.g., `fopen()` vs. `open()`
 - Program statements that invoke each call
 - Only a subset of these may be adversary accessible
- E.g., consider the “open” system call
 - May be invoked via “open” or “fopen” library call
 - `fopen(input_pathname, ...)` vs. `fopen(“/bin/sh”)`
- E.g., consider the “read” system call
- How many system calls access adversary-controlled data?

System Calls as Attack Surface

- Which system calls should constitute a program's attack surface?

System Calls as Attack Surface

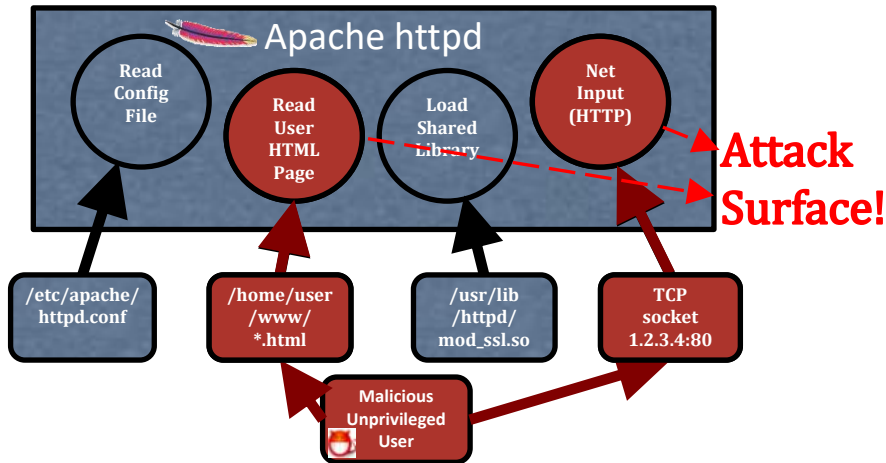
- Which system calls should constitute a program's attack surface?
 - All of them
 - At some point, any call may access adversary-controlled data
 - So test them all

System Calls as Attack Surface

- Which system calls should constitute a program's attack surface?
 - All of them
 - At some point, any call may access adversary-controlled data
 - So test them all
 - Only ones that actually may access **adversary-controlled** resources
 - Only need to test a subset of such each program's entry points to evaluate the attack surface
 - How do you determine which may access adversary-controlled resources?

Program Attack Surface

- Program system calls accessible to an adversary



System TCB Attack Surface

- Only **13.8%** of total entry points for Linux system services were accessible to adversaries at all
 - Only **3.8%** for read/write operations
 - Listing all entry points as attack surface would be a huge over-approximation

| Total Entrypoints | Accessible to Adversaries | Potentially Vulnerable (overt permissions) | Previously Known Bugs |
|-------------------|---------------------------|--|-----------------------|
| 2138 | 295 | 81 | 35 |

- Found via runtime testing with Linux package test suites – lower bound

An example - E-voting

- Who are the principals?
 - Voters, Admins, Talliers, Others
- Who are adversaries?
- Which commands may be threatened (attack surface)?
 - Start Program (by admins)
 - Submit a vote (by voter)
 - Count votes (by tallier)


Access control

- Give “users” permissions to access “resources”

```
-rwxrwxr-x 1 zhiyung zhiyung 7.2K Jan 9 18:04 test
```

- Attack surface computed with respect to a **threat model**
 - Local attacker (unprivileged)
 - Local attacker (system privilege)
 - Local attack (with root)
 - Remote attacker
- ... **and an access control policy**
 - What can a (unprivileged/system/root) user do on a system? What files can the user write? What services can they contact?
- Why do we need access control?
 - Hint: think of it from the attack surface point of view
 - (Android story)
- Challenge: how do we know if our policy is good enough? The right policy? The right implementation?





Policies and Mechanisms

- Policy says what is, and is not, allowed
 - This defines “security” for the site/system/*etc.*
- Mechanisms enforce policies
- Composition of policies
 - If policies conflict, discrepancies may create security vulnerabilities

Policies and Mechanisms (example)

- Policy says what is, and is not, allowed
 - A file should be readable by only the root user
- Mechanisms enforce policies
 - Kernel checks through file `open()` syscall.


Agenda

- Access Control Matrix
 - Overview
 - Access Control Matrix Model
 - Protection State Transitions
 - Commands
 - Conditional Commands
- Foundational Results



Alice and Bob

- Standard names for “agents” in a security or crypto scenario
- Also known as “A” and “B”



An Access Control Scenario



- Alice:

1. New Secret foo


Intent:

- Bob's **cp** is attempting to violate Alice's expected access policy
- If **cp** succeeds then the principle of **confidentiality** is not satisfied

- Bob:


2. If (cp foo afoo)
3. then echo "success"
4. else echo "fail"

Characterizing the Violation



Basic Abstraction: States and Transitions


Characterizing the Violation



Basic Abstraction: States and Transitions

Q: What are the States?

Characterizing the Violation



Basic Abstraction: States and Transitions

Q: What are the States?

Q: What determines if we reach State 2 or 4 from State 1?

States

- State of a system
 - A collection of the current values of all memory locations, storages, registers, etc.
 - A subset of this collection that deals with protection is the protection state of the system

Secure and non-Secure States

Characterize states in a system as “Secure” and “non-Secure”

A system is **Secure** if every transition maps Secure states to Secure states

Consequence: In the scenario, security is compromised if Alice’s “New secret foo” yields a state in which Bob can access foo.

Protection States

- An **abstraction** that focuses on security properties
 - ✓ Primarily interested in characterizing Safe states
 - ✓ Goal is to prove that all operations in the system preserve “security” of the protection state
 - ✓ **Access Control Matrix** is our first Protection State model

Questions

