

CS165 – Computer Security

Access Control, BLP and Biba Policies
Nov 2, 2021

Agenda

- Access Control Matrix
 - Overview
 - Access Control Matrix Model
 - Protection State Transitions
 - Commands
 - Conditional Commands
- Foundational Results



Protection States

- An **abstraction** that focuses on security properties
 - ✓ Primarily interested in characterizing Safe states
 - ✓ Goal is to prove that all operations in the system preserve “security” of the protection state
 - ✓ **Access Control Matrix** is our first Protection State model
- Protection state of system
 - Describes current settings, values of system relevant to protection

Confidentiality Scenario

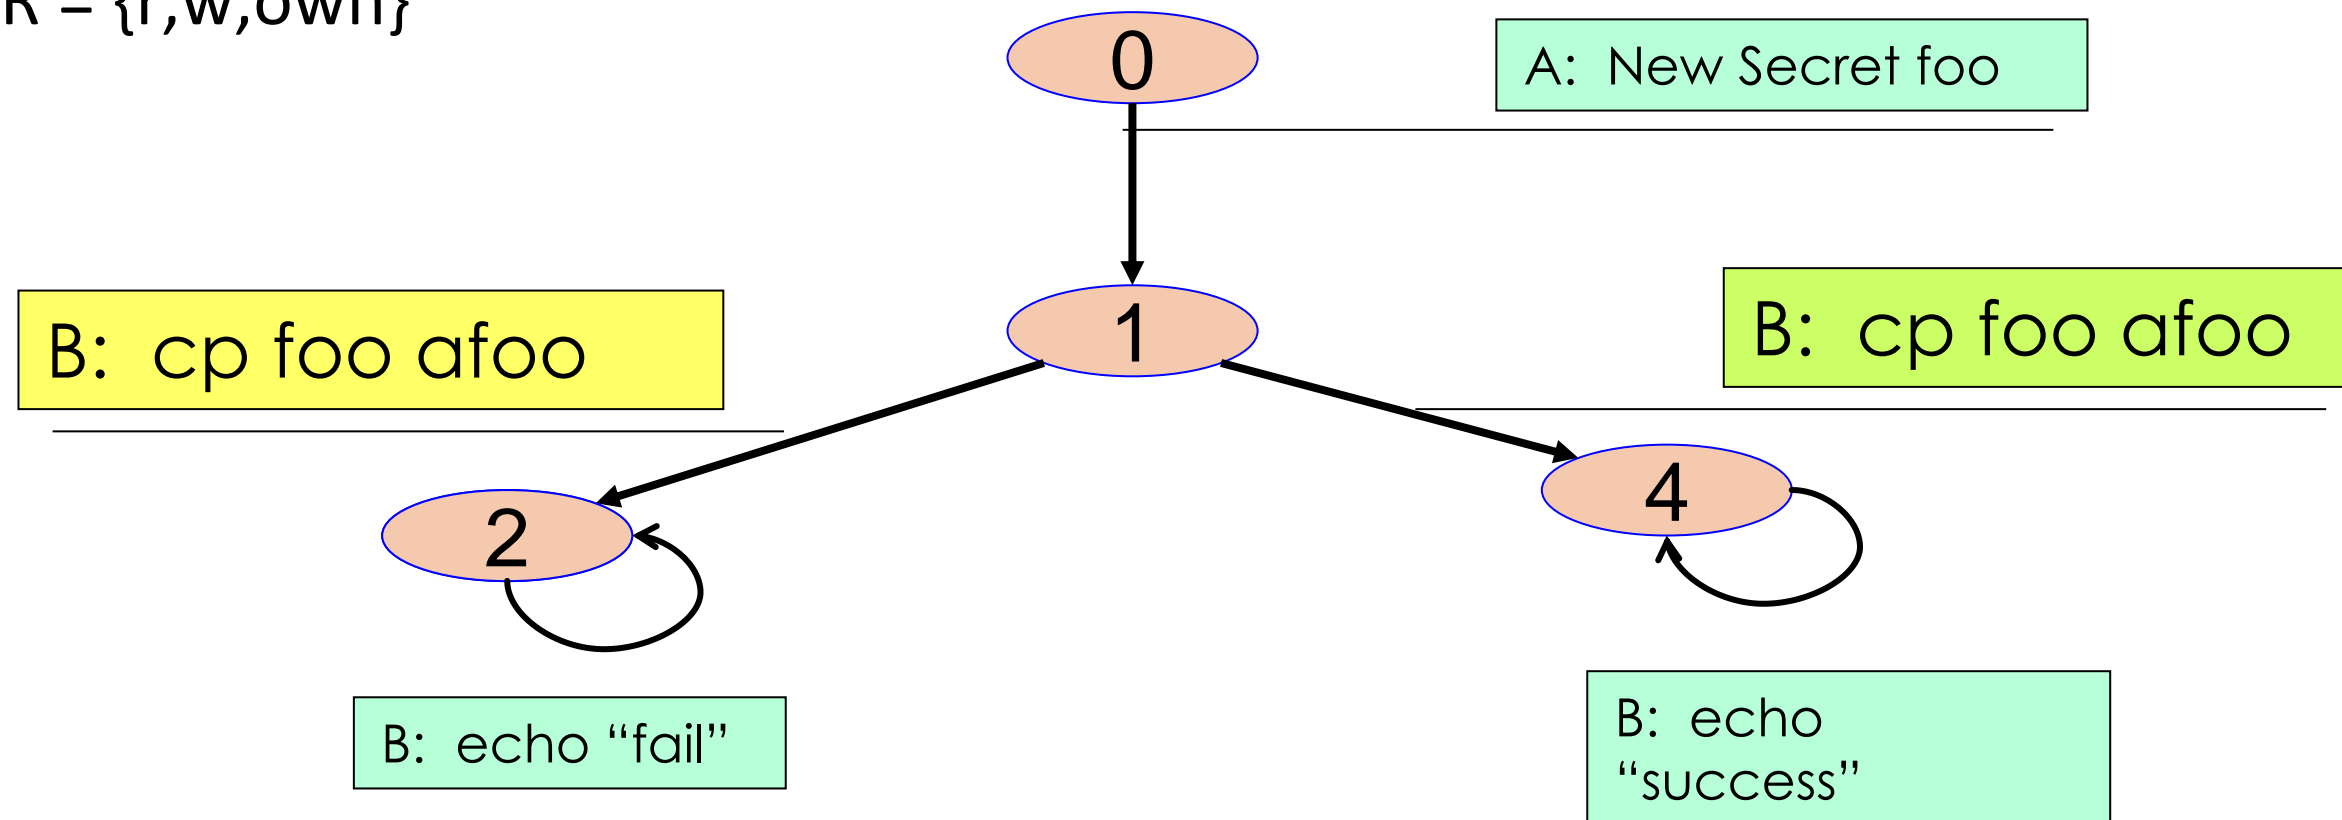
Initial State

Subjects $S_0 = \{A, B\}$

Objects $O_0 = \{\}$

AC Matrix $A_0 = \{\}$

Rights $R = \{r, w, own\}$



Confidentiality Scenario

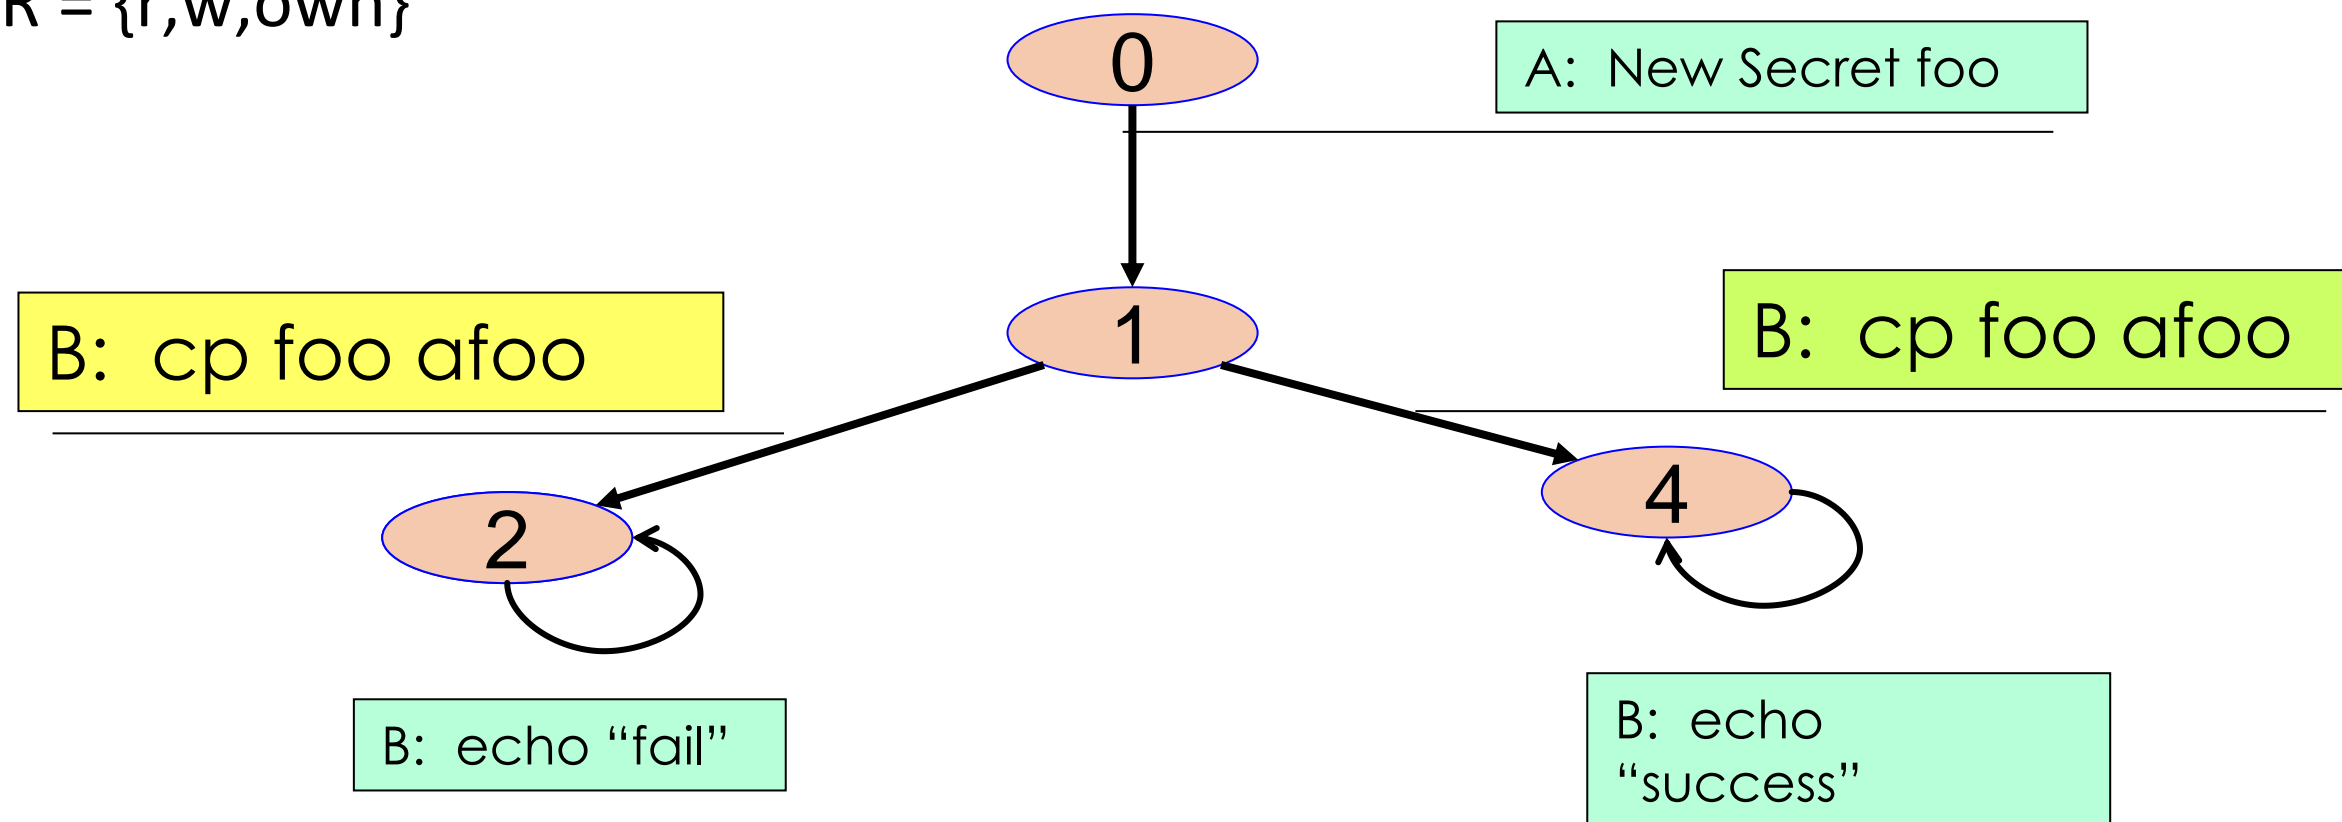
Initial State

Subjects $S_0 = \{A, B\}$

Objects $O_0 = \{\}$

AC Matrix $A_0 = \{\}$

Rights $R = \{r, w, own\}$



Intended State 1

Subjects $S_1 = \{A, B\}$

Objects $O_1 = \{\text{foo}\}$

AC Matrix $A_1 = \{ (A, \text{foo}, [r, w, own]), (B, \text{foo}, []) \}$

Confidentiality Scenario

Initial State

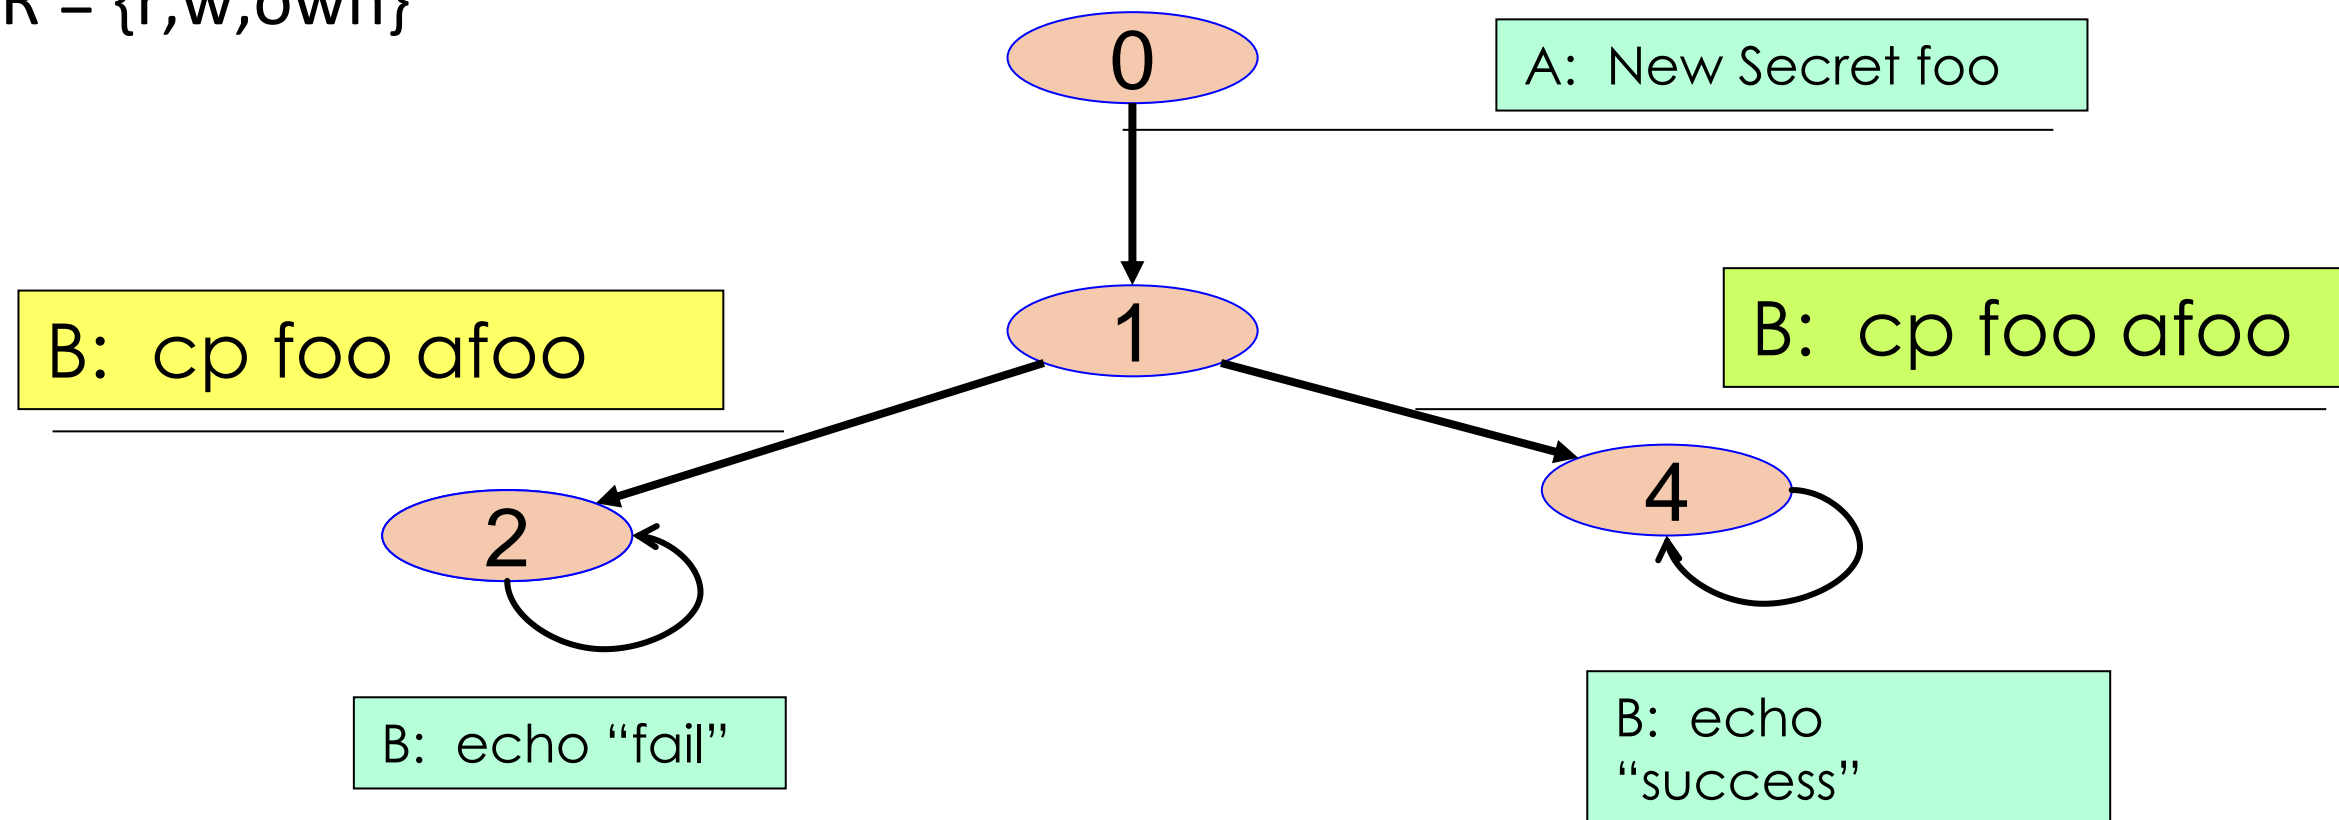
Subjects $S_0 = \{A, B\}$

Objects $O_0 = \{\}$

AC Matrix $A_0 = \{\}$

Rights $R = \{r, w, own\}$

(S_0, O_0, A_0) A: New Secret foo (S_1, O_1, A_1)



Intended State 1

Subjects $S_1 = \{A, B\}$

Objects $O_1 = \{\text{foo}\}$

AC Matrix $A_1 = \{ (A, \text{foo}, [r, w, own]), (B, \text{foo}, []) \}$

Confidentiality Scenario

Initial State

Subjects $S_0 = \{A, B\}$

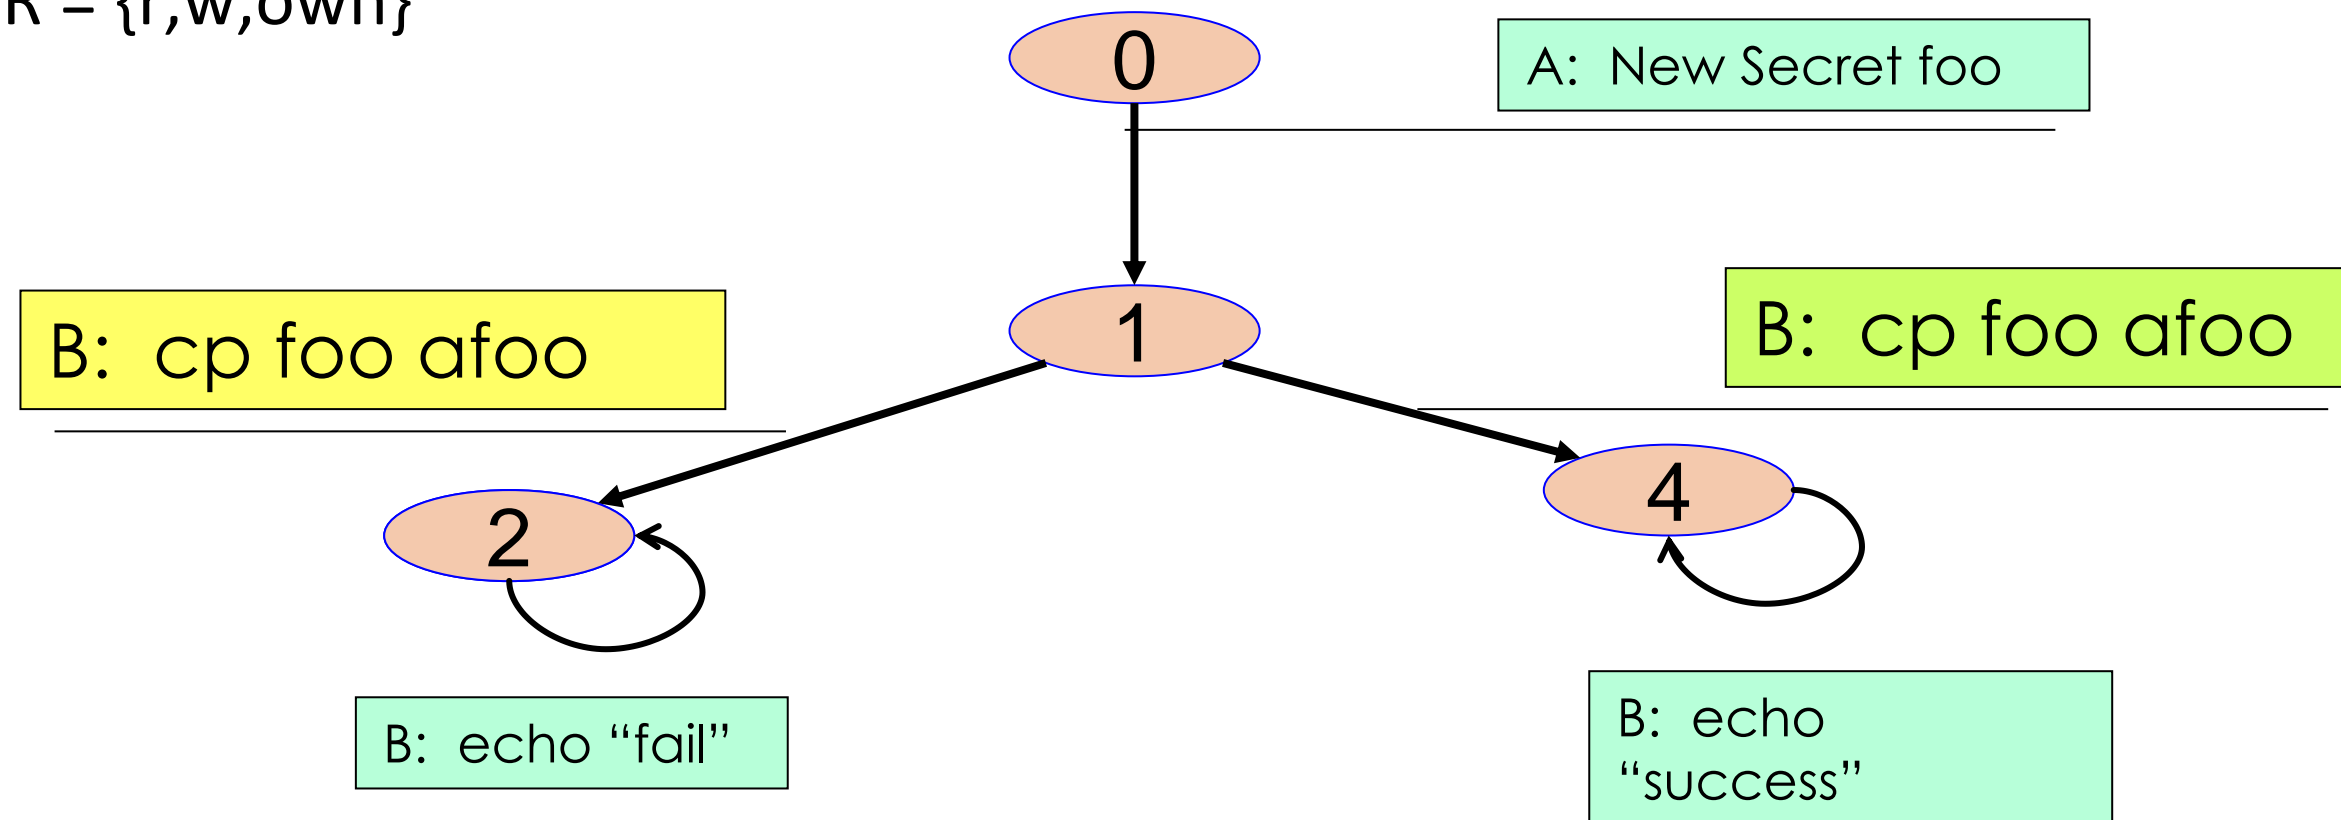
Objects $O_0 = \{\}$

AC Matrix $A_0 = \{\}$

Rights $R = \{r, w, own\}$

(S_0, O_0, A_0) A: New Secret foo (S_1, O_1, A_1)

(S_1, O_1, A_1) B: cp foo afoo



Intended State 1

Subjects $S_1 = \{A, B\}$

Objects $O_1 = \{\text{foo}\}$

AC Matrix $A_1 = \{ (A, \text{foo}, [r, w, own]), (B, \text{foo}, []) \}$

Confidentiality Scenario

Initial State

Subjects $S_0 = \{A,B\}$

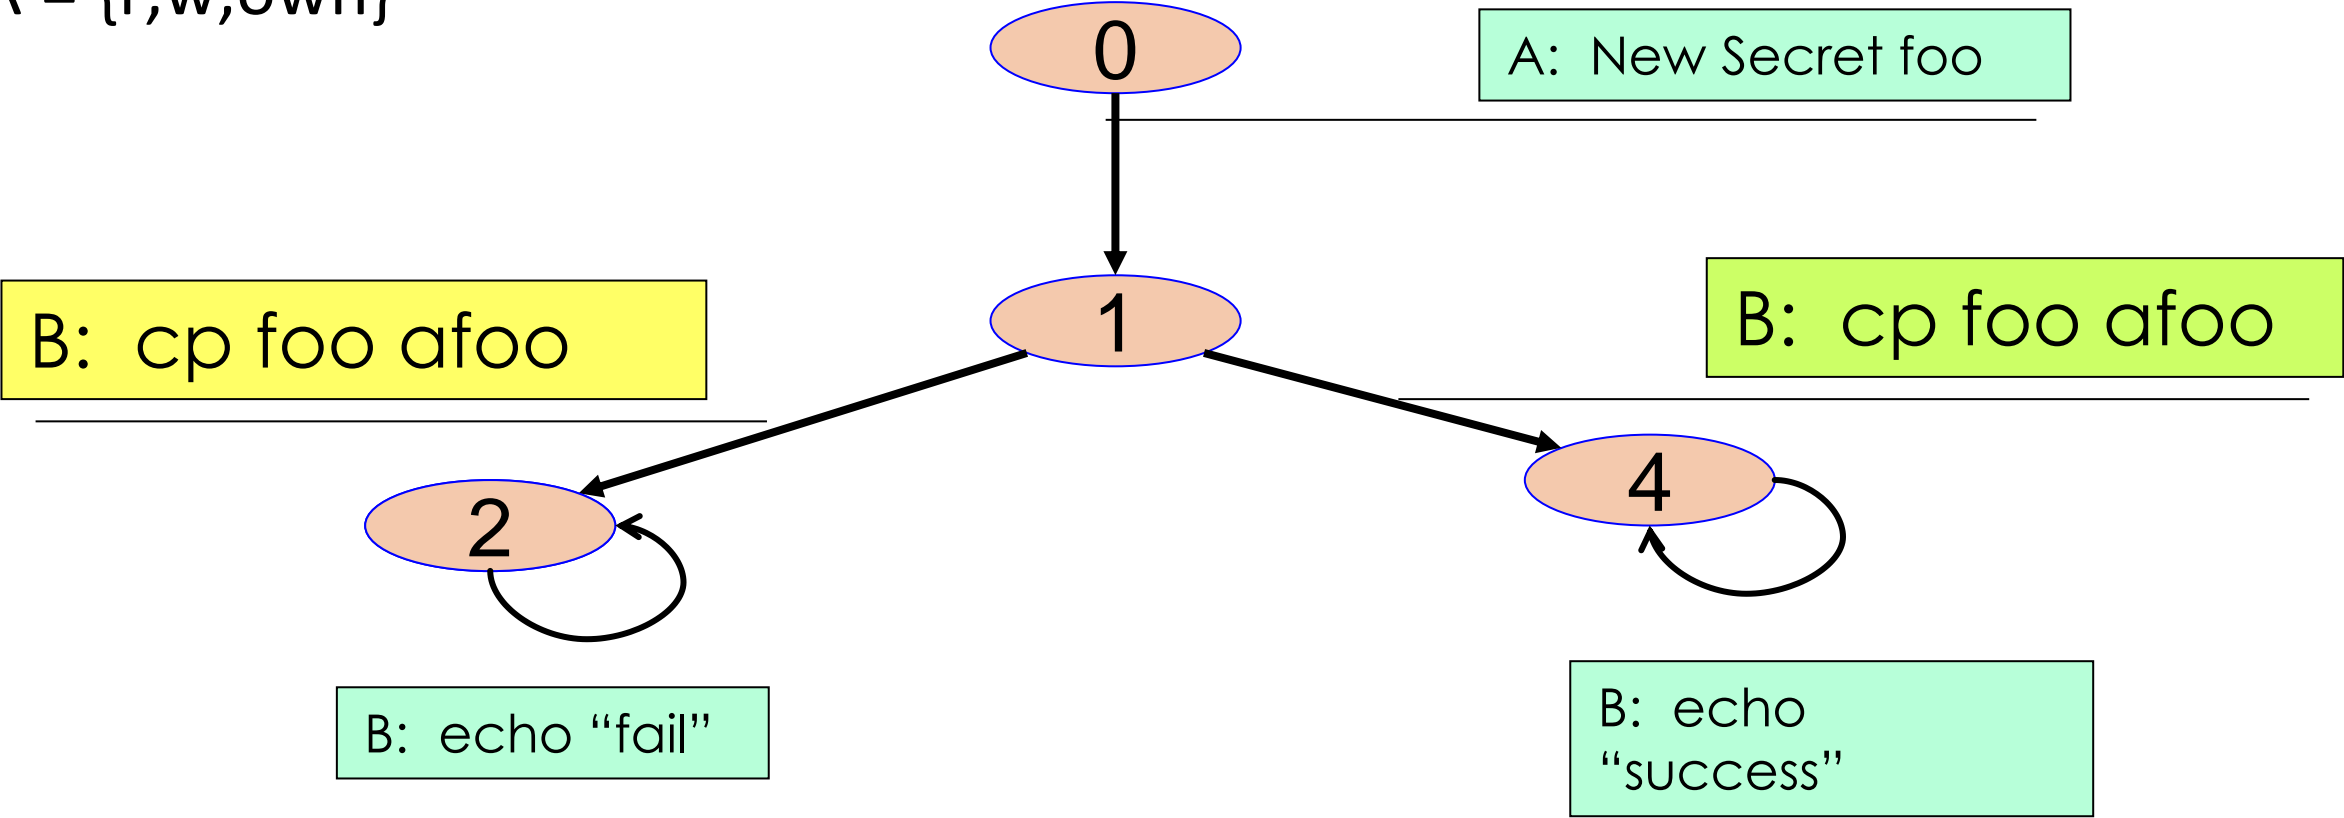
Objects $O_0 = \{\}$

AC Matrix $A_0 = \{\}$

Rights $R = \{r,w,own\}$

(S_0, O_0, A_0) A: New Secret foo (S_1, O_1, A_1)

(S_1, O_1, A_1) B: cp foo afoo $(S_?, O_?, A_?)$



Intended State 1

Subjects $S_1 = \{A,B\}$

Objects $O_1 = \{\text{foo}\}$

AC Matrix $A_1 = \{ (A,\text{foo},[r,w,own]), (B,\text{foo},[]) \}$

Confidentiality Scenario

Initial State

Subjects $S_0 = \{A, B\}$

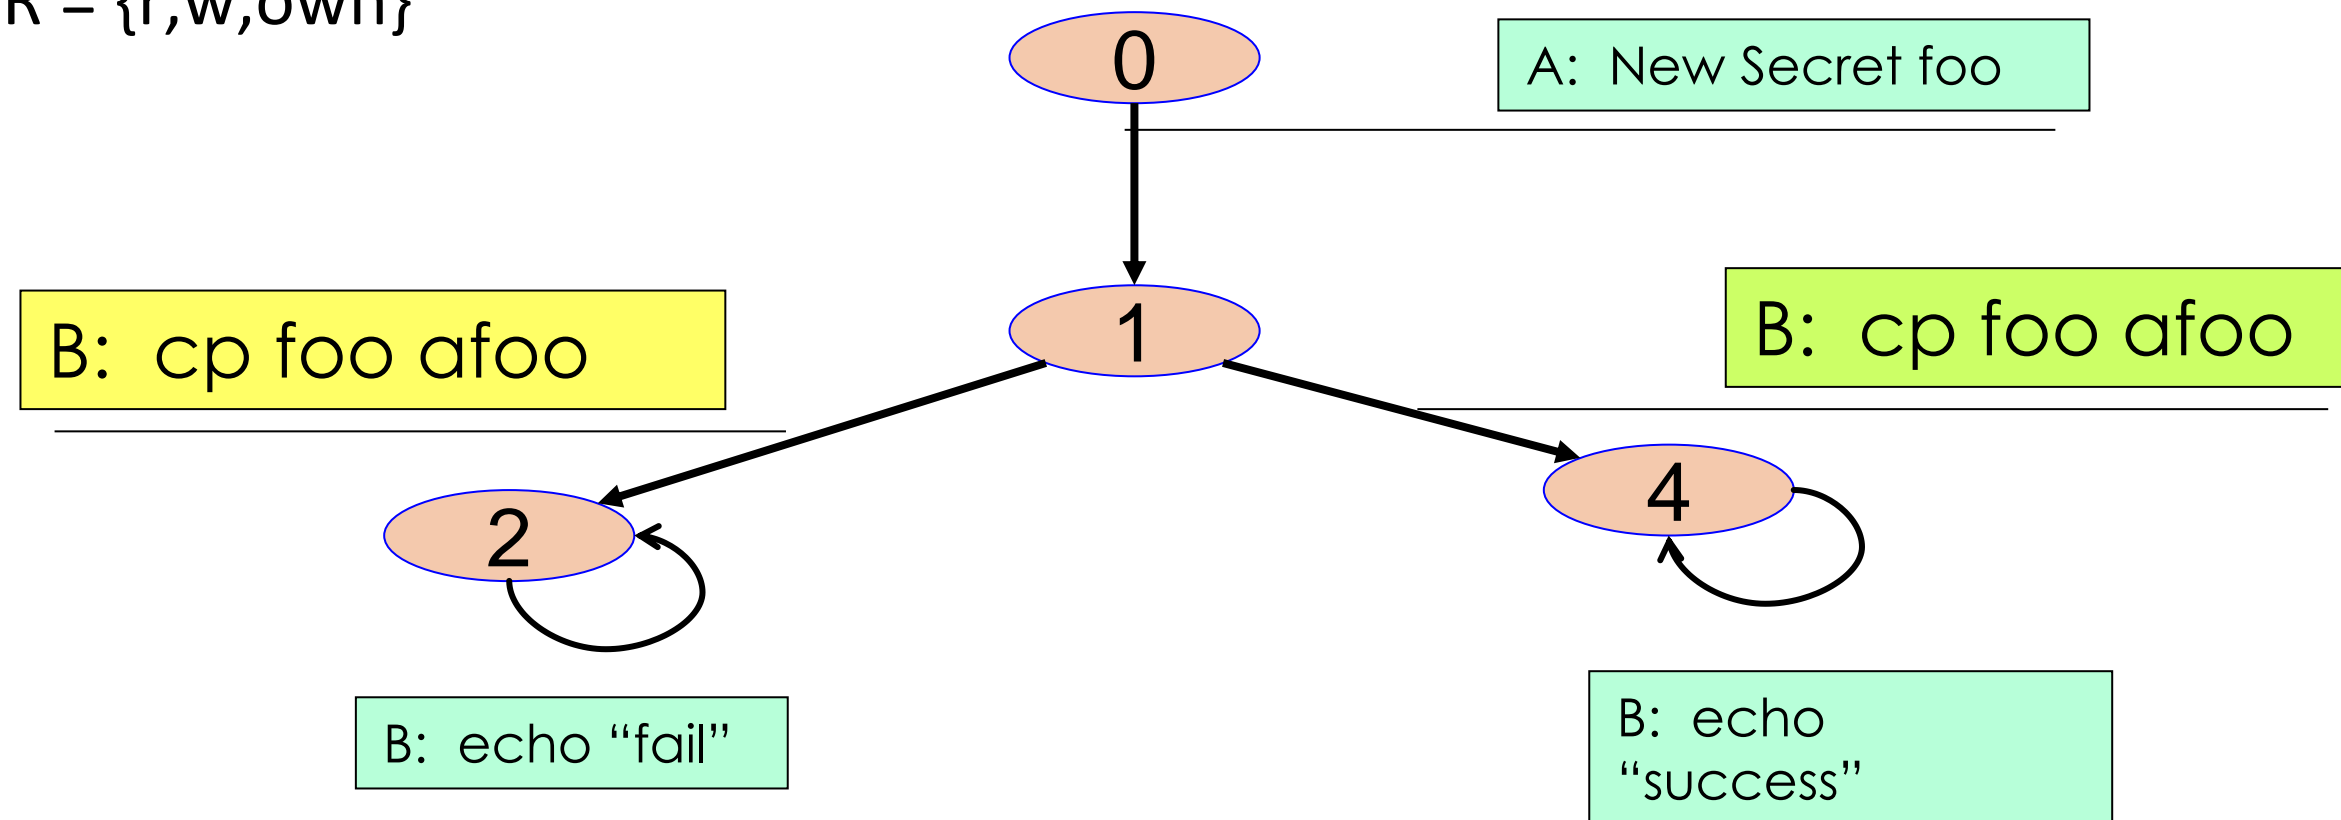
Objects $O_0 = \{\}$

AC Matrix $A_0 = \{\}$

Rights $R = \{r, w, own\}$

(S_0, O_0, A_0) A: New Secret foo (S_1, O_1, A_1)

(S_1, O_1, A_1) B: cp foo afoo $(S_?, O_?, A_?)$



Intended State 1

Subjects $S_1 = \{A, B\}$

Objects $O_1 = \{\text{foo}\}$

AC Matrix $A_1 = \{ (A, \text{foo}, [r, w, own]), (B, \text{foo}, []) \}$

Critical issue is
definition of cp

Agenda

- Access Control Matrix
 - Overview
 - Access Control Matrix Model
 - Protection State Transitions
 - Commands
 - Conditional Commands
- Foundational Results



Access Control Matrix Model

- Lampson '71 (<http://www.computerhistory.org/fellowawards/hall/bios/Butler,Lampson/>), refined by Graham and Denning ('71, '72)
- Key Concepts
 - **Objects**, the protected entities, O
 - **Subjects**, the active entities acting on the objects, S
 - **Rights**, the controlled operations subjects can perform on objects, R
 - **Access Control Matrix**, A , maps Objects and Subjects to sets of Rights
 - State: (S, O, A)

Access Control Matrix Model

- Access control matrix
 - Describes protection state *precisely*
 - Matrix describing rights of subjects
 - State transitions change elements of matrix

Description

		objects (entities)					
		o_1	...	o_m	s_1	...	s_n
subjects	s_1						
	s_2						
	...						
	s_n						

- Subjects $S = \{ s_1, \dots, s_n \}$
- Objects $O = \{ o_1, \dots, o_m \}$
- Rights $R = \{ r_1, \dots, r_k \}$
- Entries $A[s_i, o_j] \subseteq R$
- $A[s_i, o_j] = \{ r_x, \dots, r_y \}$ means subject s_i has rights r_x, \dots, r_y over object o_j

Example

- Processes p, q
- Files f, g
- Rights $r, w, x, a(ppend), o(wn)$

	f	g	p	q
p	rwo	r	$rwxo$	w
q	a	ro	r	$rwxo$

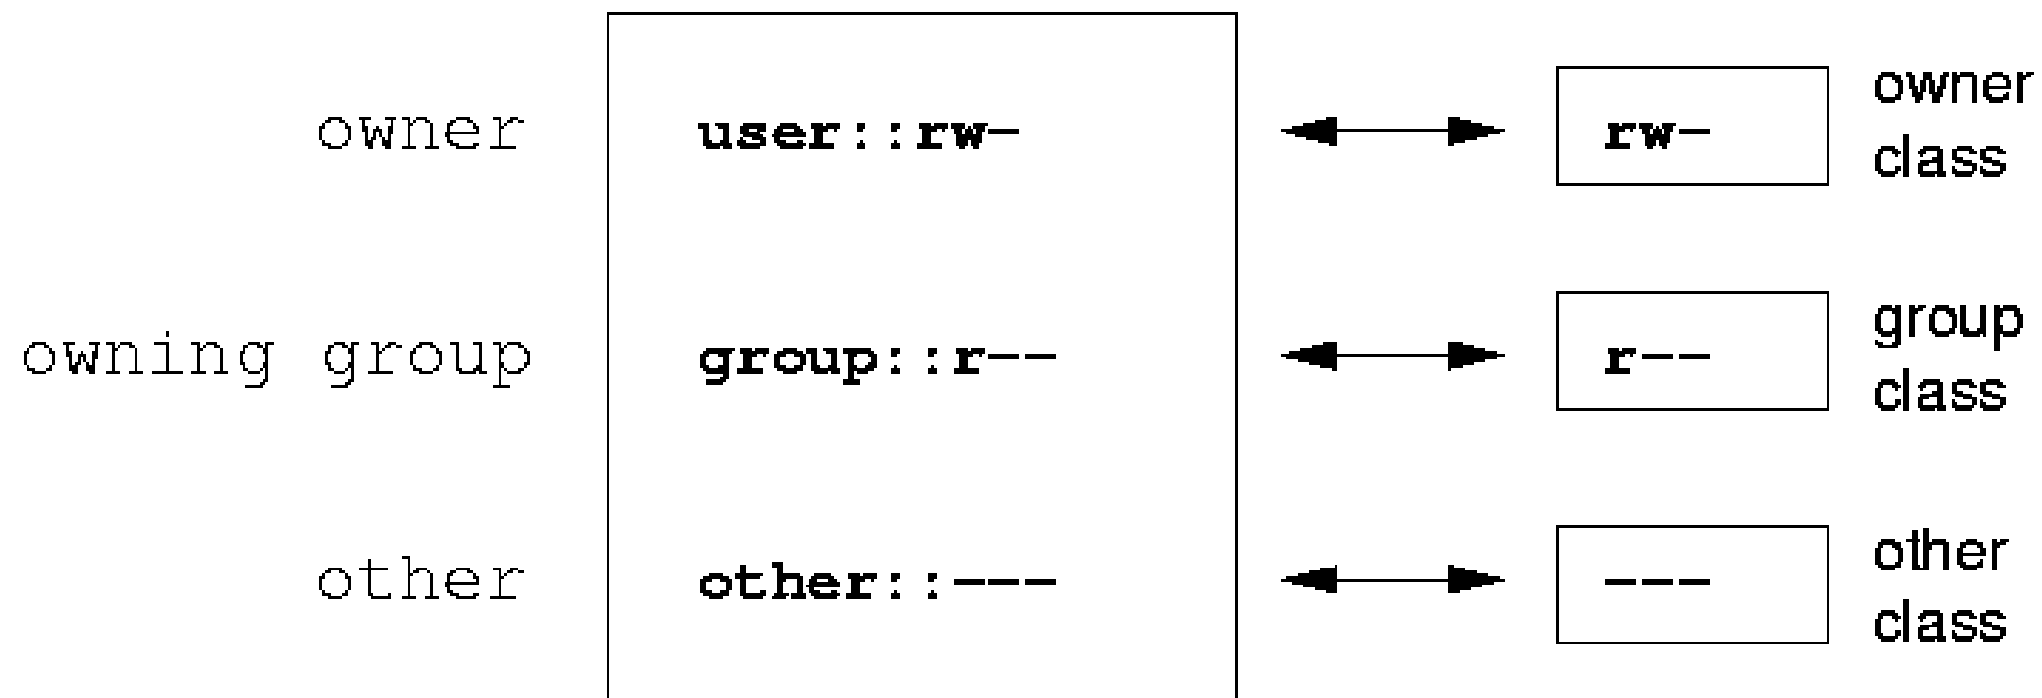
Example 2

- Procedures *inc_ctr*, *dec_ctr*, *manage*
- Variable *counter*
- Rights *+*, *−*, *call*

	<i>counter</i>	<i>inc_ctr</i>	<i>dec_ctr</i>	<i>manage</i>
<i>inc_ctr</i>	<i>+</i>			
<i>dec_ctr</i>	<i>−</i>			
<i>manage</i>		<i>call</i>	<i>call</i>	<i>call</i>

Example in Linux

```
user@user-desktop:~/cs165$ ls -l
total 32
-rw-r--r-- 1 user group 26498 Feb 17 15:19 code.tar.gz
```



Example in Linux

- `[root@target /]# ls -l /home/`
`drwxrwxr-x 2 team-temp admin 36864 ... admin`

Agenda

- Access Control Matrix
 - Overview
 - Access Control Matrix Model
 - Protection State Transitions
 - Commands
 - Conditional Commands
- Foundational Results



Protection State Transitions

- Sequences of state transitions are represented by commands that update the access control matrix (ACM)
- Primitive commands
 - **create subject s ; create object o**
 - Creates new row, column in ACM; creates new column in ACM
 - **destroy subject s ; destroy object o**
 - Deletes row, column from ACM; deletes column from ACM
 - **enter r into $A[s, o]$**
 - Adds r rights for subject s over object o
 - **delete r from $A[s, o]$**
 - Removes r rights from subject s over object o

Creating File

- Process p creates file f with r and w permission

```
command create•file( $p$ ,  $f$ )  
    create object  $f$ ;  
    enter own into  $A[p, f]$ ;  
    enter  $r$  into  $A[p, f]$ ;  
    enter  $w$  into  $A[p, f]$ ;  
end
```

Mono-Operational Commands

- Make process p the owner of file g

```
command make•owner( $p$ ,  $g$ )  
  enter own into  $A[p, g]$ ;  
end
```

- Mono-operational command
 - Single primitive operation in this command

Conditional Commands

- Let p give q r rights over f , if p owns f
command $grant \cdot read \cdot file \cdot 1(p, f, q)$
 if own **in** $A[p, f]$
 then
 enter r **into** $A[q, f];$
 end
- Mono-conditional command
 - Single condition in this command

Multiple Conditions

- Let p give q (r and w) rights over f , if p owns f and p has c rights over q

```
command grant•read•file•2( $p, f, q$ )  
  if own in  $A[p, f]$  and  $c$  in  $A[p, q]$   
  then  
    enter  $r$  into  $A[q, f]$ ;  
    enter  $w$  into  $A[q, f]$ ;  
end
```

Special Rights

- The copy right (or grant right) allows the possessor to grant rights to another
- The own right enables the possessors to add or delete privileges for themselves and others
- Principle of attenuation of Privilege
 - A subject may not give rights it does not possess to another

Short Summary: Key Points

- Access control matrix simplest abstraction mechanism for representing protection state
- Transitions alter protection state
- 6 primitive operations alter matrix
 - Transitions can be expressed as commands composed of these operations and, possibly, conditions

Outline

- Access Control Matrix
 - Overview
 - Access Control Matrix Model
 - Protection State Transitions
 - Commands
 - Conditional Commands
- **Foundational Results**



Safety



- Let R be the set of generic (primitive) rights of the system

Safety



- Let R be the set of generic (primitive) rights of the system
- Definition: when a generic right r is added to an element of the access control matrix not already containing r , that right r is said to be *leaked*



Safety



- Let R be the set of generic (primitive) rights of the system
- Definition: when a generic right r is added to an element of the access control matrix not already containing r , that right r is said to be *leaked*
- Definition: If a system can never leak right r , the system is called *safe* with respect to the right r . If the system can leak right r , the system is called *unsafe* with respect with the right r



Example: project 3

- On cs165-internal,
- Unprivileged users: team0, team1, ...
- Privileged user: admin
- File: /home/admin/

	/home/admin/
team0	read
team1	read
...	read
admin	read/write

Outline

- Access Control Matrix
- Other Security Policy Models



Other Types of Security Policies

- Military (governmental) security policy
 - Policy primarily protecting **confidentiality**
- Commercial security policy
 - Policy primarily protecting **integrity**
- **Confidentiality** policy
- **Integrity** policy
- More abstract and conceptual

Confidentiality Policy

- Goal: prevent the unauthorized disclosure of information
 - Deals with information flow
 - Extensive redundancy in military makes integrity/availability less of a problem

Confidentiality Policy

- Goal: prevent the unauthorized disclosure of information
 - Deals with information flow
 - Extensive redundancy in military makes integrity/availability less of a problem
- Multi-level security (MLS) models are best-known examples
 - **Bell-LaPadula Model (BLP model), 1970s**

Bell-LaPadula Model

- Security levels arranged in linear ordering




- Subjects have *security clearance* $L(s)$
- Objects have *security classification* $L(o)$

Summary of BLP Model

Policy (Write/Create): Information flows up, not down

“Write up” allowed, “write down” disallowed

“Read down” allowed, “read up” disallowed



Example

<i>security level</i>	<i>subject</i>	<i>object</i>
Top Secret	Terry	Personnel Files
Secret	Sam	E-Mail Files
Confidential	Charlie	Activity Logs
Unclassified	Umed	Telephone Lists

- Terry can read all files
- Charlie cannot read Personnel or E-Mail Files
- Umed can only read Telephone Lists

Lattice Model



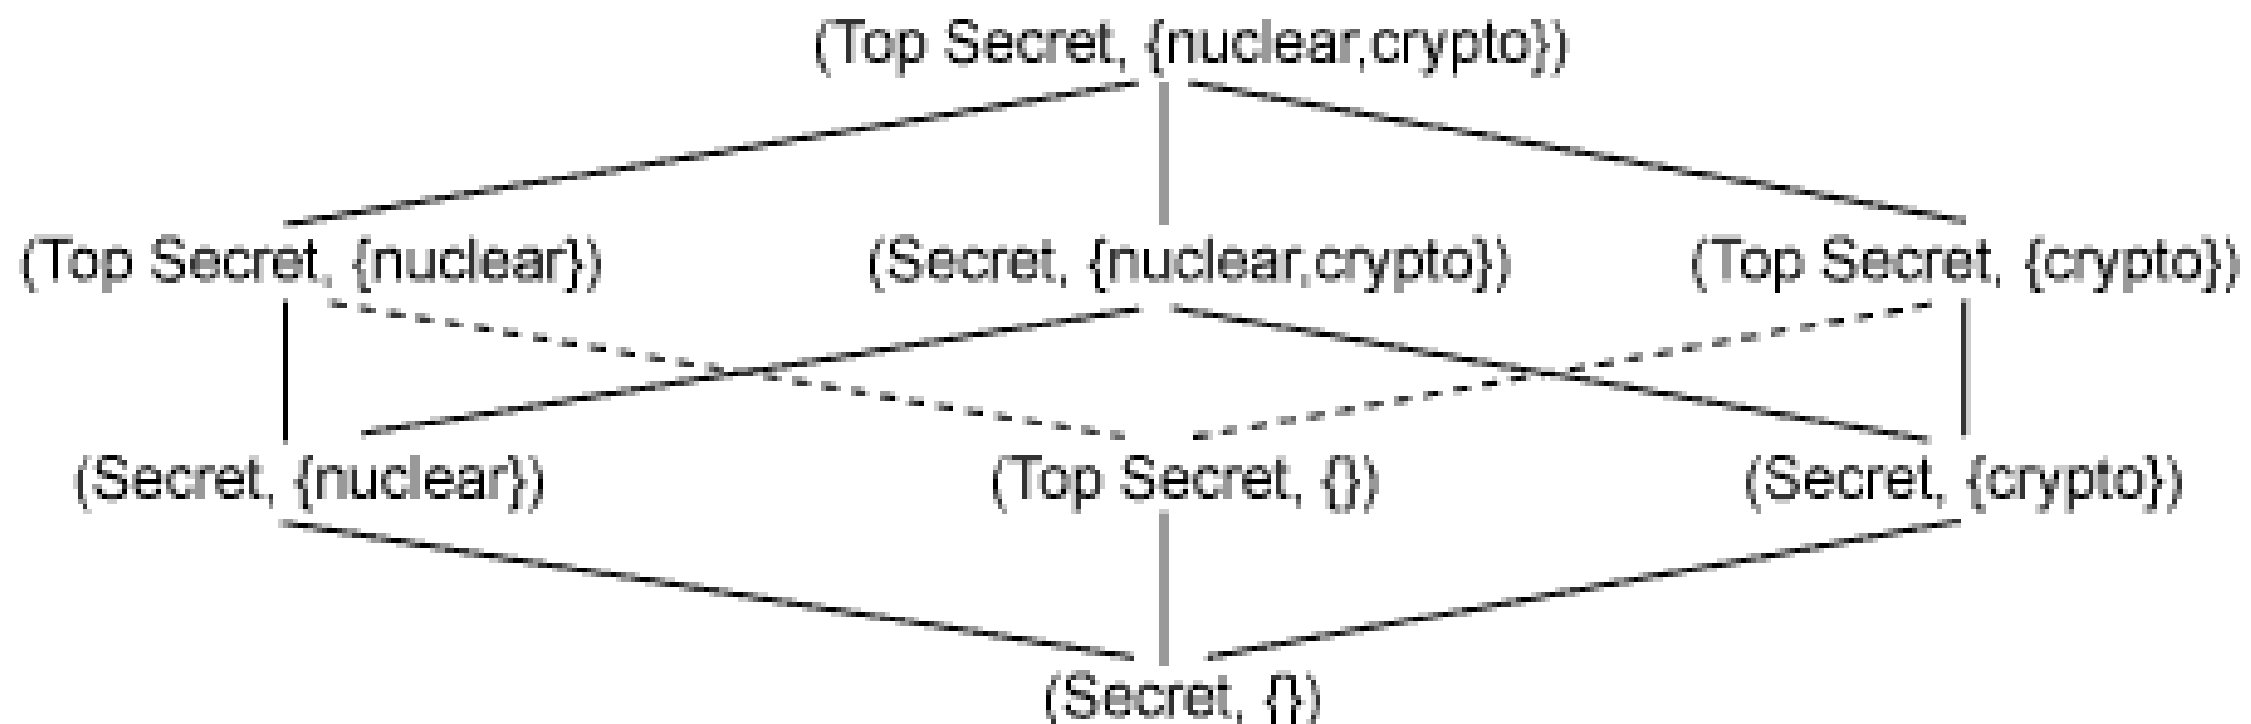
- Used by the US military (and many others), the Lattice model uses MLS to define policy
- Expand notion of security level to include categories
 - Categories (actually unbounded set)
 - NUC(lear), INTEL(igence), CRYPTO(graphy)
 - Note that these levels are used for physical documents in the governments as well.
- Security level is (*clearance, category set*), or formally (L, C) where L is the clearance level, and C is the set of categories

Lattice Model – Label-based Access Control

- Examples
 - Alice: (TOP SECRET, {CRYPTPO})
 - Bob: (SECRET, {NUC})
 - Charlie: (TOP SECRET, {CRYPTO, NUC})
 - DocA: (SECRET, {CRYPTO, NUC})
 - DocB: (SECRET, {CRYPTO})

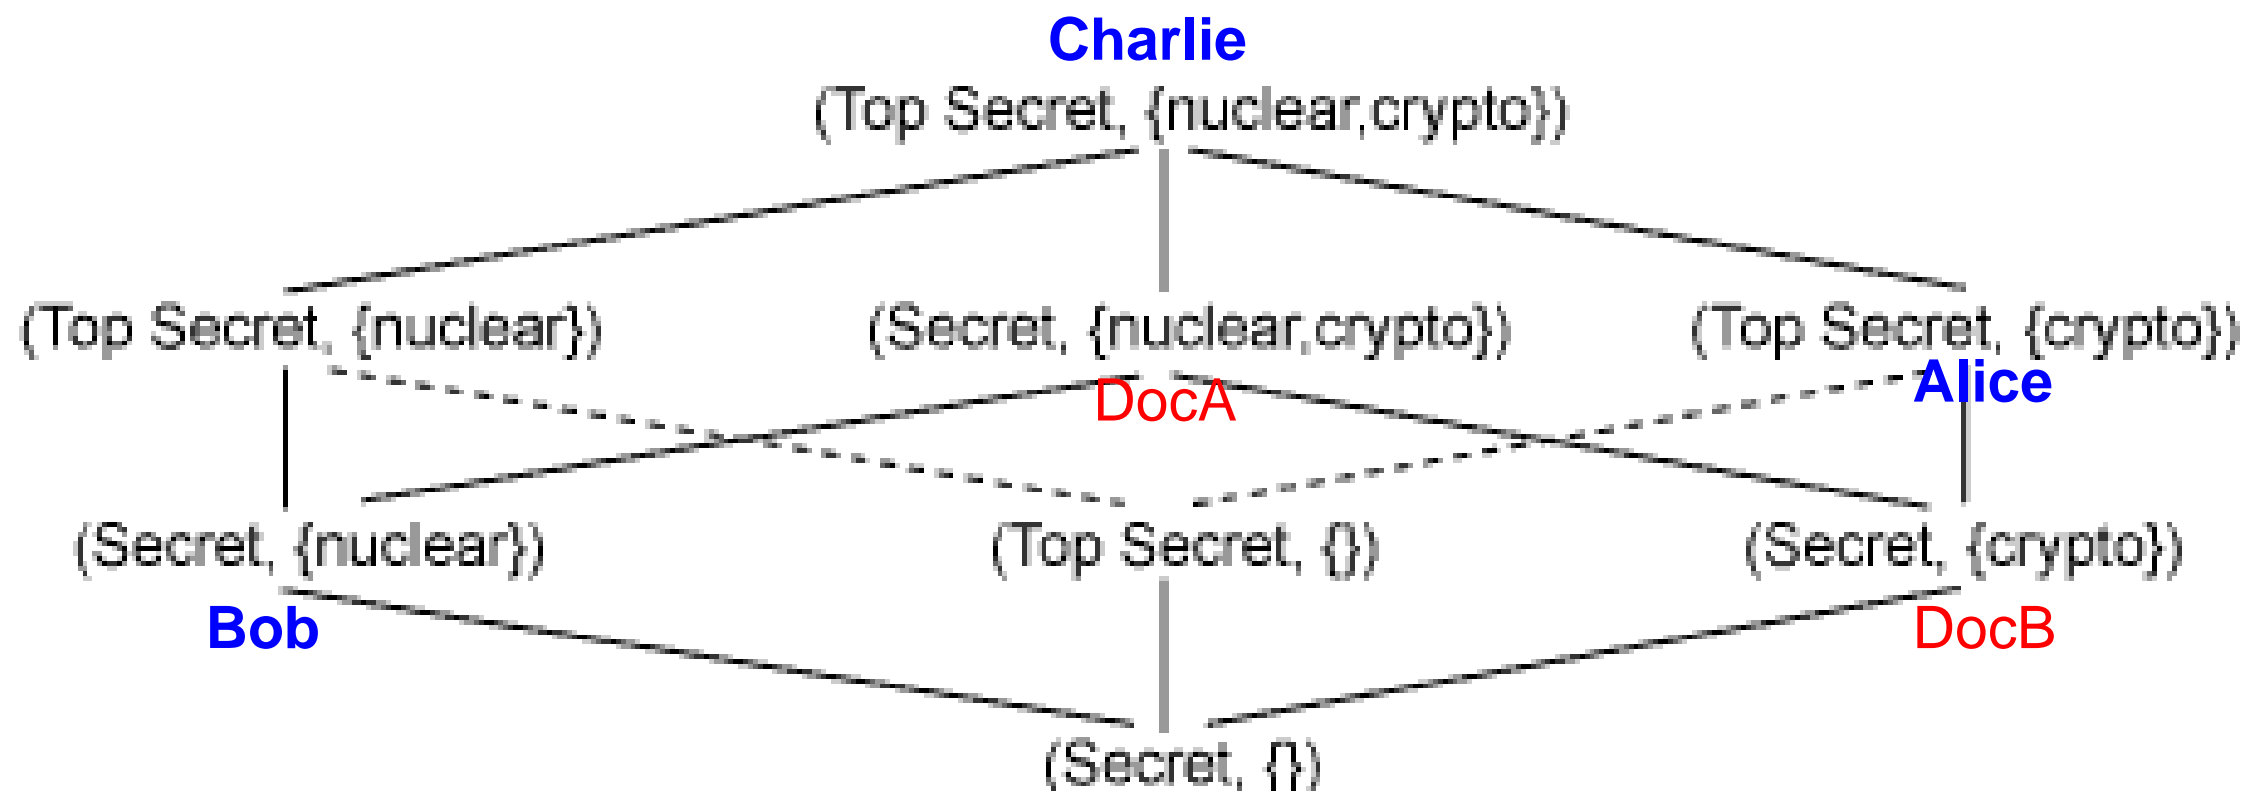
Lattice Model – Label-based Access Control

- Examples
 - Alice: (TOP SECRET, {CRYPTPO})
 - Bob: (SECRET, {NUC})
 - Charlie: (TOP SECRET, {CRYPTO, NUC})
 - DocA: (SECRET, {CRYPTO, NUC})
 - DocB: (SECRET, {CRYPTO})



Lattice Model – Label-based Access Control

- Examples
 - Alice: (TOP SECRET, {CRYPTO})
 - Bob: (SECRET, {NUC})
 - Charlie: (TOP SECRET, {CRYPTO, NUC})
 - DocA: (SECRET, {CRYPTO, NUC})
 - DocB: (SECRET, {CRYPTO})




Integrity Policy

- MLS talks about who can “read” a document (*confidentiality*)
- *Integrity* is considered who can “write or make change to” a document
 - Thus, who can effect the integrity (content) of a document
 - Example: You may not care who can read DNS records, but you better care who changes them!

BIBA Model

- Biba defined a dual of secrecy for integrity
 - “no read down, no write up”



In Reality

- High-integrity programs may need access to low-integrity data
 - Command line input
 - Web server accepting data from the Internet
 - Attacker-controlled files
 - SSH daemon runs as root and reads `/home/[user]` directory
- Or ...
 - Sometimes programs may expect high-integrity objects mistakenly

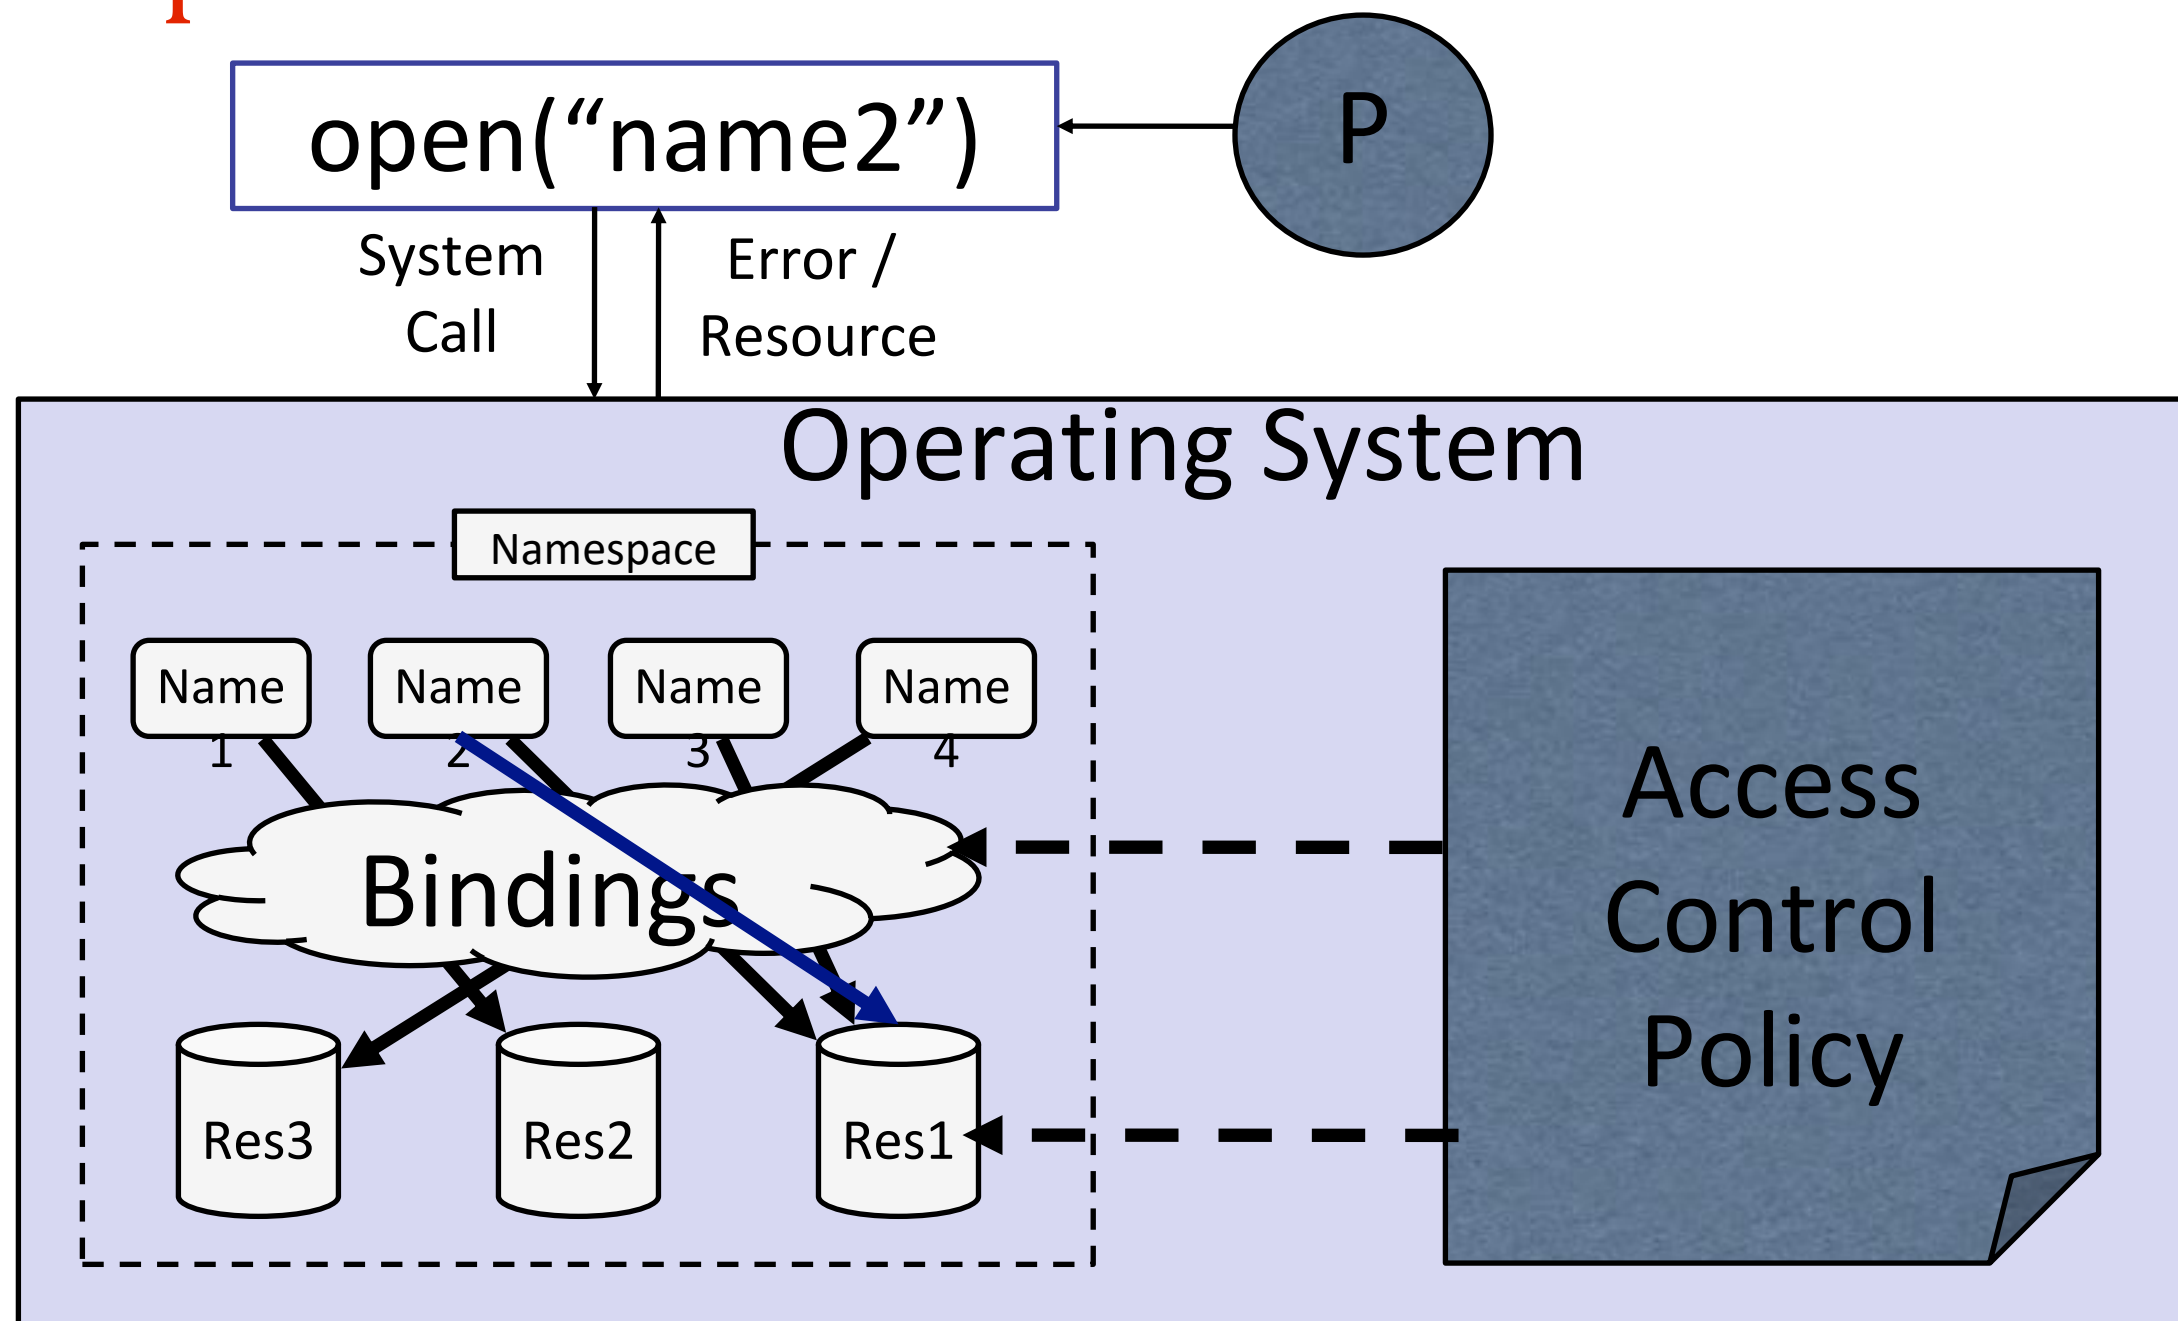
Example Integrity Problem: Resource Access

- Programs require a variety of **resources** to function
 - *Files*
 - *Inter-process communication channels*
 - *Signals*
- How hard can accessing resources securely be?
 - Just a simple **open(filename)**, right?
 - **Wrong!**



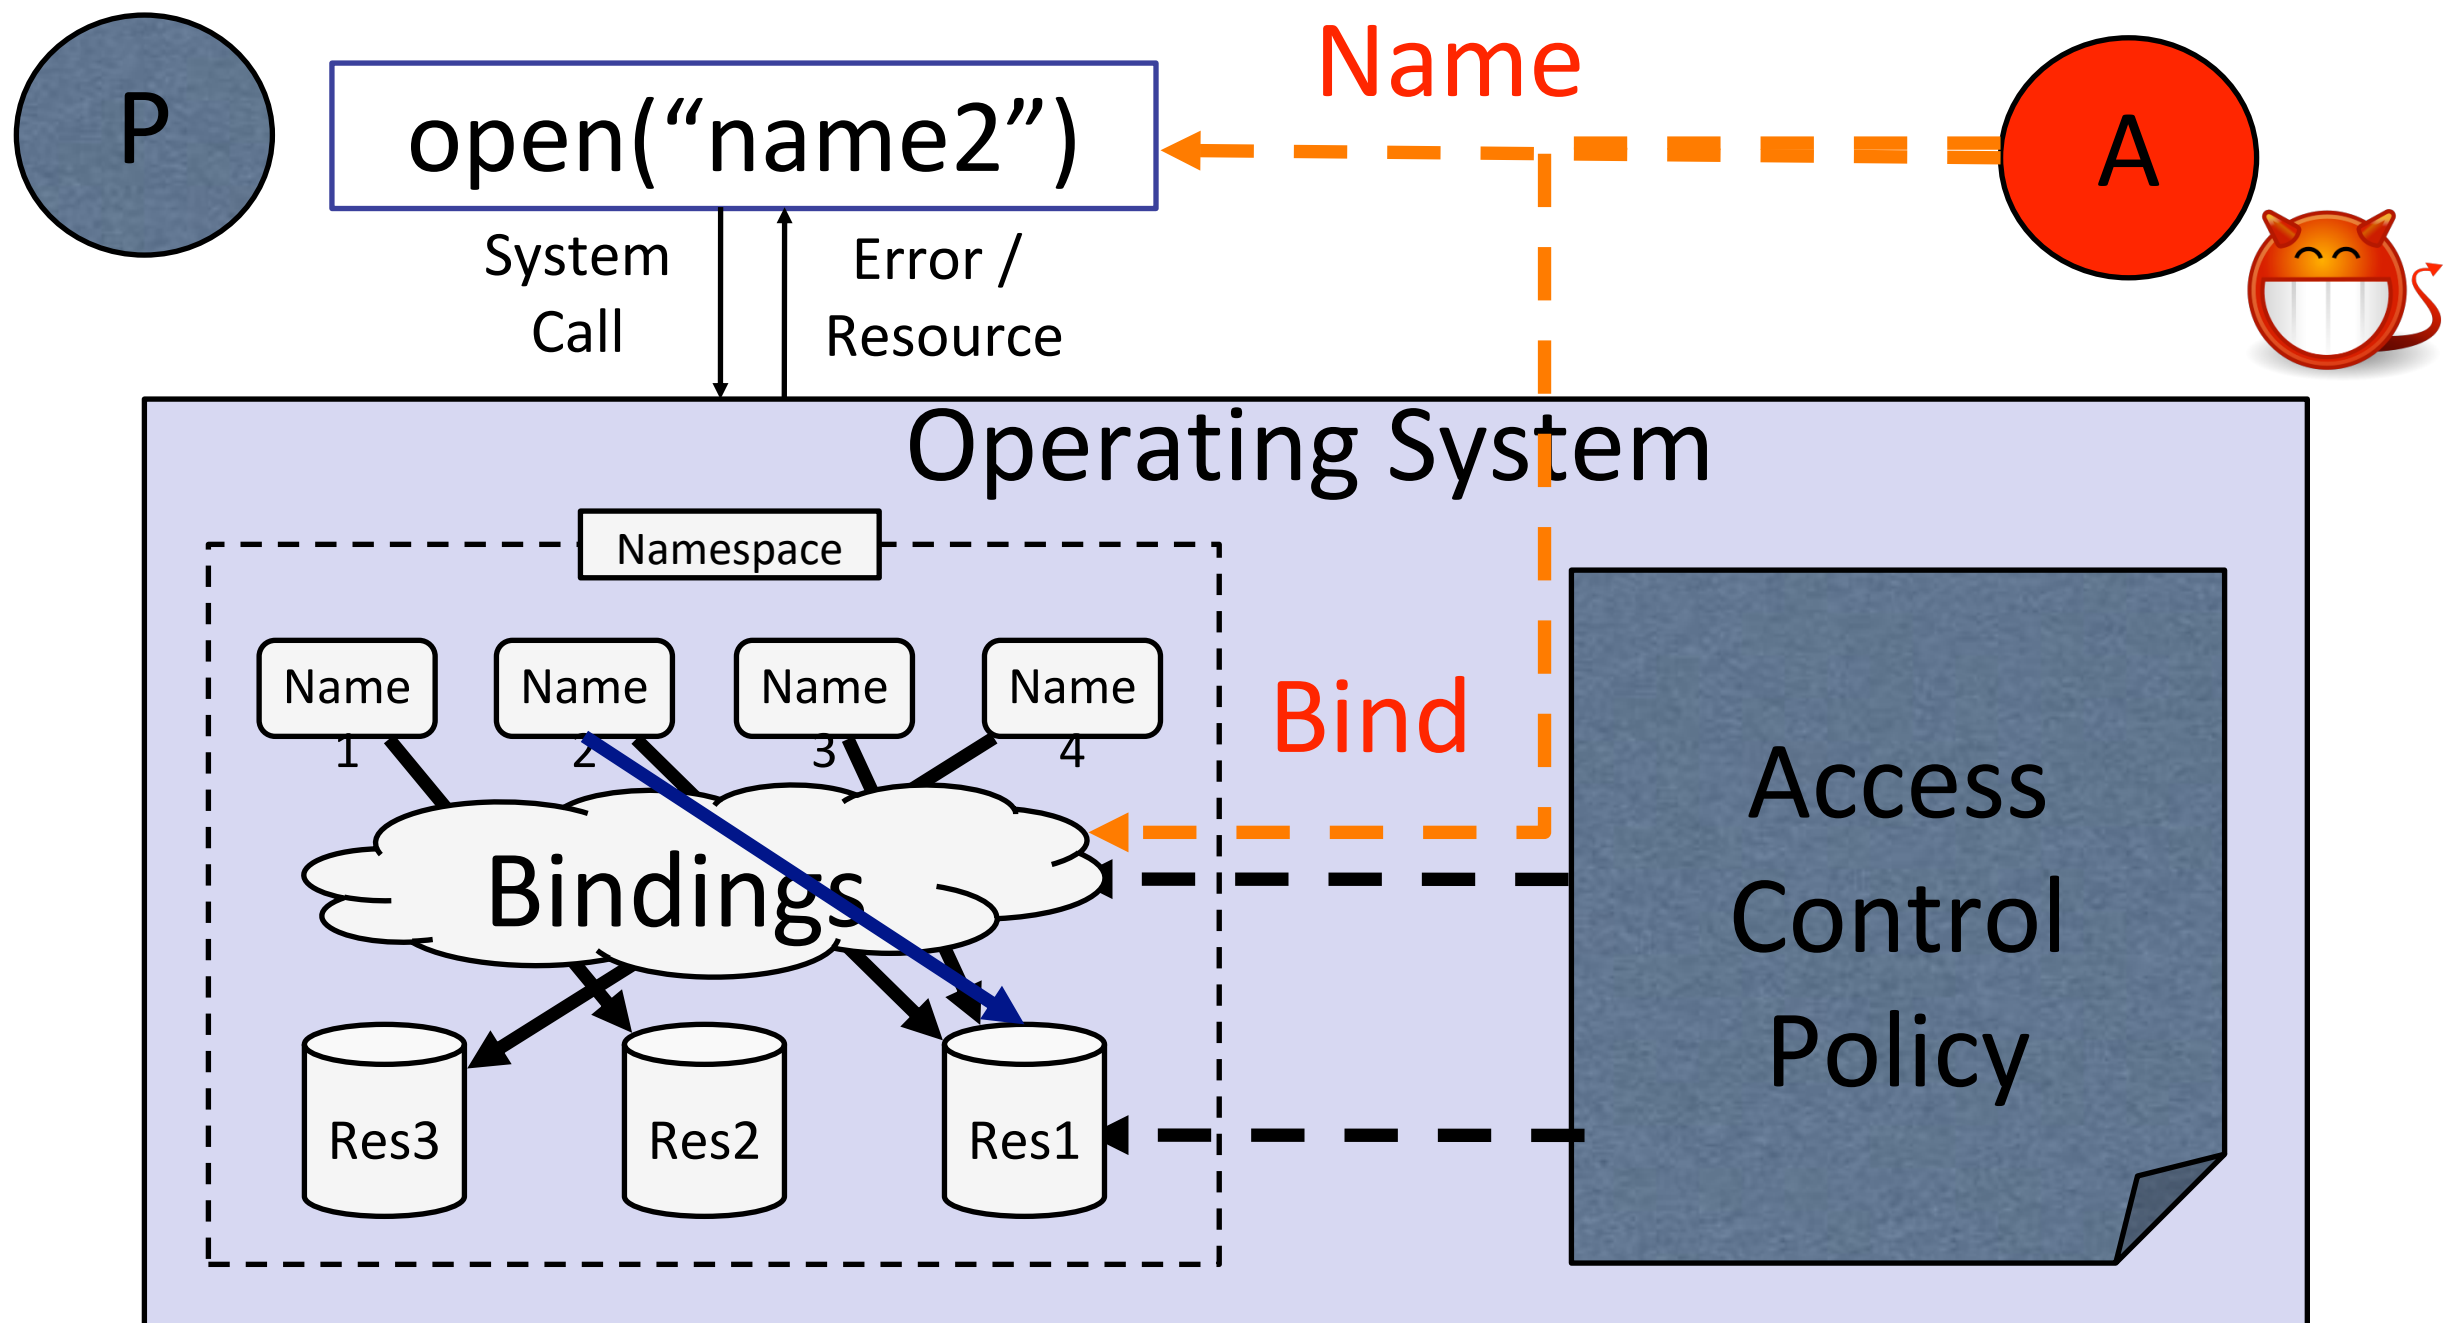
Resource Access

- **Inputs:** name, bindings (dirs, symbolic links)
- **Output:** resource



Resource Access Attacks

- Adversaries control inputs -- name, bindings to direct victim programs to malicious resources



Adversary Control


- Who is an adversary?
 - Determined by **threat model** and **access control policy**
 - Based on access control policy
 - UID x has as adversary any UID $y \neq x$ (except superuser *root*)
- Adversary control of name or binding



Link Following

- Link following attack
 - Adversary controls **bindings** to **redirect** victim to a resource **not under** adversary's control (confused deputy)

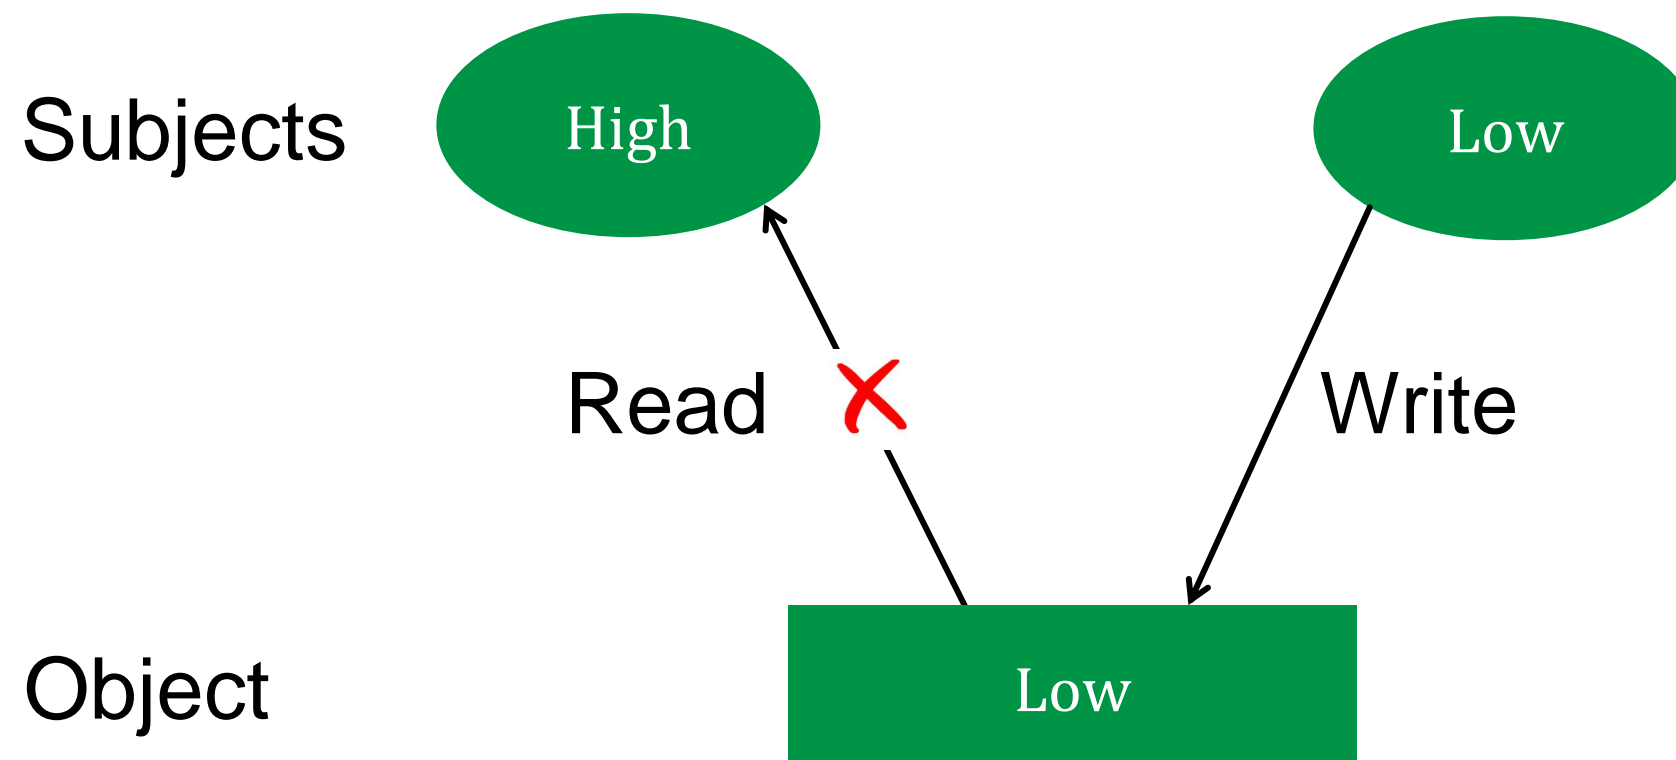
```
drwxrwsr-x  2 root mail  4096 2011-01-13 10:06 mail
```



Link Following

<i>Integrity level</i>	<i>Subject</i>	<i>Object</i>
High	root process	
Low	mail process	/var/mail/root

Victim expects high integrity, gets low integrity instead



Link Following

Access control matrix point-of-view: read access leaked

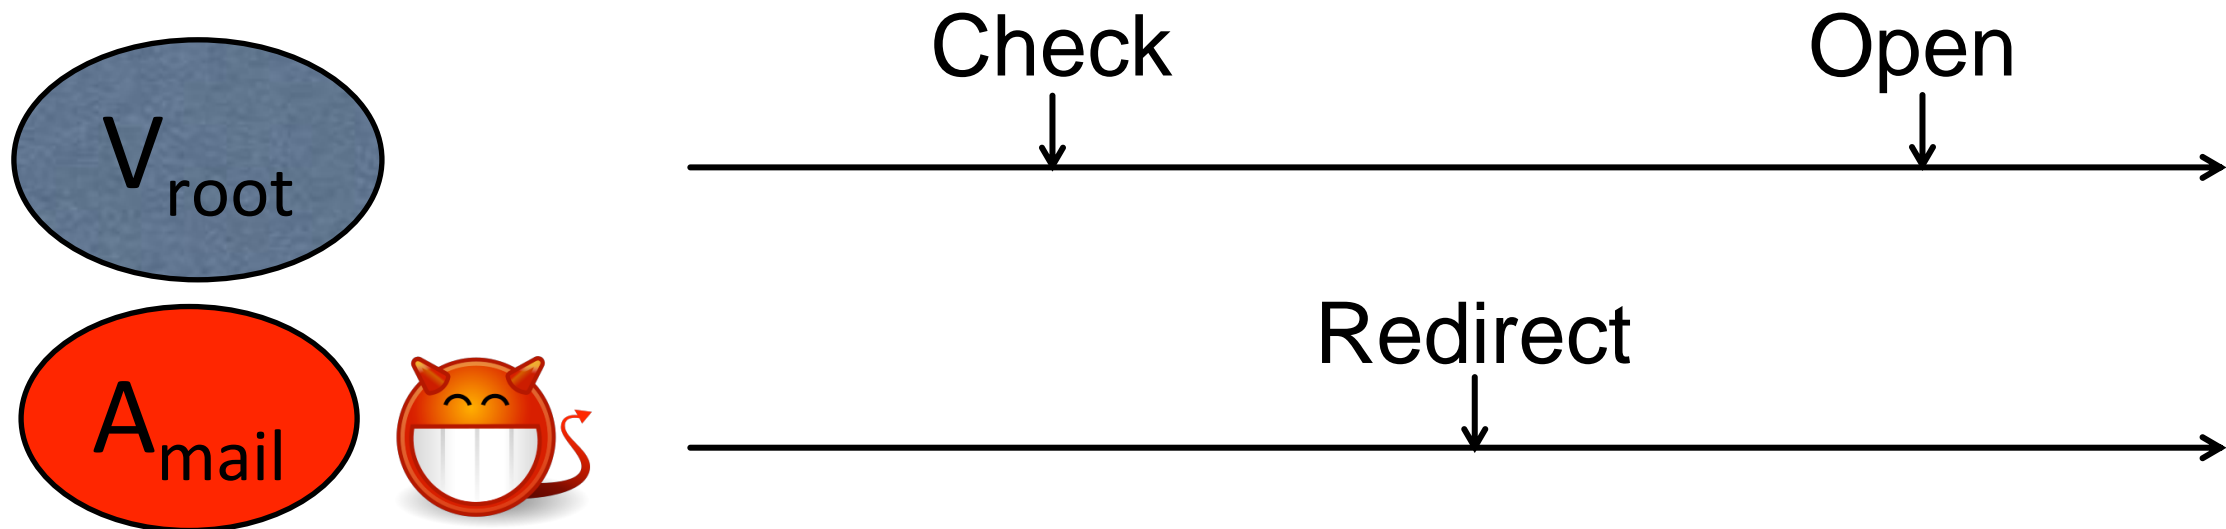
	/var/mail/root	/etc/passwd
root (high-privilege)	read/write	read/write
mail	read/write	no access -> read

Achieved through two commands:

1. `rm /var/mail/root`
2. `ln -s /var/mail/root /etc/passwd`


Link Following

- Can the root process protect itself by checking the low-integrity input?
 - Whether `/var/mail/root` is a *symbolic link*?
 - Unfortunately, it doesn't work




Race Conditions

- Race Conditions
 - Adversary exploits non-atomicity in “check” and “use” of resource
 - Time-of-check, time-of-use (TOCTTOU) attacks



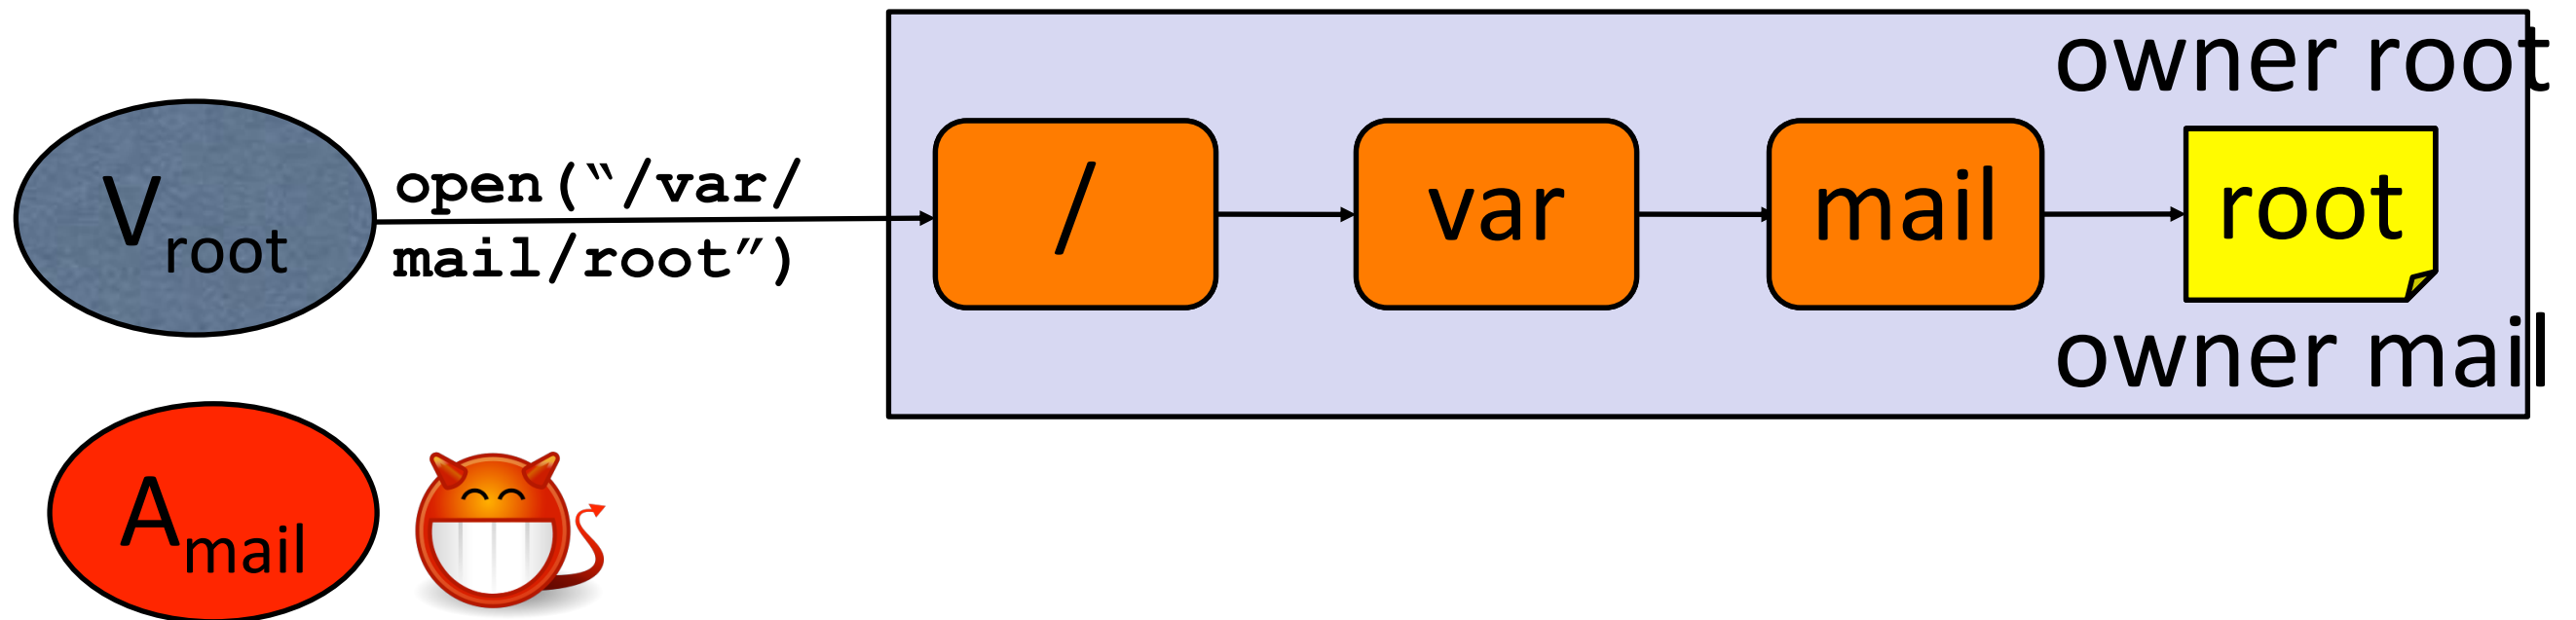
Race Conditions

- Race Conditions
 - Adversary exploits non-atomicity in “check” and “use” of resource
 - Time-of-check, time-of-use (TOCTTOU) attacks




Another attack: File Squatting

- Adversary controls **bindings** to direct victim to an adversary accessible (low integrity) resource
 - Can trick the victim to read fake emails!
- Victim **expects** high integrity



Yet another attack: Directory Traversal

- Adversary controls the **name** (as opposed to **binding**) to direct victim to a high-secrecy resource
- **Root cause?**



Fundamental Problem

- When a program expects a high-integrity resource but gets low-integrity one
- Or when a program expects low-integrity resource but did not perform adequate checks
- They are all attack surface!

Questions

