

CS165 – Computer Security

Vulnerability discovery & static analysis
Nov 4, 2021

Our Goal

- How to exploit a vulnerability?
- How to find them?



Our Goal

- How to exploit a vulnerability?
- How to find them?



Vulnerability

- How do you define computer “vulnerability”?
 - *Flaw*
 - *Accessible to adversary*
 - *Adversary has ability to exploit*



Vulnerability

- How do you define computer “vulnerability”?
 - *Flaw – Can we find flaws in source code?*
 - *Accessible to adversary – Can we find what is accessible*
 - *Adversary has ability to exploit – Can we find how to exploit?*



Vulnerability

- How do you define computer “vulnerability”?
 - *Flaw – Can we find flaws in source code?*
 - *Accessible to adversary – Can we find what is accessible (attack surface)?*
 - *Adversary has ability to exploit – Can we find how to exploit?*



Vulnerability

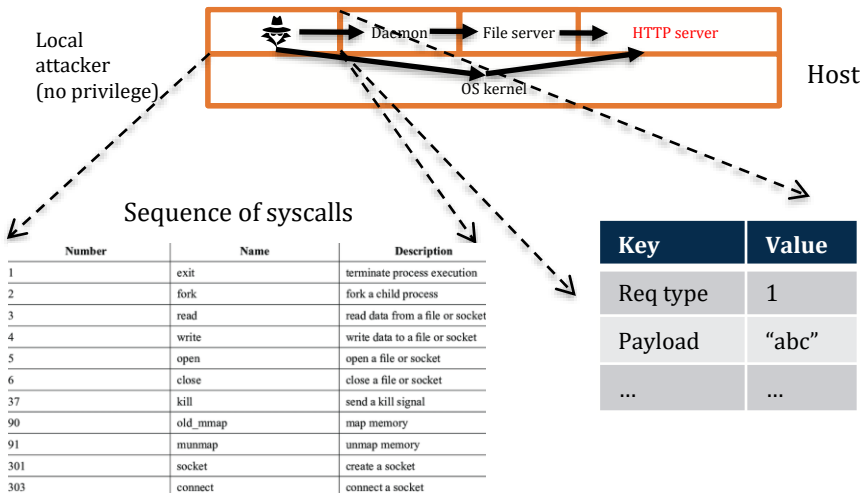
- How do you define computer “vulnerability”?
 - *Flaw – Can we find flaws in source code?*
 - *Accessible to adversary – Can we find what is accessible (attack surface)?*
 - *Adversary has ability to exploit – Can we find how to exploit (highly dependent on types of flaws, e.g., buffer overflow)?*



One Approach

- Run the program on various **low-integrity** inputs (i.e., through identified **attack surface**)
 - See what happens
 - Maybe you will find a flaw
- How should you choose inputs?

Analyzing Attack Surface



Dynamic Analysis Options

- Regression Testing
 - Run program on many **normal** inputs and look for bad behavior in the responses
 - Typically looking for behavior that differs from expected – e.g., a previous version of the program
- Fuzz Testing
 - Run program on many **abnormal** inputs and look for bad behavior in the responses
 - Looking for behaviors that may be triggered by adversaries
 - Bad behaviors are typically crashes caused by memory errors

Dynamic Analysis Options

- Which approach is more likely to find vulnerabilities?
- Why?

Fuzz Testing

- Fuzz Testing
 - Idea proposed by Bart Miller at Wisconsin in 1988
- **Problem:** People assumed that utility programs could correctly process any input values
 - Available to all
- Found that they could crash 25-33% of UNIX utility programs

Fuzz Testing

- Fuzz Testing
 - Idea proposed by Bart Miller at Wisconsin in 1988
- Approach
 - Generate random inputs
 - Run lots of programs using random inputs
 - Identify crashes of these programs
 - Correlate with the random inputs that caused the crashes
- **Problems:** Not checking returns, Array indices...

```
1. inp=`perl -e '{print "A"x8000}'`  
2. for program in /usr/bin/*; do  
3.     for opt in {a..z} {A..Z}; do  
4.         timeout -s 9 1s  
           $program -$opt $inp  
5.     done  
6. done
```

1009 Linux programs. 13 minutes. 52
new bugs in 29 programs.

Example

```
#include <string.h>
int main(int argc, char **argv) {
    char buf[64];
    strcpy(buf, argv[1]);
}
```

Along as the input is longer than 64, there will be a bug found!

Challenges

- **Idea**: Search for possibly accessible and exploitable flaws in a program by running the program under a variety of inputs
- **Challenge**: Selecting input values for the program
 - What should be the goals in choosing input values for dynamic analysis?

Challenges

- **Idea**: Search for possibly accessible and exploitable flaws in a program by running the program under a variety of inputs
- **Challenge**: Selecting input values for the program
 - What should be the goals in choosing input values for dynamic analysis?
 - *Find all exploitable flaws*
 - *With the fewest possible inputs*
- How should these goals impact input choices?

Black Box Fuzzing

- Like Miller – Feed the program random inputs and see if it crashes
- **Pros:** Easy to configure
- **Cons:** May not search efficiently
 - May re-run the same path over again (low coverage)
 - May be very hard to generate inputs for certain paths (checksums, hashes, restrictive conditions)
 - May cause the program to terminate for logical reasons – fail format checks and stop

Black Box Fuzzing

- Example

```
function( int type, char *buf )  
{  
    if ( type == MAGIC_NUMBER1) {  
        if ( check_format( buf )) {  
            update( buf );  
        }  
    }  
}
```

Mutation-Based Fuzzing

- Supply a well-formed input
 - Generate random changes to that input
- No assumptions about input
 - Only assumes that variants of well-formed input may be problematic
- Example: zzuf
 - <http://sam.zoy.org/zzuf/>
 - Reading: The Fuzzing Project Tutorial
 - <https://fuzzing-project.org/tutorials.html>

Mutation-Based Fuzzing

- Example: zzuf
 - <http://sam.zoy.org/zzuf/>
- The Fuzzing Project Tutorial
 - `zzuf -s 0:1000000 -c -C 0 -q -T 3 objdump -x win9x.exe`
 - Fuzzes the program `objdump` using the sample input `win9x.exe`
 - Try 1M seed values (-s) from command line (-c) and keep running if crashed (-C 0) with timeout (-T 3)

Mutation-Based Fuzzing

- Easy to setup, and not dependent on program details
- But may be strongly biased by the initial input
- Still prone to some problems
 - May re-run the same path over again (same test)
 - May be very hard to generate inputs for certain paths (checksums, hashes, restrictive conditions)

Generation-Based Fuzzing


- Generational fuzzer generate inputs “from scratch” rather than using an initial input and mutating
- However, to overcome problems of naïve fuzzers they often need a format or protocol spec to start
 - GET /blabla HTTP/1.0
{Opcode: [3-5 bytes] [space] [n bytes] [4 bytes of characters] ... }
 - Examples fuzzers include
 - SPIKE, Peach Fuzz
- However format-aware fuzzing is cumbersome, because you'll need a fuzzer specification for every input format you are fuzzing

Generation-Based Fuzzing

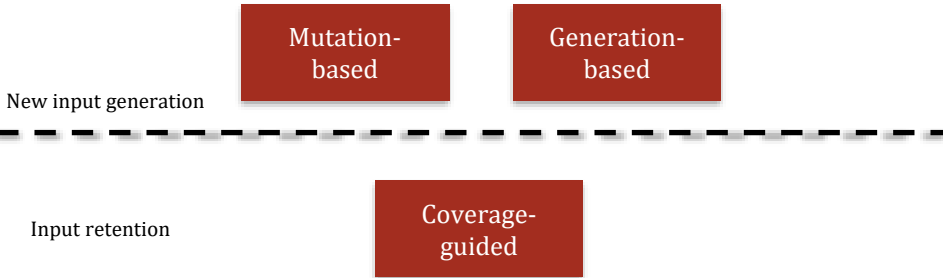
- Can be more accurate, but at a cost
- **Pros:** More complete search
 - Values more specific to the program operation
 - Can account for dependencies between inputs
- **Cons:** More work
 - Get the specification
 - Write the generator – ad hoc
- Need to do for each program

Grey Box Fuzzing

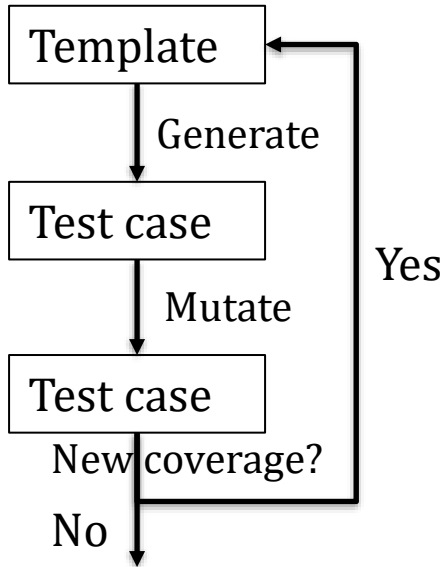
- Rather than treating the program as a black box, instrument the program to track the paths run
 - Also called **coverage-guided**
- Save inputs that lead to new paths
 - Associated with the paths they exercise
- Example
 - American Fuzzy Lop (AFL)
- “State of the practice” at this time



Relationship of fuzzing techniques



Syzkaller



```
1  # Copyright 2015 syzkaller project authors. All rights reserved.
2  # Use of this source code is governed by Apache 2 LICENSE that can be found
3
4  include <linux/kvm.h>
5  include <linux/kvm_host.h>
6  include <uapi/linux/fcntl.h>
7  include <asm/kvm.h>
8  include <asm/mce.h>
9
10 resource fd_kvm[fd]
11 resource fd_kvmvm[fd]
12 resource fd_kvmcpu[fd]
13 resource fd_kvmdev[fd]
14 resource fd_sgx_provision[fd]
15
16 openat$KVM(fd const[AT_FDCMD], file ptr[in, string["/dev/kvm"]], flags flag
17 openat$sgx_provision(fd const[AT_FDCMD], file ptr[in, string["/dev/sgx_prov
18
19 ioctl$KVM_CREATE_VM(fd fd_kvm, cmd const[KVM_CREATE_VM], type const[0]) fd
20 ioctl$KVM_GET_MSR_INDEX_LIST(fd fd_kvm, cmd const[KVM_GET_MSR_INDEX_LIST],
21 ioctl$KVM_CHECK_EXTENSION(fd fd_kvm, cmd const[KVM_CHECK_EXTENSION], arg ir
22 ioctl$KVM_GET_VCPU_MMAP_SIZE(fd fd_kvm, cmd const[KVM_GET_VCPU_MMAP_SIZE])
```

Sample template

AFL

- Provides compiler wrappers for gcc to instrument target program to collect fuzzing stats



AFL

- Provides compiler wrappers for gcc to instrument target program to collect fuzzing stats
- See
 - <http://lcamtuf.coredump.cx/afl/>

AFL Build

- Provides compiler wrappers for gcc to instrument target program to collect fuzzing stats
- Replace the gcc compiler in your build process with afl-gcc
- For example, in the Makefile
`CC=path-to/afl-gcc`
- Then build your target program with afl-gcc
 - Generates a binary instrumented for AFL fuzzing

AFL Use

- Provides compiler wrappers for gcc to instrument target program to collect fuzzing stats

- Run the fuzzer using afl-fuzz

```
afl-fuzz -i <input-dir> -o <output-dir> <path-to-bin>  
[args]
```

- `output` is the directory where the AFL results will be placed

AFL Demo

AFL Output

- Shows the results of the fuzzer
 - E.g., provides inputs that will cause the crash
- File “**fuzzer_stats**” provides summary of stats – UI
- File “**plot_data**” shows the progress of fuzzer
- Directory “**queue**” shows inputs that led to paths
- Directory “**crashes**” contains input that caused crash
- Directory “**hangs**” contains input that caused hang

AFL Operation

- How does AFL work?
 - http://lcamtuf.coredump.cx/afl/technical_details.txt
- Fuzzing strategies
 - Highly deterministic at first – bit flips, add/sub integer values, and choose interesting integer values
 - Then, non-deterministic choices – insertions, deletions, and combinations of test cases

Grey Box Fuzzing


- Finds flaws, but still does not understand the program
- **Pros:** Much better than black box testing
 - Essentially no configuration
 - Lots of crashes have been identified
- **Cons:** Still a bit of a stab in the dark
 - May not be able to execute some paths
 - Searches for inputs independently from the program
- Need to improve the effectiveness further

White Box Fuzzing

- Combines **test generation** with fuzzing
 - Test generation based on static analysis and/or symbolic execution – more later
 - Rather than generating new inputs and hoping that they enable a new path to be executed, compute inputs that will execute a desired path
 - And use them as fuzzing inputs
- **Goal:** Given program with a set of input parameters, generate a set of inputs that maximizes code coverage

Symbolic execution

```
void func(int a) {  
  if(a > 0) {  
    if(b == a) {  
      // bug is here  
    }  
  }  
}
```



$a > 0 \wedge b == a$  $a=1, b=1$
SMT solver

Take Away

- Goal is to discover vulnerabilities in our programs before adversaries exploit them
- One approach is dynamic testing of the program
 - Fuzz testing aims to achieve good program coverage with little effort for the programmer
 - Challenge is to generate the right inputs
- Black box (Mutational and generation), Grey box, and White box approaches are being investigated
 - AFL (Grey box) is now commonly used

Dynamic Analysis Limits

- Major advantage
 - When we produce a crash, it is a real crash
- Issue
 - However, may not be exploitable (i.e., not a vulnerability)
 - On the other hand, want to fix memory errors
 - But often not assertion failures
- Major limitation
 - We cannot find all vulnerabilities in a program
- Why not?

Questions

