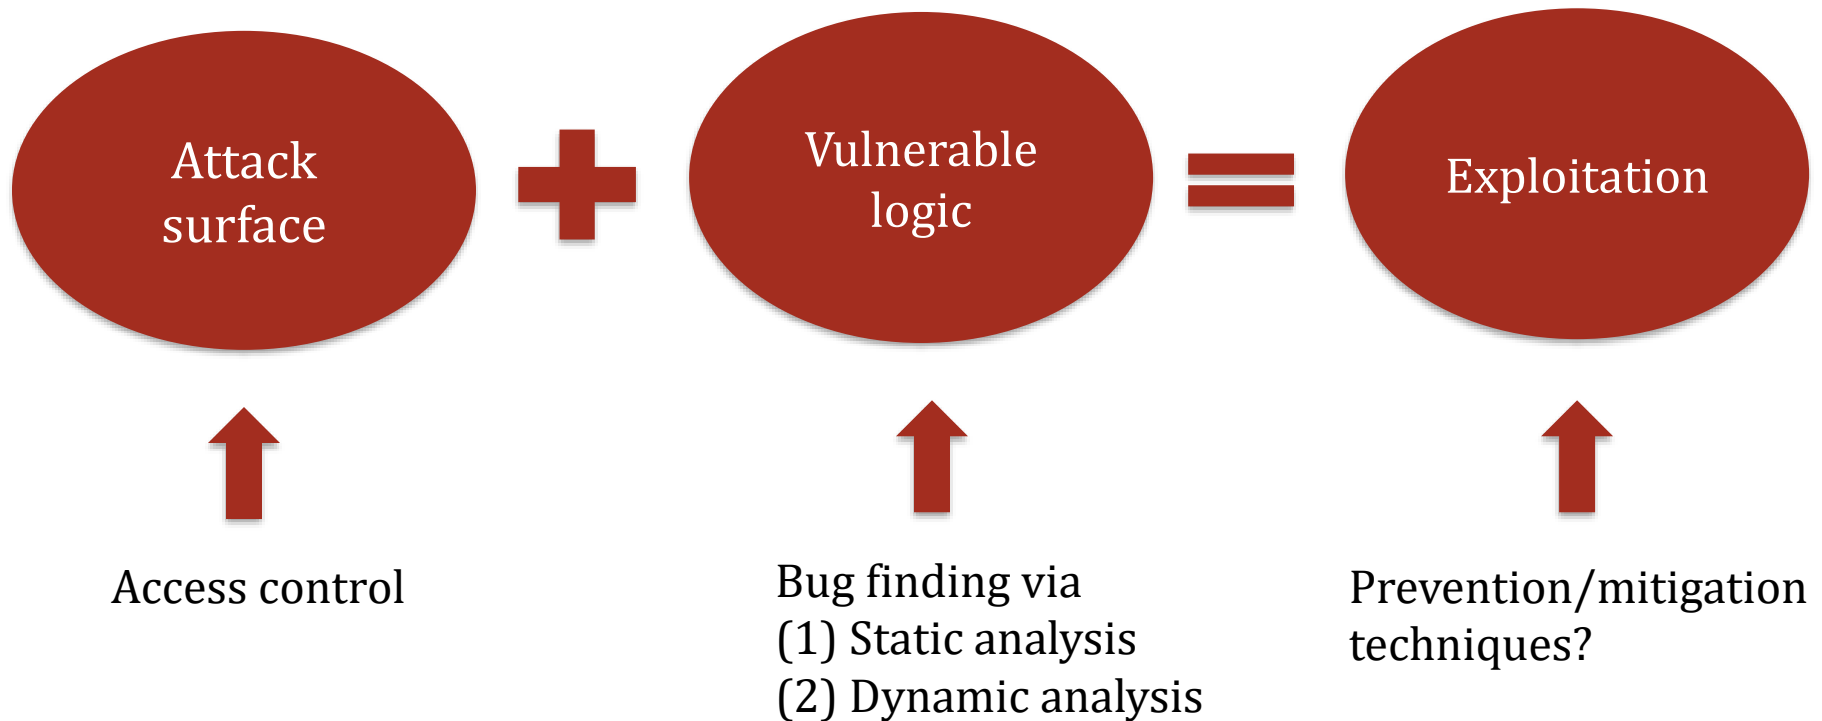


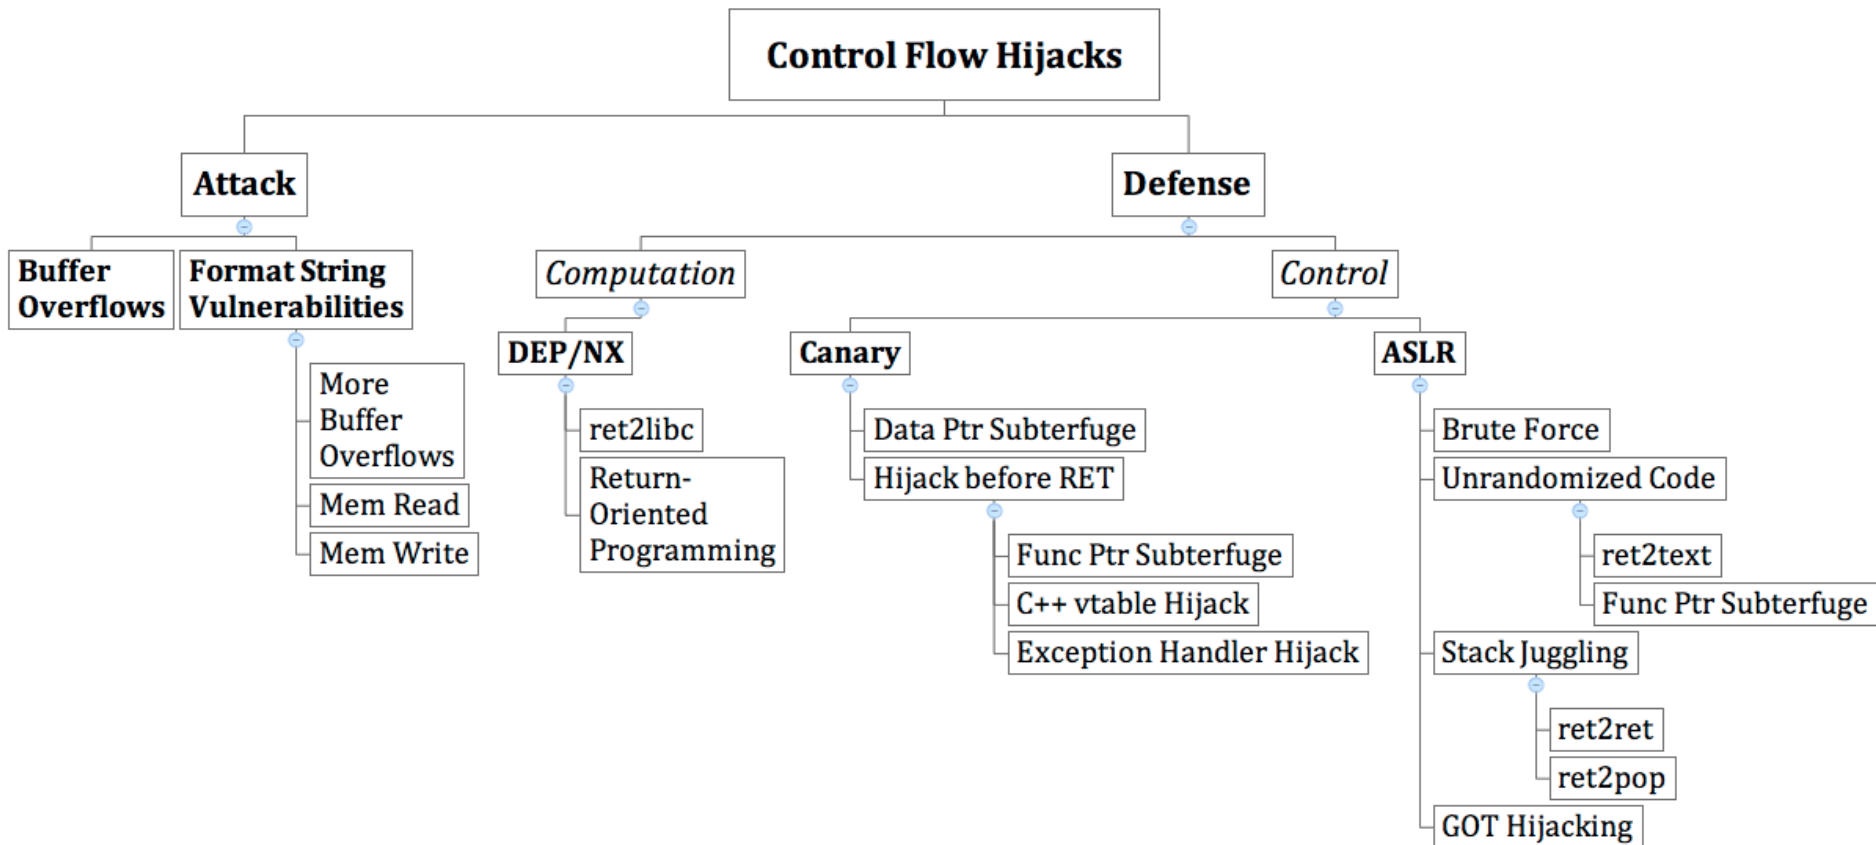
# CS165 – Computer Security

Control Flow Integrity and Software Fault Isolation  
Nov 16, 2021

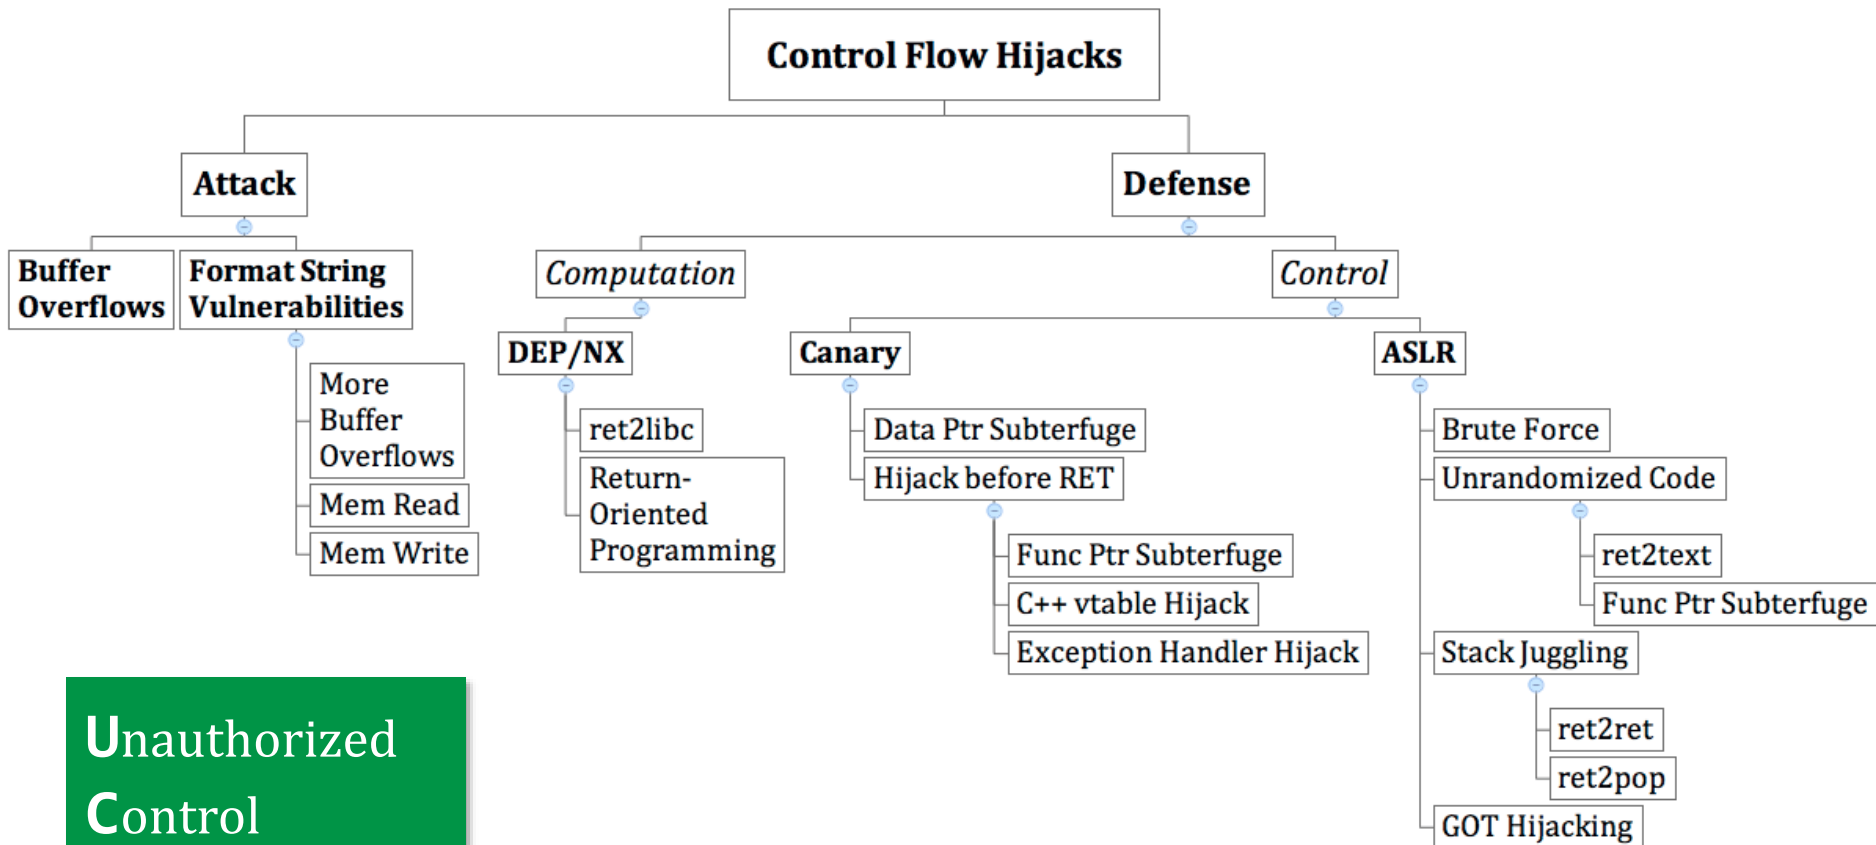
# Our story so far...



# Our story so far...



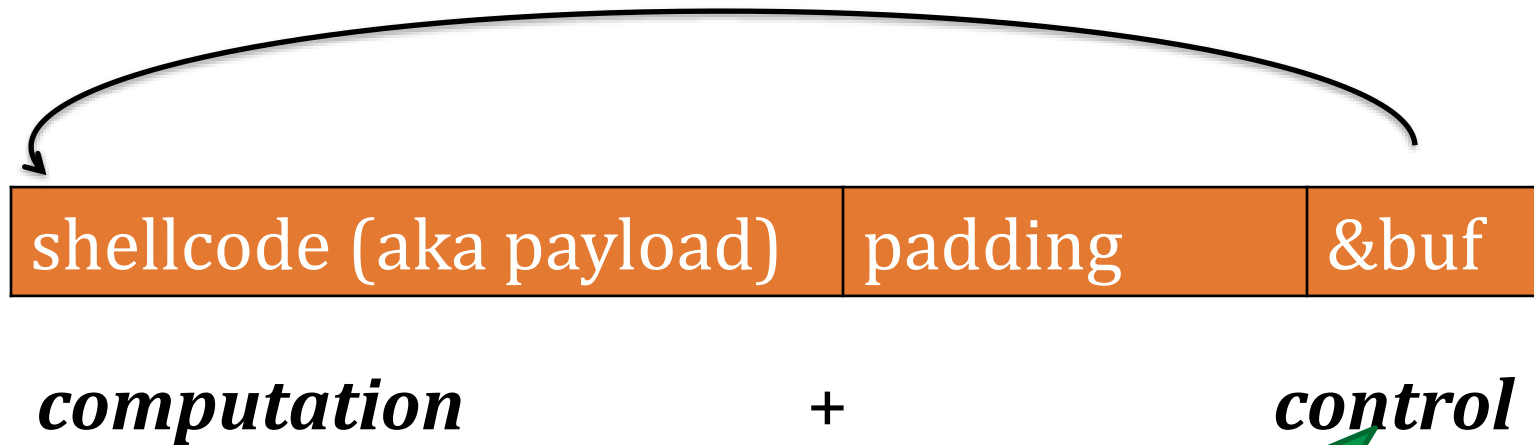
# Our story so far...



Unauthorized  
Control  
Information  
Tampering

# Control Flow Hijack:

## Always control + computation



Stack buffer overflow, return-to-libc  
ROP, functional pointer subterfuge:  
Hijacking the control

# Can we prevent control manipulation?

## Control Flow Integrity!

# Agenda


Reference Monitors

Control Flow Integrity

Software Fault Isolation

# Agenda

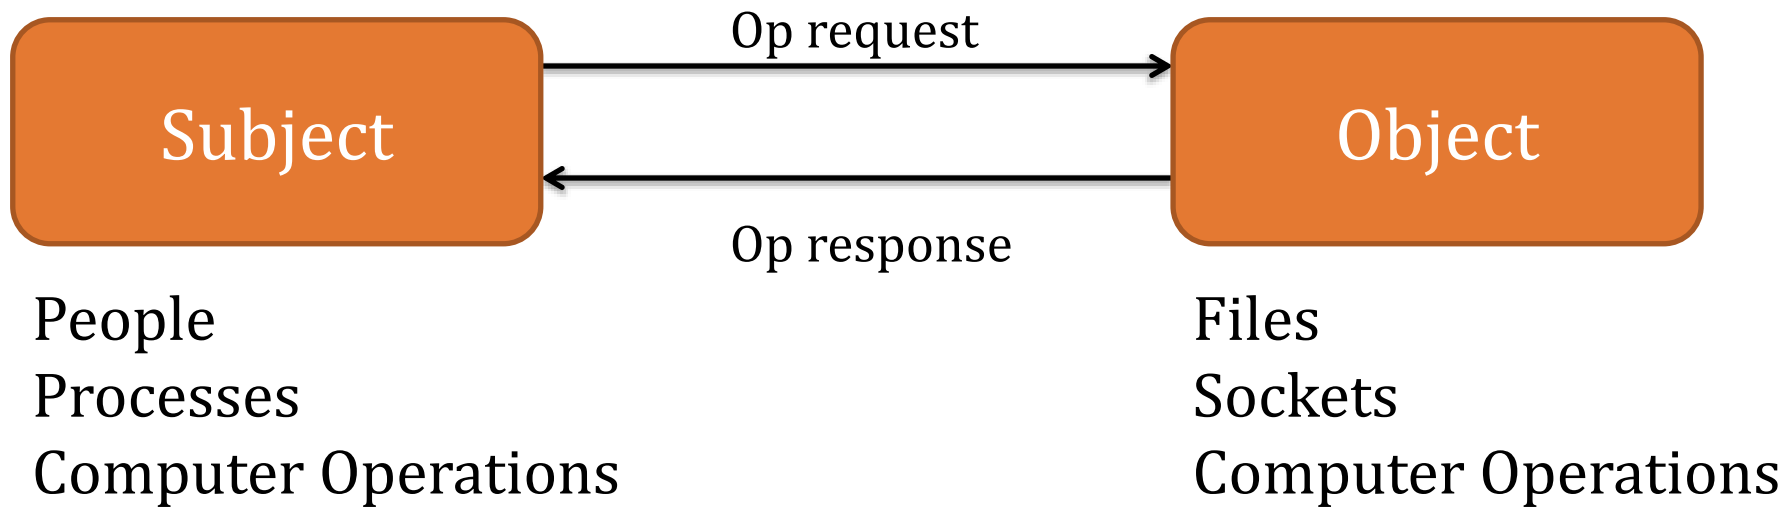
Reference Monitors




Control Flow Integrity

Software Fault Isolation







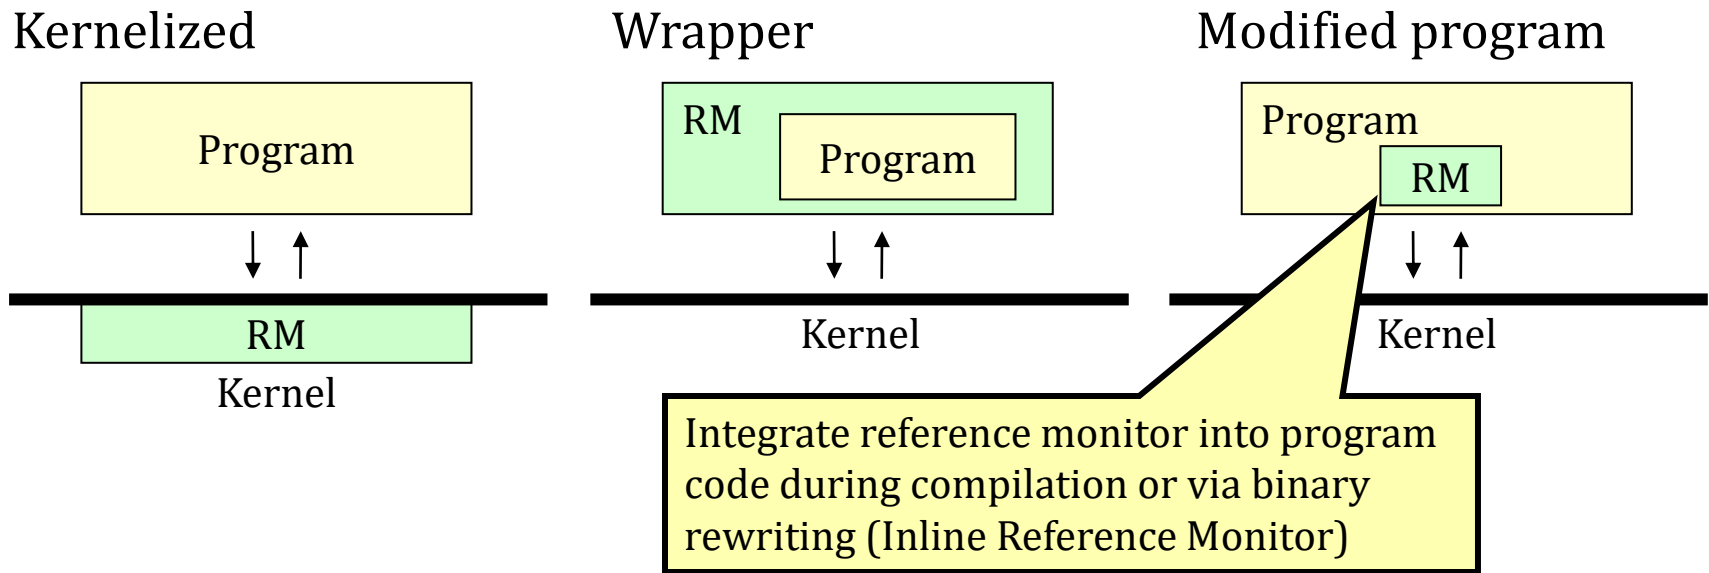
## Principles:

1. Complete Mediation: The reference monitor must always be invoked
2. Tamper-proof: The reference monitor cannot be changed by unauthorized subjects or objects
3. Verifiable: The reference monitor is small enough to thoroughly understand, test, and ultimately, verify.

# OS As a Reference Monitor

- OS enforces a variety of policies
  - File accesses are checked against file's **Access Control List** (ACL)
  - Process cannot write into memory of another process
  - Some operations require superuser privileges
    - But may need to switch back and forth (e.g., setuid in Unix)
  - Enforce CPU sharing, disk quotas, etc.

# Reference Monitor Implementation

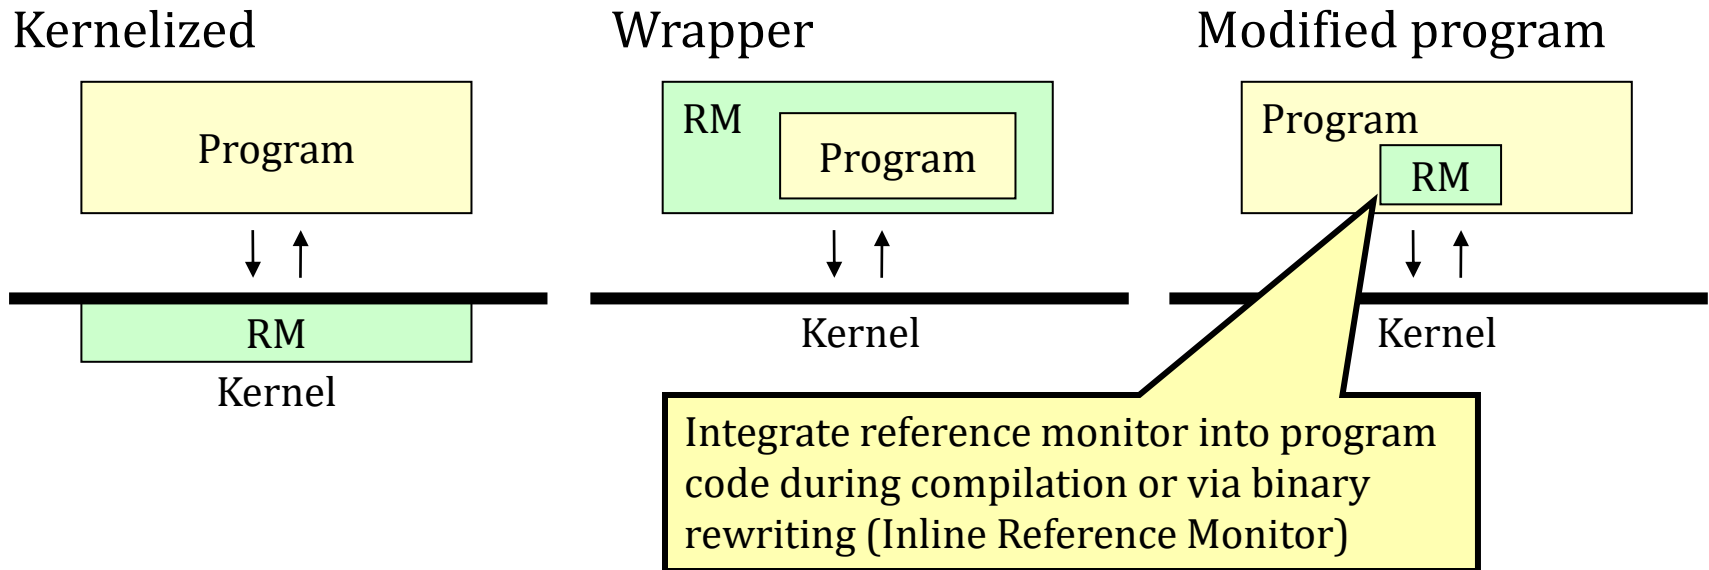


Today's Example: Inlining a control flow policy into a program

# What Makes a Process itself Safe?

- **Memory safety:** all memory accesses are “correct”
  - Respect array bounds, separation of code and data
- **Type safety:** all function calls and operations have arguments of correct type
- **Control-flow safety:** all control transfers are envisioned by the original program
  - No arbitrary jumps, no calls to library routines that the original program did not call

# Reference Monitor Implementation



- Policies can depend on application semantics
- Enforcement doesn't require context switches in the kernel
- Lower performance overhead

# Agenda

Reference Monitors



Control Flow Integrity

Software Fault Isolation

# Agenda

Reference Monitors



Control Flow Integrity



Software Fault Isolation



# Control Flow Integrity

## **Assigned Reading:**

*Control-Flow Integrity: Principles,  
Implementation and Applications*  
by Abadi, Budiu, Erlingsson, and Ligatti

# Control Flow Integrity

# Control Flow Integrity

- **protects against powerful adversary**
  - with full control over entire data memory

# Control Flow Integrity

- **protects against powerful adversary**
  - with full control over entire data memory
- **widely-applicable**
  - language-neutral; requires binary only

# Control Flow Integrity

- **protects against powerful adversary**
  - with full control over entire data memory
- **widely-applicable**
  - language-neutral; requires binary only
- **provably-correct & trustworthy**
  - formal semantics; small verifier

# Control Flow Integrity

- **protects against powerful adversary**
  - with full control over entire data memory
- **widely-applicable**
  - language-neutral; requires binary only
- **provably-correct & trustworthy**
  - formal semantics; small verifier
- **efficient**
  - hmm... 0-45% in experiments; average 16%

# Control Flow Integrity

- **protects against powerful adversary**
  - with full control over entire data memory
- **widely-applicable**
  - language-neutral; requires binary only
- **provably-correct & trustworthy**
  - formal semantics; small verifier
- **efficient**
  - hmm... 0-45% in experiments; average 16%

# CFI Adversary Model

## CAN

- Overwrite any data memory at any time
  - stack, heap, data segs
- Overwrite registers in current context



# CFI Adversary Model

## CAN

- Overwrite any data memory at any time
  - stack, heap, data segs
- Overwrite registers in current context

## CANNOT

- Execute Data
  - NX takes care of that
- Modify Code
  - text seg usually read-only
- Write to %ip
  - true in x86
- Overwrite registers in other contexts
  - kernel will restore regs


# CFI Overview

**Invariant:** Execution must follow a path in a control flow graph (CFG) created ahead of run time.

“static”

Most control flow transfer targets are hard-coded

```
jnz  short loc_18002C19E  0F 85 B4 00 00 00  
call  __scrt_initialize_crt  EB 12 0B 00 00
```



But what about indirect jumps and ret?

e.g., func pointer, ret address


# CFI Overview

## **Method to check indirect control transfers:**

- build CG and CFG statically, e.g., at compile time
  - **call, jmp, ret** instructions
- instrument (rewrite) binary, e.g., at install time
  - add IDs and ID checks; maintain ID uniqueness
- verify CFI instrumentation at load time
  - indirect jump targets, presence of IDs and ID checks, ID uniqueness
- perform ID checks at run time
  - indirect jumps have matching IDs

# Call Graph – Checking Return Address

```
void orange()    void red(int x)    void green()
{
1. red(1);      {
2. red(2);      green();
3. green();     ...
}              }
}
```



- Upon return (e.g., red)

ret



```
mov ecx, [esp] //check ret addr
cmp [ecx], AABBCDDh
jne error_label
jmp ecx+4
```

- At the caller (e.g., orange)


call  
...



```
0x00 call red
0x04 AABBCDDh #read-only data
0x08 ...
```


# Call Graph – Checking Return Address

```
void orange()    void red(int x)    void green()
{
1. red(1);      {
2. red(2);      green();
3. green();     ...
}              }
}
```




- Upon return (e.g., red)

- At the caller (e.g., orange)



# Call Graph – Checking Return Address

```
void orange()    void red(int x)    void green()
{
1. red(1);      {
2. red(2);      green();
3. green();     ...
}              }
}
```



- Upon return

ret

```
mov ecx, [esp] //check ret addr
cmp [ecx], AABBCDDh
jne error_label
jmp ecx+4
```

- At the caller (e.g., orange)

call  
...

```
0x00 call red
0x04 AABBCDDh #read-only data
0x08 ...
```

# Checking Function Pointer Deference

```
FF 53 08          call  [ebx+8]          ; call a function pointer
```

is instrumented using `prefetchnta` destination IDs, to become:

```
8B 43 08          mov  eax, [ebx+8]      ; load pointer into register
3E 81 78 04 78 56 34 12  cmp  [eax+4], 12345678h ; compare opcodes at destination
75 13             jne  error_label      ; if not ID value, then fail
FF D0            call  eax              ; call function pointer
3E 0F 18 05 DD CC BB AA prefetchnta [AABBCCDDh] ; label ID, used upon the return
```

# Checking Function Pointer Deference

```
FF 53 08          call  [ebx+8]          ; call a function pointer
```

is instrumented using `prefetchnta` destination IDs, to become:

```
8B 43 08          mov  eax, [ebx+8]      ; load pointer into register
3E 81 78 04 78 56 34 12  cmp  [eax+4], 12345678h ; compare opcodes at destination
75 13             jne  error_label      ; if not ID value, then fail
FF D0            call  eax              ; call function pointer
3E 0F 18 05 DD CC BB AA prefetchnta [AABBCCDDh] ; label ID, used upon the return
```



Check dest label



# Performance

**Size:** increase 8% avg

**Time:** increase 0-45%; 16% avg

– I/O latency helps hide overhead





Fig. 6. Execution overhead of inlined CFI enforcement on SPEC2000 benchmarks.

# Context-Sensitive CFI

- Previous assumption: destination is fixed (a single target or a group)

```
void orange()    void red(int x)    void green()
{
1. red(1);      {
2. red(2);      ..
3. green();     }
}              }
```



A more precise CFI for orange lines 1 and 2.

# Context-Sensitive CFI

Whether different calling contexts are distinguished

<pre>void orange() { 1. red(1); 2. red(2); 3. green(); }</pre>	<pre>void red(int x) { .. }</pre>	<pre>void green() {     green();     yellow(); }</pre>
--	---	--



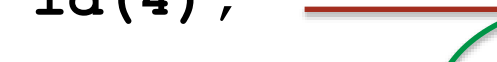
Context sensitive  
distinguishes 2 different  
calls to red(-)

# Context Sensitive Example

```
a = id(4);
```

```
void id(int z)
{ return z; }
```

```
b = id(5);
```



## Context-Sensitive

(color denotes matching call/ret)

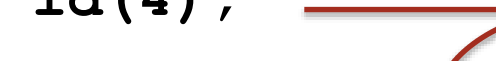
Context sensitive can tell one call returns 4, the other 5

```
a = id(4);
```

```
void id(int z)
```

```
{ return z; }
```

```
b = id(5);
```



## Context-Insensitive (note merging)

Context insensitive will say both calls return  $\{4,5\}$

# Context Sensitivity Problems

Suppose A and B both call C.

- CFI uses same return label in A and B.

How to prevent C from returning to B when it was called from A?

- **Solution: Shadow Call Stack**
  - a protected memory region for call stack
  - each call/ret instrumented to update shadow
  - CFI ensures instrumented checks will be run

# Security Guarantees

Effective against attacks based on illegitimate control-flow transfer

- buffer overflow, ret2libc, ROP, pointer subterfuge, etc.

Any check becomes non-circumventable.

# Security Guarantees

Effective against attacks based on illegitimate control-flow transfer

- buffer overflow, ret2libc, ROP, pointer subterfuge, etc.

Any check becomes non-circumventable.

Allow data-only attacks since they respect CFG!

- incorrect usage (e.g. printf can still dump mem)
- substitution of data (e.g. replace file names)

# CFI an active area of research

CCS 2015:

- Per-Input Control-Flow Integrity
- Practical Context-Sensitive CFI
- CCFI: Cryptographically Enforced Control Flow Integrity
- Losing Control: On the Effectiveness of Control-Flow Integrity under Stack Attacks



# CFI Summary

Control Flow Integrity ensures that control flow follows a path in CFG

- Accuracy of CFG determines level of enforcement
- Can build other security policies on top of CFI
- Simple version now deployed in Windows 10 (a slow but continuing trend)

# Agenda

Reference Monitors



Control Flow Integrity



Software Fault Isolation

# Agenda


Reference Monitors



Control Flow Integrity



Software Fault Isolation



# Software Fault Isolation

## **Optional Reading:**

*Efficient Software-Based Fault Isolation*

by Wahbe, Lucco, Anderson, Graham

# Motivation: Running untrusted code

- We often need to run buggy/untrusted code:
  - programs from untrusted Internet sites:
    - toolbars, viewers, codecs for media player
  - old or insecure applications: ghostview, outlook
  - legacy daemons: sendmail, bind
  - honeypots
- Goal: if application “misbehaves,” kill it

# Isolation Mechanisms

# Isolation Mechanisms

- Hardware
  - Memory Protection (virtual address translation, x86 segmentation)

# Isolation Mechanisms

- Hardware
  - Memory Protection (virtual address translation, x86 segmentation)
- Software
  - Sandboxing
  - Language-Based



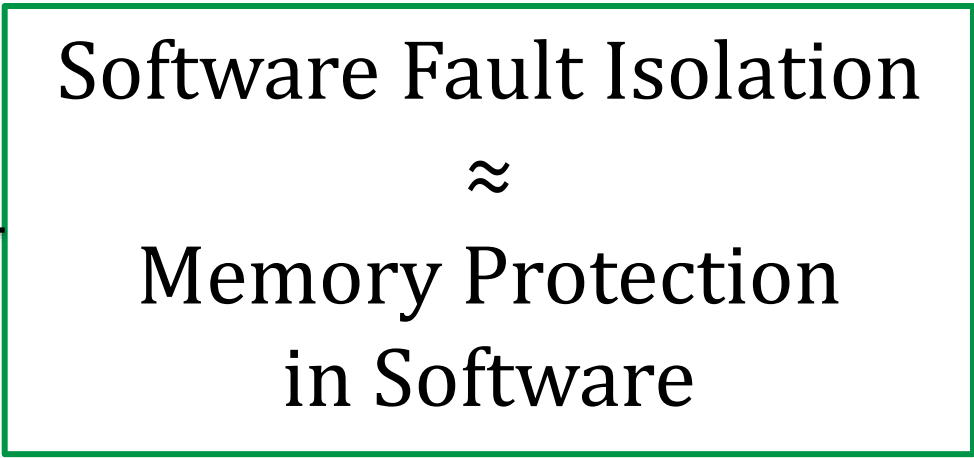
# Isolation Mechanisms

- Hardware
  - Memory Protection (virtual address translation, x86 segmentation)
- Software
  - Sandboxing
  - Language-Based
- Hardware + Software
  - Virtual machines

# Isolation Mechanisms

- Hardware
  - Memory Protection (virtual address translation, x86 segmentation)
- Software
  - Sandboxing
  - Language-Based
- Hardware + Software
  - Virtual machines

Software Fault Isolation  
 $\approx$   
Memory Protection  
in Software



# Software Fault Isolation

- SFI ensures that a module only accesses memory within its region by adding *checks* (also a type of Inline Reference Monitor)
  - e.g., a plugin can access only its own memory

```
if (module_lower < x < module_upper)
    z = load[x];
```



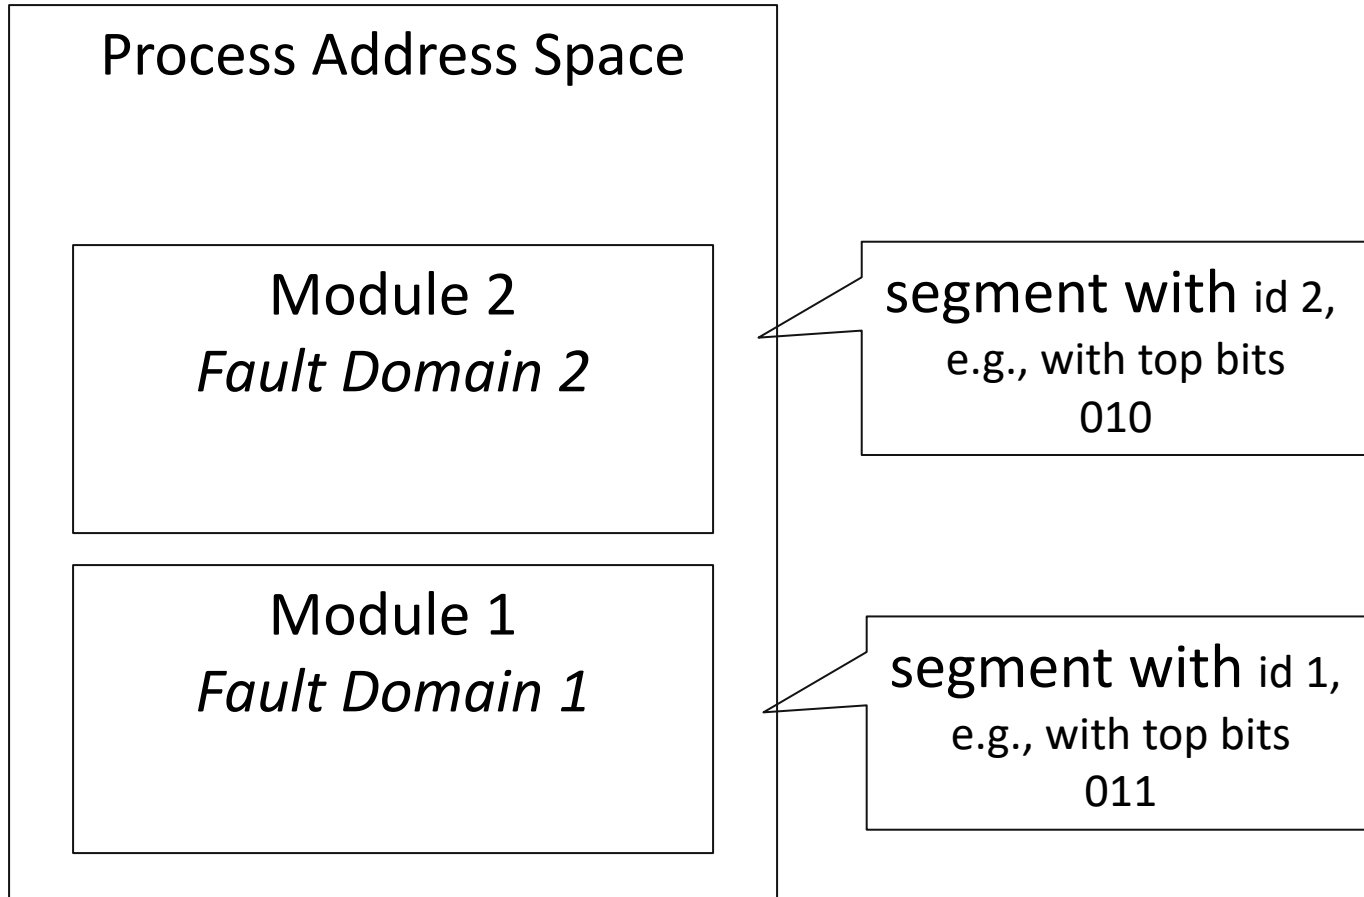
SFI Check

- CFI as a building block can ensure inserted memory checks are executed

# SFI Goals

- Confine faults inside distrusted extensions
  - codec shouldn't compromise media player
  - device driver shouldn't compromise kernel
  - plugin shouldn't compromise web browser
- Allow for efficient cross-domain calls
  - numerous calls between media player and codec
  - numerous calls between device driver and kernel

# Main Idea



# SFI Example

```
int compute_sum( int a[], int len )
{
    int sum = 0;
    for(int i = 0; i < len; ++i) {
        sum += a[i];
    }
    return sum;
}
```

```
...
mov    ecx, 0h                ; int i = 0
mov    esi, [esp+8]           ; a[] base ptr
LOOP:  and    esi, 20FFFFFFh    ; SFI masking
      add    eax, [esi+ecx*4]   ; sum += a[i]
      inc    ecx               ; ++i
      cmp    ecx, edx          ; i < len
      jl     LOOP
```

# Optimizing SFI using CFI

```
int compute_sum( int a[], int len )
{
    int sum = 0;
    for(int i = 0; i < len; ++i) {
        sum += a[i];
    }
    return sum;
}
```

```
...
mov     ecx, 0h                ; int i = 0
mov     esi, [esp+8]           ; a[] base ptr
and     esi, 20FFFFFFh        ; SFI masking
LOOP:   add     eax, [esi+ecx*4] ; sum += a[i]
        inc     ecx            ; ++i
        cmp     ecx, edx        ; i < len
        jl      LOOP
```

# Agenda

Reference Monitors



Control Flow Integrity



Software Fault Isolation





# Questions

